

# CS101 Algorithms and Data Structures

Graphs

Textbook Ch B.4, B.5.1, 22.1

# Outline

- Definitions
  - Undirected graphs
  - Directed graph
- Representation
  - Adjacency matrix
  - Adjacency list

# Undirected Graphs

We will define an Undirected Graph ADT as a collection of *vertices*

$$V = \{v_1, v_2, \dots, v_n\}$$

- The number of vertices is denoted by

$$|V| = n$$


- Associated with this is a collection  $E$  of unordered pairs  $\{v_i, v_j\}$  termed *edges* which connect the vertices

# Undirected Graphs

Consider this collection of vertices

$$V = \{v_1, v_2, \dots, v_9\}$$

where  $|V| = n = 9$




# Undirected graphs

Associated with these vertices are  $|E| = 5$  edges

$$E = \{\{v_1, v_2\}, \{v_3, v_5\}, \{v_4, v_8\}, \{v_4, v_9\}, \{v_6, v_9\}\}$$

- The pair  $\{v_j, v_k\}$  indicates that both vertex  $v_j$  is adjacent to vertex  $v_k$  and vertex  $v_k$  is adjacent to vertex  $v_j$



# Undirected graphs

We will assume that a vertex is never adjacent to itself

- For example,  $\{v_1, v_1\}$  will not define an edge

The maximum number of edges in an undirected graph is

$$|E| \leq \binom{|V|}{2} = \frac{|V|(|V|-1)}{2} = O(|V|^2)$$


# An undirected graph

Example: given the  $|V| = 7$  vertices

$$V = \{A, B, C, D, E, F, G\}$$

and the  $|E| = 9$  edges

$$E = \{\{A, B\}, \{A, D\}, \{A, E\}, \{B, C\}, \{B, D\}, \{B, E\}, \{C, E\}, \{C, F\}, \{D, E\}\}$$



# Degree


The degree of a vertex is defined as the number of adjacent vertices

$$\text{degree}(A) = \text{degree}(D) = \text{degree}(C) = 3$$

$$\text{degree}(B) = \text{degree}(E) = 4$$

$$\text{degree}(F) = 1$$

$$\text{degree}(G) = 0$$





Those vertices adjacent to a given vertex are its *neighbors*





# Sub-graphs

A *sub-graph* of a graph contains a **subset** of the vertices and a subset of the edges that connect the **subset** of the vertices in the original graph





# Sub-graphs

A *sub-graph* of a graph contains a subset of the vertices and a subset of the edges that connect the subset of the vertices in the original graph



# Vertex-induced sub-graphs

A *vertex-induced sub-graph* contains a subset of the vertices and all the edges in the original graph between those vertices



# Paths

A path in an undirected graph is an ordered sequence of vertices

$(v_0, v_1, v_2, \dots, v_k)$


where  $\{v_{j-1}, v_j\}$  is an edge for  $j = 1, \dots, k$

- Termed *a path from  $v_0$  to  $v_k$*
- The length of this path is  $k$

# Paths

A path of length 4:


(A, B, E, C, F)



# Paths

A path of length 5:


(A, B, E, C, B, D)



# Paths


*A trivial path of length 0:*

(A)




# Simple path

A *simple path* has no repetitions (other than perhaps the first and last vertices)



(A, B, E, C, F)




(A, B, E, C, B, D)



# Simple cycle

A *simple cycle* is a simple path of **at least two vertices** with the first and last vertices equal




(A, B, D, A)


# Connectedness

Two vertices  $v_i, v_j$  are said to be *connected* if there exists a path from  $v_i$  to  $v_j$

A graph is connected if there exists a path between any two vertices



A connected graph




An unconnected graph

# Weighted graphs

A weight may be associated with each edge in a graph


- This could represent distance, energy consumption, cost, etc.
- Such a graph is called a *weighted graph*



# Weighted graphs

The *length* of a path within a weighted graph is the sum of all of the edges which make up the path


- The length of the path (A, D, G) in the following graph is  $5.1 + 3.7 = 8.8$



# Weighted graphs

Different paths may have different weights


- Another path is (A, C, F, G) with length  $1.2 + 1.4 + 4.5 = 7.1$



# Weighted graphs

Problem: find the shortest path between two vertices


- Here, the shortest path from A to G is (A, C, F, D, E, G) with length 5.7



# Trees

A graph is a tree if it is connected and there is a unique path between any two vertices

- Example: three trees on the same eight vertices



Properties:

- The number of edges is  $|E| = |V| - 1$
- The graph is *acyclic*, that is, it does not contain any cycles
- Adding one more edge must create a cycle
- Removing any one edge creates two unconnected sub-graphs


# Trees

Any tree can be converted into a rooted tree by:

- Choosing any vertex to be the root
- Defining its neighboring vertices as its children

and then recursively defining:

- All neighboring vertices other than that one designated its parent to be its children





# Forests


A forest is any graph that has no cycles

Consequences:

- The number of edges is  $|E| < |V|$
- The number of trees is  $|V| - |E|$
- Removing any one edge adds one more tree to the forest

Here is a forest with 22 vertices and 18 edges

- There are four trees



# Outline

- Definitions
  - Undirected graphs
  - Directed graph
- Representation
  - Adjacency matrix
  - Adjacency list

# Directed graphs

In a *directed graph*, the **edges** on a graph are be **associated with a direction**

- Edges are ordered pairs  $(v_j, v_k)$  denoting a connection from  $v_j$  to  $v_k$
- The edge  $(v_j, v_k)$  is different from the edge  $(v_k, v_j)$

Streets are directed graphs:


- In most cases, you can go two ways unless it is a one-way street

# Directed graphs

Given a graph of nine vertices  $V = \{v_1, v_2, \dots, v_9\}$

- These six pairs  $(v_j, v_k)$  are *directed edges*

$$E = \{(v_1, v_2), (v_3, v_5), (v_5, v_3), (v_6, v_9), (v_8, v_4), (v_9, v_4)\}$$



# Directed graphs

The maximum number of directed edges in a directed graph is

$$|E| \leq 2 \binom{|V|}{2} = 2 \frac{|V|(|V|-1)}{2} = |V|(|V|-1) = O(|V|^2)$$

# In and out degrees


The degree of a vertex in a directed graph:

- The *out-degree* of a vertex is the number of outward edges from the vertex
- The *in-degree* of a vertex is the number of inward edges to the vertex

In this graph:

$$\text{in\_degree}(v_1) = 0 \quad \text{out\_degree}(v_1) = 2$$

$$\text{in\_degree}(v_5) = 2 \quad \text{out\_degree}(v_5) = 3$$




# Sources and sinks

Definitions:

- Vertices with an in-degree of zero are described as *sources*
- Vertices with an out-degree of zero are described as *sinks*

In this graph:

- Sources:  $v_1, v_6, v_7$
- Sinks:  $v_2, v_9$



# Paths

A path in a directed graph is an ordered sequence of vertices

$$(v_0, v_1, v_2, \dots, v_k)$$


where  $(v_{j-1}, v_j)$  is an edge for  $j = 1, \dots, k$

A path of length 5 in this graph is

$$(v_1, v_4, v_5, v_3, v_5, v_2)$$

A simple cycle of length 3 is

$$(v_8, v_4, v_5, v_8)$$






# Connectedness

Two vertices  $v_j, v_k$  are said to be *connected* if there exists a path from  $v_j$  to  $v_k$

- A graph is *strongly connected* if there exists a directed path between any two vertices
- A graph is *weakly connected* there exists a path between any two vertices that ignores the direction

In this graph:


- The sub-graph  $\{v_3, v_4, v_5, v_8\}$  is strongly connected
- The sub-graph  $\{v_1, v_2, v_3, v_4, v_5, v_8\}$  is weakly connected



# Weighted directed graphs

In a weighted directed graphs, each edge is associated with a value

If both  $(v_j, v_k)$  and  $(v_k, v_j)$  are edges, it is not required that they have the same weight



# Outline

- Definitions
  - Undirected graphs
  - Directed graph
- Representation
  - Adjacency matrix
  - Adjacency list

# The Graph ADT

The Graph ADT describes a container storing an adjacency relation

- Queries include:
  - The number of vertices
  - The number of edges
  - List the vertices adjacent to a given vertex
  - Are two vertices adjacent?
  - Are two vertices connected?
- Modifications include:
  - Inserting or removing an edge
  - Inserting or removing a vertex (and all edges containing that vertex)

The run-time of these operations will depend on the representation


# Binary-relation list

The most inefficient is a relation list:

- A container storing the edges

$\{(1, 2), (1, 4), (3, 5), (4, 2), (4, 5), (5, 2), (5, 3), (5, 8), (6, 9), (7, 9), (8, 4)\}$

- Requires  $\Theta(|E|)$  memory
- Determining if  $v_j$  is adjacent to  $v_k$  is  $O(|E|)$
- Finding all neighbors of  $v_j$  is  $\Theta(|E|)$



# Adjacency matrix

Requiring more memory but also faster, an adjacency matrix


- The matrix entry  $(j, k)$  is set to true if there is an edge  $(v_j, v_k)$

|   | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| 1 |   | T |   | T |   |   |   |   |   |
| 2 |   |   |   |   |   |   |   |   |   |
| 3 |   |   |   |   | T |   |   |   |   |
| 4 |   | T |   |   | T |   |   |   |   |
| 5 |   | T | T |   |   |   |   | T |   |
| 6 |   |   |   |   |   |   |   |   | T |
| 7 |   |   |   |   |   |   |   | T |   |
| 8 |   |   |   |   |   |   |   |   |   |
| 9 |   |   |   |   |   |   |   |   |   |

Requires  $\Theta(|V|^2)$  memory

- Determining if  $v_j$  is adjacent to  $v_k$  is  $O(1)$

- Finding all neighbors of  $v_j$  is  $\Theta(|V|)$



# Adjacency list

Most efficient for algorithms is an adjacency list

- Each vertex is associated with a list of its neighbors

```

1  • → 2 → 4
2  •
3  • → 5
4  • → 2 → 5
5  • → 2 → 3 → 8
6  • → 9
7  • → 9
8  • → 4
9  •
    
```

- Requires  $\Theta(|V| + |E|)$  memory


- On average:

- Determining if  $v_j$  is adjacent to  $v_k$  is

$$O\left(\frac{|E|}{|V|}\right)$$

- Finding all neighbors of  $v_j$  is

$$\Theta\left(\frac{|E|}{|V|}\right)$$



# Adjacency Matrix

A graph of  $n$  vertices may have up to

$$\binom{n}{2} = \frac{n(n-1)}{2} = \mathbf{O}(n^2)$$

edges

The first straight-forward implementation is an adjacency matrix




# Adjacency Matrix

Define an  $n \times n$  matrix  $\mathbf{A} = (a_{ij})$  and if the vertices  $v_i$  and  $v_j$  are connected with weight  $w$ , then set  $a_{ij} = w$  and  $a_{ji} = w$

That is, the matrix is symmetric, e.g.,

|   | 0    | 1    | 2    | 3    | 4    | 5    |
|---|------|------|------|------|------|------|
| 0 |      | 0.83 |      | 0.65 |      | 0.51 |
| 1 | 0.83 |      | 0.67 | 0.72 | 0.46 | 0.90 |
| 2 |      | 0.67 |      | 0.39 |      |      |
| 3 | 0.65 | 0.72 | 0.39 |      | 0.24 |      |
| 4 |      | 0.46 |      | 0.24 |      |      |
| 5 | 0.51 | 0.90 |      |      |      |      |




# Adjacency Matrix

An unweighted graph may be saved as an array of Boolean values

- vertices  $v_i$  and  $v_j$  are connected then set

$$a_{ij} = a_{ji} = \text{true}$$


|   | 0 | 1 | 2 | 3 | 4 | 5 |
|---|---|---|---|---|---|---|
| 0 |   | T | F | T | F | T |
| 1 | T |   | T | T | T | T |
| 2 | F | T |   | T | F | F |
| 3 | T | T | T |   | T | F |
| 4 | F | T | F | T |   | F |
| 5 | T | T | F | F | F |   |



# Adjacency Matrix

If the graph was directed, then the matrix would not necessarily be symmetric

|   | 0    | 1    | 2 | 3    | 4    | 5    |
|---|------|------|---|------|------|------|
| 0 |      |      |   | 0.65 |      |      |
| 1 | 0.83 |      |   | 0.72 | 0.46 | 0.90 |
| 2 |      | 0.67 |   | 0.39 |      |      |
| 3 |      |      |   |      | 0.24 |      |
| 4 |      |      |   | 0.24 |      |      |
| 5 | 0.51 |      |   |      |      |      |



# Default Values

Question: what do we do about vertices which are not connected?

- the value 0
- a negative number, e.g.,  $-1$
- positive infinity:  $\infty$

The last is the most logical, in that it makes sense that two vertices which are not connected have an infinite distance between them

The distance from a node to itself is  $0$



# Sparse Matrices

- The memory required for creating an  $n \times n$  matrix using a 2D array is  $\Theta(n^2)$  bytes
- This could potentially waste a significant amount of memory:
  - Consider a friendship graph: nodes represent persons and edges represent friendship
  - The world population is 7.4 billion  $\Rightarrow$  the size of the matrix is  $(7.4 \times 10^9)^2 \approx 55 \times 10^{18}$
  - However, each person on average has, say, 100 friends. Hence only  $\frac{100}{7.4 \times 10^9}$  of the matrix elements are true. The other elements are the default value: false.

# Adjacency list

- For an undirected graph, use an array of linked lists to store edges
  - Each vertex has a linked list that stores all the edges connected to the vertex
  - Each node in a linked list must store two items of information: the connecting vertex and the weight

# Adjacency list





# Adjacency list

- To store a **directed graph**
  - Each vertex has a linked list that stores all the edges originated from the vertex
  - Each node in a linked list stores two items of information: the vertex that the edge connects to, the weight



# Adjacency list



# Summary

- Definitions
  - Undirected graphs
  - Directed graph
- Representation
  - Adjacency matrix
  - Adjacency list