



CS120: Computer Networks

Lecture 19. Other Topics in Transportation Layer


Zhice Yang

Outline

- TCP Fairness
- QoS
- QUIC

Evaluation Criteria

- Defining fairness is hard
 - In terms of a host, a TCP link, or an application ?
- TCP fairness goal: if **n** TCP sessions share same bottleneck link of bandwidth **R**, each should have average rate of ***R/n***



- Fairness Index



$$\bullet f(x_1 \dots x_n) = \frac{(\sum x_i)^2}{n * \sum x_i^2}$$

Fairness in TCP

- Consider the steady state, TCP uses a (linear) scheme to adjust its window cwnd
 - $\text{cwnd}' = b * \text{cwnd} + a$
- Possible Designs
 - Additive increase, additive decrease
 - Additive increase, multiplicative decrease (AIMD)
 - Multiplicative increase, additive decrease
 - Multiplicative increase, multiplicative decrease


Fairness in TCP

- Consider a case with two TCP connections



Fairness in TCP


- Consider a case with two TCP connections



Increase the throughput to
utilize more bandwidth


Fairness in TCP

- Consider a case with two TCP connections



Fairness in TCP


- Consider a case with two TCP connections



Decrease the throughput
to avoid congestion


Fairness in TCP

- Consider a case with two TCP connections




Fairness in TCP

- Consider a case with two TCP connections
 - Behavior of additive increase additive decrease
 - Stable but not fair




Fairness in TCP

- Consider a case with two TCP connections
 - Behavior of multiplicative increase multiplicative decrease
 - Stable but not fair




Fairness in TCP

- Consider a case with two TCP connections
 - Behavior of multiplicative increase additive decrease
 - Not stable



Fairness in TCP

- Consider a case with two TCP connections
 - Behavior of AIMD
 - Stable and fair



Fairness and RTT

- TCP connection with smaller RTT occupies more bandwidth
 - When congestion happens, they recover more quickly
 - TCP adjust cwnd in RTT basis

Fairness and Parallel TCP Connections

- Application can open multiple parallel connections between two hosts
 - web browsers do this , e.g., link of rate R with 9 existing connections:
 - new app asks for 1 TCP, gets rate $R/10$
 - new app asks for 11 TCPs, gets $R/2$

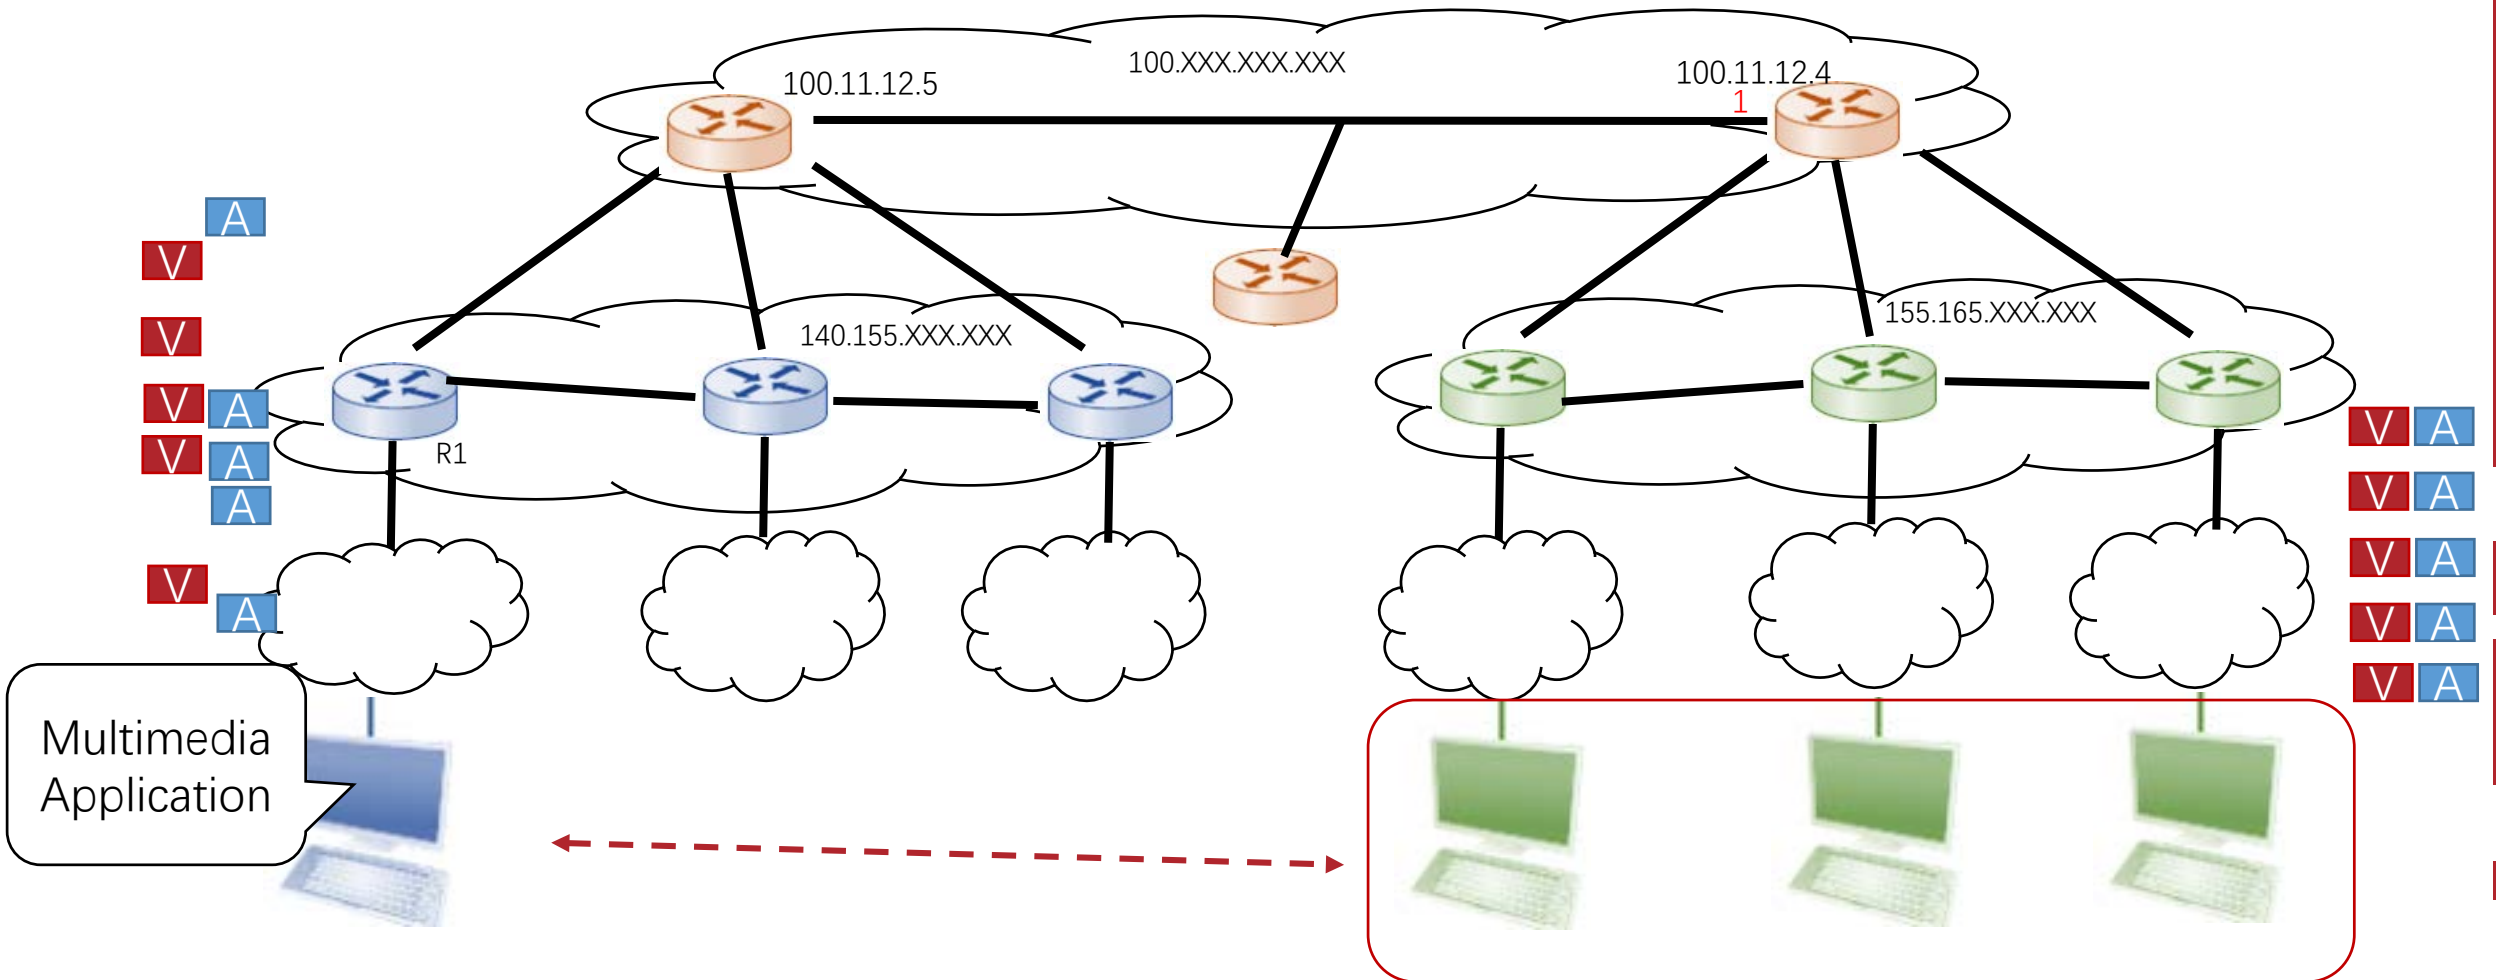
Fairness and UDP

- Some apps do not use TCP
 - do not want rate throttled by congestion control
- Instead, use UDP:
 - send audio/video at constant rate, tolerate packet loss
- There is no “Internet police” policing use of congestion control

Outline

- TCP Fairness
 - QoS
- QUIC

Realtime Application

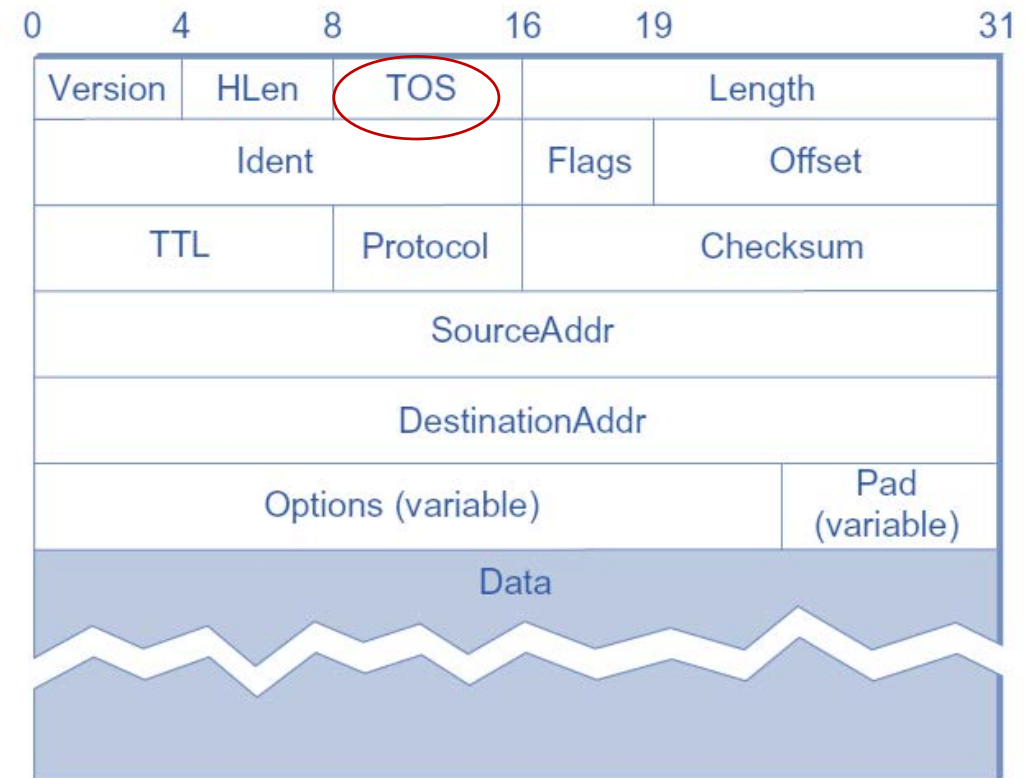


Packet Classification

- Classify Packets into Flows according to
 - Protocol
 - Source Address
 - Destination Address
 - Source Port
 - Destination Port
 - ToS

Differentiated Service Code Point

- Reuse ToS Field
 - 0-5bit: Differentiated Service Code Point (DSCP) Field
 - 6-7bit: Explicit Congestion Notification
- DSCP field encodes Per-Hop Behavior
 - Expedited Forwarding (all packets receive minimal delay & loss)
 - Assured Forwarding (packets marked with low/high drop probabilities)



Set Packet Class



- DSCP Field in Practice
 - Edge Routers
 - Set Differentiated Service (DS) Field in IP header
 - e.g., because of subscription
 - Core Routers
 - Implement Per Hop Behavior
 - According to DS Field of packets

Commonly used DSCP values

DSCP value	Hex value	Decimal value	Meaning	Drop probability	Equivalent IP precedence value
101 110	0x2e	46	Expedited forwarding (EF)	N/A	101 Critical
000 000	0x00	0	Best effort	N/A	000 - Routine
001 010	0x0a	10	AF11	Low	001 - Priority
001 100	0x0c	12	AF12	Medium	001 - Priority
001 110	0x0e	14	AF13	High	001 - Priority
010 010	0x12	18	AF21	Low	010 - Immediate
010 100	0x14	20	AF22	Medium	010 - Immediate
010 110	0x16	22	AF23	High	010 - Immediate
011 010	0x1a	26	AF31	Low	011 - Flash
011 100	0x1c	28	AF32	Medium	011 - Flash
011 110	0x1e	30	AF33	High	011 - Flash
100 010	0x22	34	AF41	Low	100 - Flash override
100 100	0x24	36	AF42	Medium	100 - Flash override
100 110	0x26	38	AF43	High	100 - Flash override


Queuing Discipline

- First-In-First-Out (FIFO)




Queuing Discipline

- Problems in FIFO
 - Too simple to provide resource allocation policies
 - e.g., yellow src does not follow congestion control (UDP), can occupy more network source



Queuing Discipline


- First-In-First-Out (FIFO) with Priority



Always dequeue the queue with higher priority
unless it is empty
Problem: starvation


Queuing Discipline

- Fair Queuing (FQ)
 - Each flow gets 1/4 output bandwidth




Queuing Discipline

- Fair Queuing (FQ)
 - Each flow gets 1/3 output bandwidth




Queuing Discipline

- Weighted Fair Queuing
 - Flows with higher weight get more output bandwidth ($2/7$, $3/7$, $1/7$, $1/7$)




Queuing Discipline

- Bit-Level Fair Queuing
 - Schedule according to earliest finish time
 - Finish time of packet i : $F_i = \max(F_{i-1}, A_i) + P_i$
 - P_i is the transmitting duration of packet i , A_i is the arriving time of packet i , F_{i-1} is the finish time of packet $i-1$ of the same flow




Queuing Discipline

- Bit-Level Fair Queuing
 - Schedule according to earliest finish time
 - Finish time of packet i : $F_i = \max(F_{i-1}, A_i) + P_i$
 - P_i is the transmitting duration of packet i , A_i is the arriving time of packet i , F_{i-1} is the finish time of packet $i-1$ of the same flow
 - Assume $A_i=0$, the finish time is virtually calculated in bit-level
 - Packets are dequeued according to the order: $F_i^1 < F_i^2 < F_i^3 < F_{i+1}^2 < F_i^4$




Queuing Discipline

- Bit-Level Fair Queuing (FQ)
 - Schedule according to earliest finish time
 - Finish time of packet i : $F_i = \max(F_{i-1}, A_i) + P_i$
 - P_i is the transmitting duration of packet i , A_i is the arriving time of packet i , F_{i-1} is the finish time of packet $i-1$ of the same flow
 - Case: when one queue is idle, idle time also count



Queuing Discipline

- Bit-Level Fair Queuing (FQ)
 - Schedule according to earliest finish time
 - Finish time of packet i : $F_i = \max(F_{i-1}, A_i) + P_i$
 - P_i is the transmitting duration of packet i , A_i is the arriving time of packet i , F_{i-1} is the finish time of packet $i-1$ of the same flow
 - Case: when one queue is idle, idle time also count




Implementation of Per-Hop Behavior

- Expedited Forwarding (EF) PHB
 - Highest priority
- Assured Forwarding (AF) PHB
 - Different levels of priorities, drop probabilities, bandwidth, etc.


Implementation of Expedited Forwarding

- First-In-First-Out (FIFO) with Priority



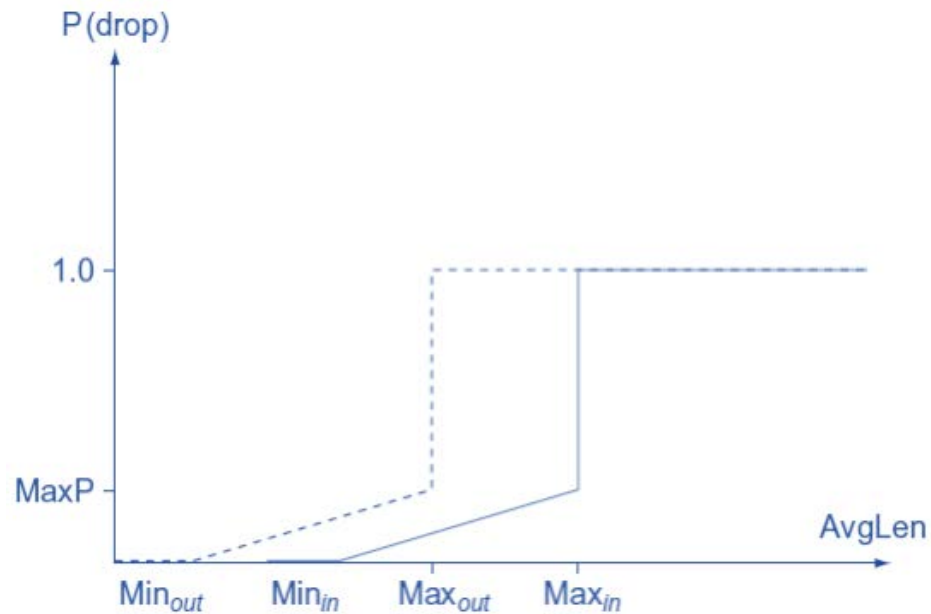
Implementation of Expedited Forwarding

- Weighted Fair Queuing (FQ)



Implementation of Assured Forwarding


- RED with In and Out (RIO)



月基本费		258元	188元
包含	国内流量	1GB	
	国内通话	800分钟	
	本地流量	本地流量无限量权益 (用满40GB后限速)	

Network Neutrality

- Network Neutrality
 - ISPs supply non-discriminated IP connectivity



Network Neutrality

- Opposite Counterpoint
 - ISPs only allows you to access their (often value-added) services



Outline


- TCP Fairness
- QoS
- QUIC

QUIC

- QUIC: Quick UDP Internet Connections
- Application-layer protocol, on top of UDP
 - Deployed by Google starting at 2014
 - Deployed on many Google servers, apps (Chrome, mobile YouTube app)
 - QUIC working group formed in Oct 2016
- Initial goal: increase performance of HTTP

QUIC

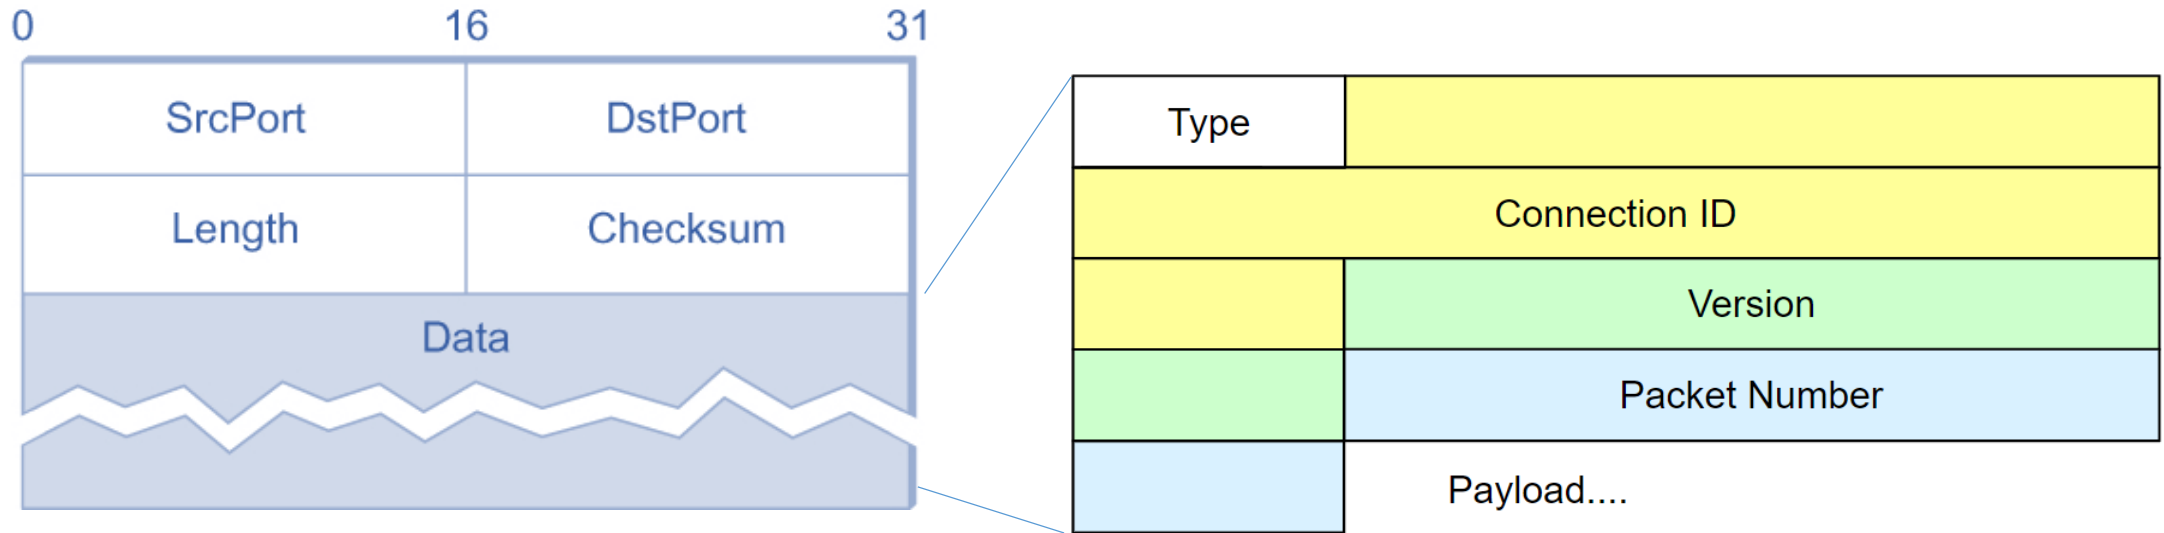
- Protocol Stack




QUIC

- Key features
 - Always encrypted
 - 0-RTT connection establishment
 - Connection migration
 - Congestion control
 - Parallel Streams

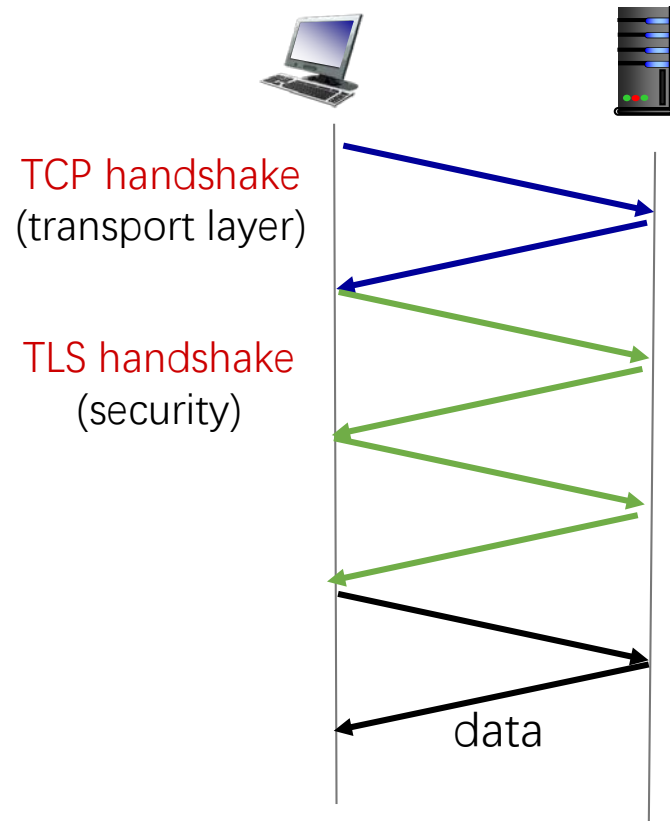
QUIC - Header




QUIC Connection Establishment



TCP
(2RTT)



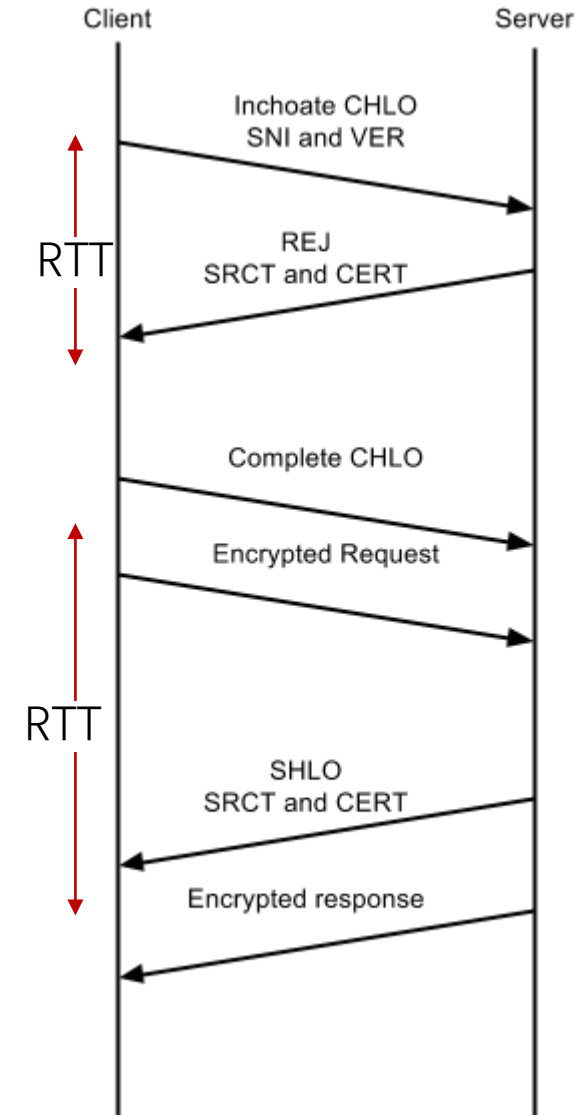
TCP+TLS 1.2
(new 4RTT
resumed 3RTT)



QUIC
(new 2RTT
Resumed 1RTT)

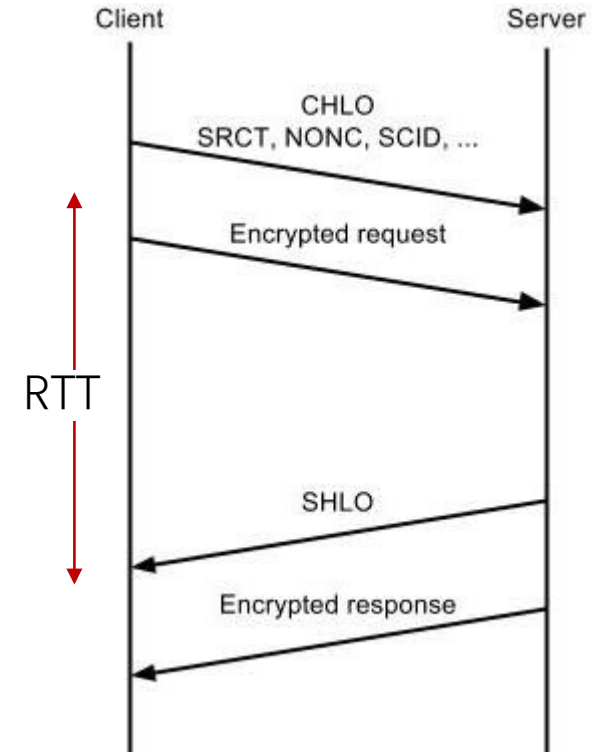
QUIC Connection Establishment

- 1-RTT (First-ever connection)
 - No cached information available
 - First CHLO is inchoate (empty)
 - Simply includes version and server name
 - Server responds with REJ
 - Includes server config, certs, etc.
 - Allows client to make forward progress
 - Second CHLO is complete
 - Followed by initially encrypted request data
 - Server responds with SHLO
 - Followed immediately by forward-secure encrypted response data



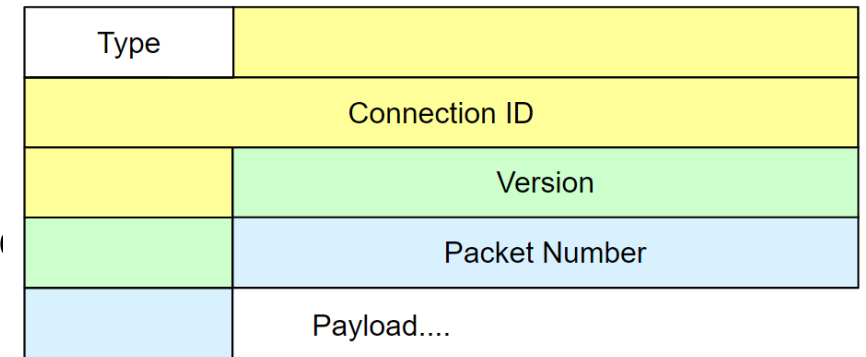
QUIC Connection Establishment

- 0-RTT (Subsequent connection)
 - Motivation: client can cache information about the *origin* it connected to
 - First CHLO is complete
 - Based on information from previous connection
 - Followed by initially encrypted data.
 - Server responds with SHLO
 - Followed immediately by forward-secure encrypted data



QUIC Connection Migration

- NAT Rebinding
 - NATs remaps port
 - Frequency (~ mins)
 - Why ? to release unused ports
 - According to TCP connection state (if they are closed)
 - UDP does not have connection state, QUIC state is encrypted
- Mobility
 - Switching between different IP
 - Wi-Fi and cellular network
- Connection Migration
 - Keep QUIC connections alive even if port and IP are changing
 - Detect connection path changes via Connection ID and IP/port
 - Connection is identified by connection ID rather than <IP, port>
 - 64-bit connection ID
 - randomly chosen by client



QUIC Congestion Control

- Incorporates TCP best practices
 - TCP Cubic, Fast Retransmission, Selective ACK, etc.
- Better signaling than TCP
 - Each packet carries a monotonically increasing packet number
 - Better RTT measurement
 - Even ACK
 - Retransmitted packets also consume new sequence numbers
 - no retransmission ambiguity
- More verbose ACK
 - support 256 NAK ranges (vs. TCP's 3SACK ranges)

QUIC - Parallel Streams

- Handle HOL blocking

Reference

- Textbook 6.5
- Some slides are adapted from http://www-net.cs.umass.edu/kurose_ross/ppt.htm by Kurose Ross
- <https://www.chromium.org/quic>