


Announcement




- TA office hour
 - Wed evening: 20:00-21:00
 - Location: SIST 1B-101
- Programming Assignment 1:
 - Released on Thursday
- Autolab registration
 - See email or BB announcement

Informed Search

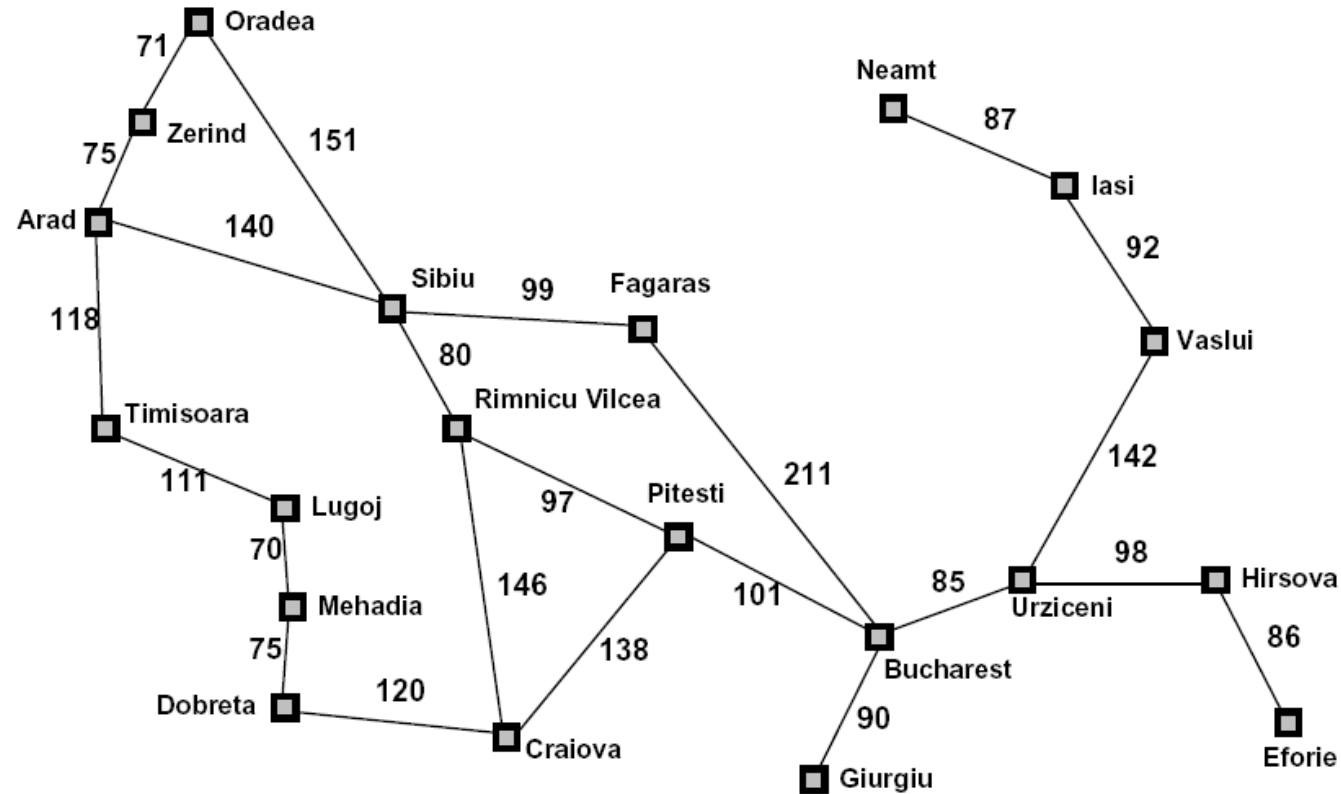


Search Heuristics

- A heuristic h is:
 - A function that *estimates* how close a state is to a goal
 - $h(\text{goal})=0$
 - Designed for a particular search problem
 - Examples: Manhattan distance, Euclidean distance for pathing



Example: Heuristic Function




Straight-line distance
to Bucharest

| | |
|----------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |

$h(x)$

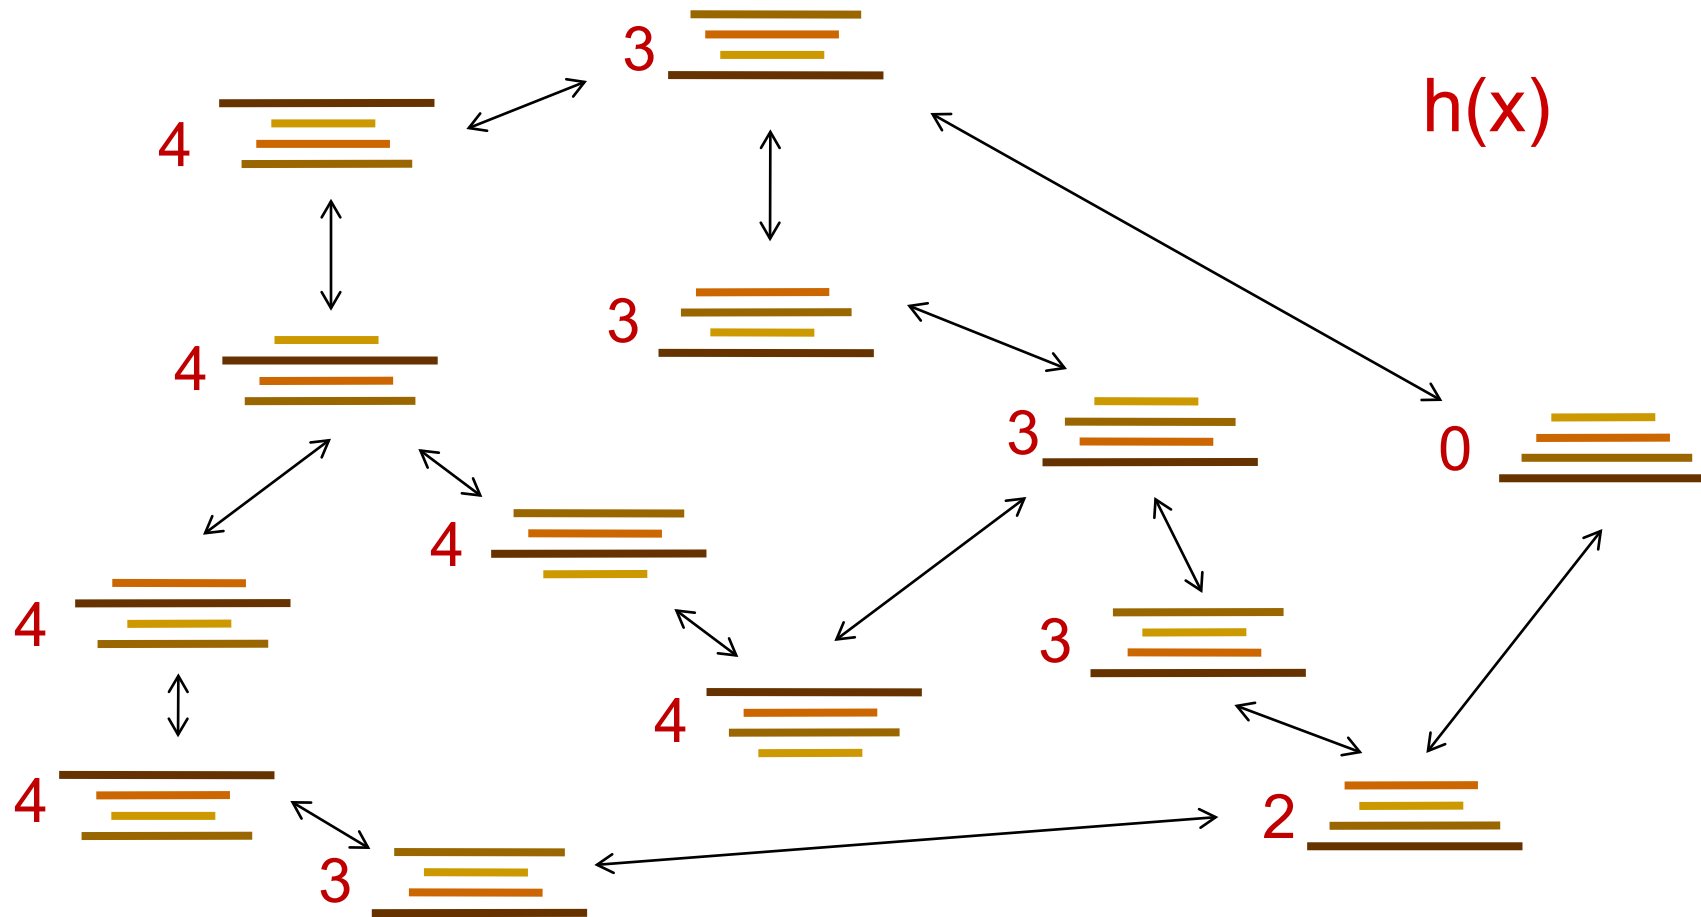
Example: Pancake Problem



Cost: Number of pancakes flipped

Example: Heuristic Function


Heuristic: the number of the largest pancake that is still out of place



Greedy Search



Example: Heuristic Function

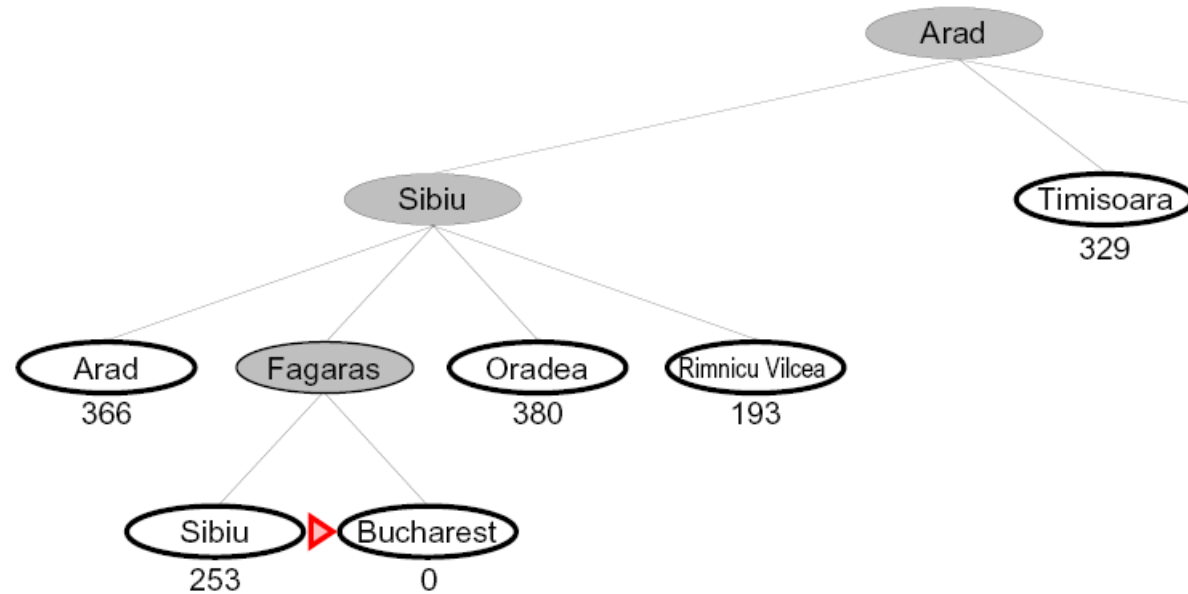


| Straight-line distance to Bucharest | |
|-------------------------------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |



$h(x)$

Greedy Search

- Expand the node that seems closest...





- What can go wrong?



Greedy Search

- Strategy: expand a node that you think is closest to a goal state
- The ideal scenario:
 - Best-first takes you straight to the goal
- Worst-case: like a badly-guided DFS



Video of Demo Contours Greedy (Pacman Small Maze)



A* Search



A* Search



UCS



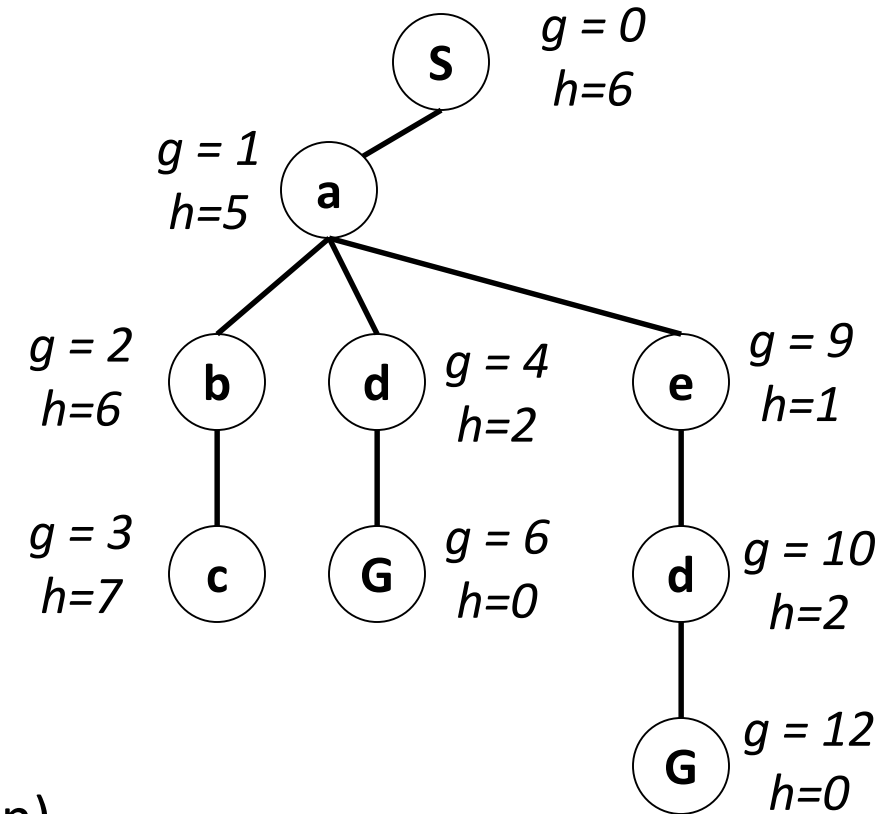

Greedy



A*

Combining UCS and Greedy


- **Uniform-cost** orders by path cost, or *backward cost* $g(n)$
- **Greedy** orders by goal proximity, or *forward cost* $h(n)$



- **A* Search** orders by the sum: $f(n) = g(n) + h(n)$


When should A* terminate?

- Should we stop when we enqueue a goal?




- No: only stop when we dequeue a goal

Is A* Optimal?




- What went wrong?
- Over-estimated goal cost

Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs


Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$


where $h^*(n)$ is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A^* in practice.

Optimality of A* Tree Search




Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:


- A will exit the fringe before B



Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f-cost


Admissibility of h

$h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$


B is suboptimal

$h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:



- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal




$$f(n) \leq f(A) < f(B)$$

UCS vs A* Contours

- Uniform-cost expands equally in all “directions”
- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Video of Demo Maze with Deep/Shallow Water --- UCS, Greedy, A*



Video of Demo Maze with Deep/Shallow Water --- UCS, Greedy, A*



Video of Demo Maze with Deep/Shallow Water --- UCS, Greedy, A*

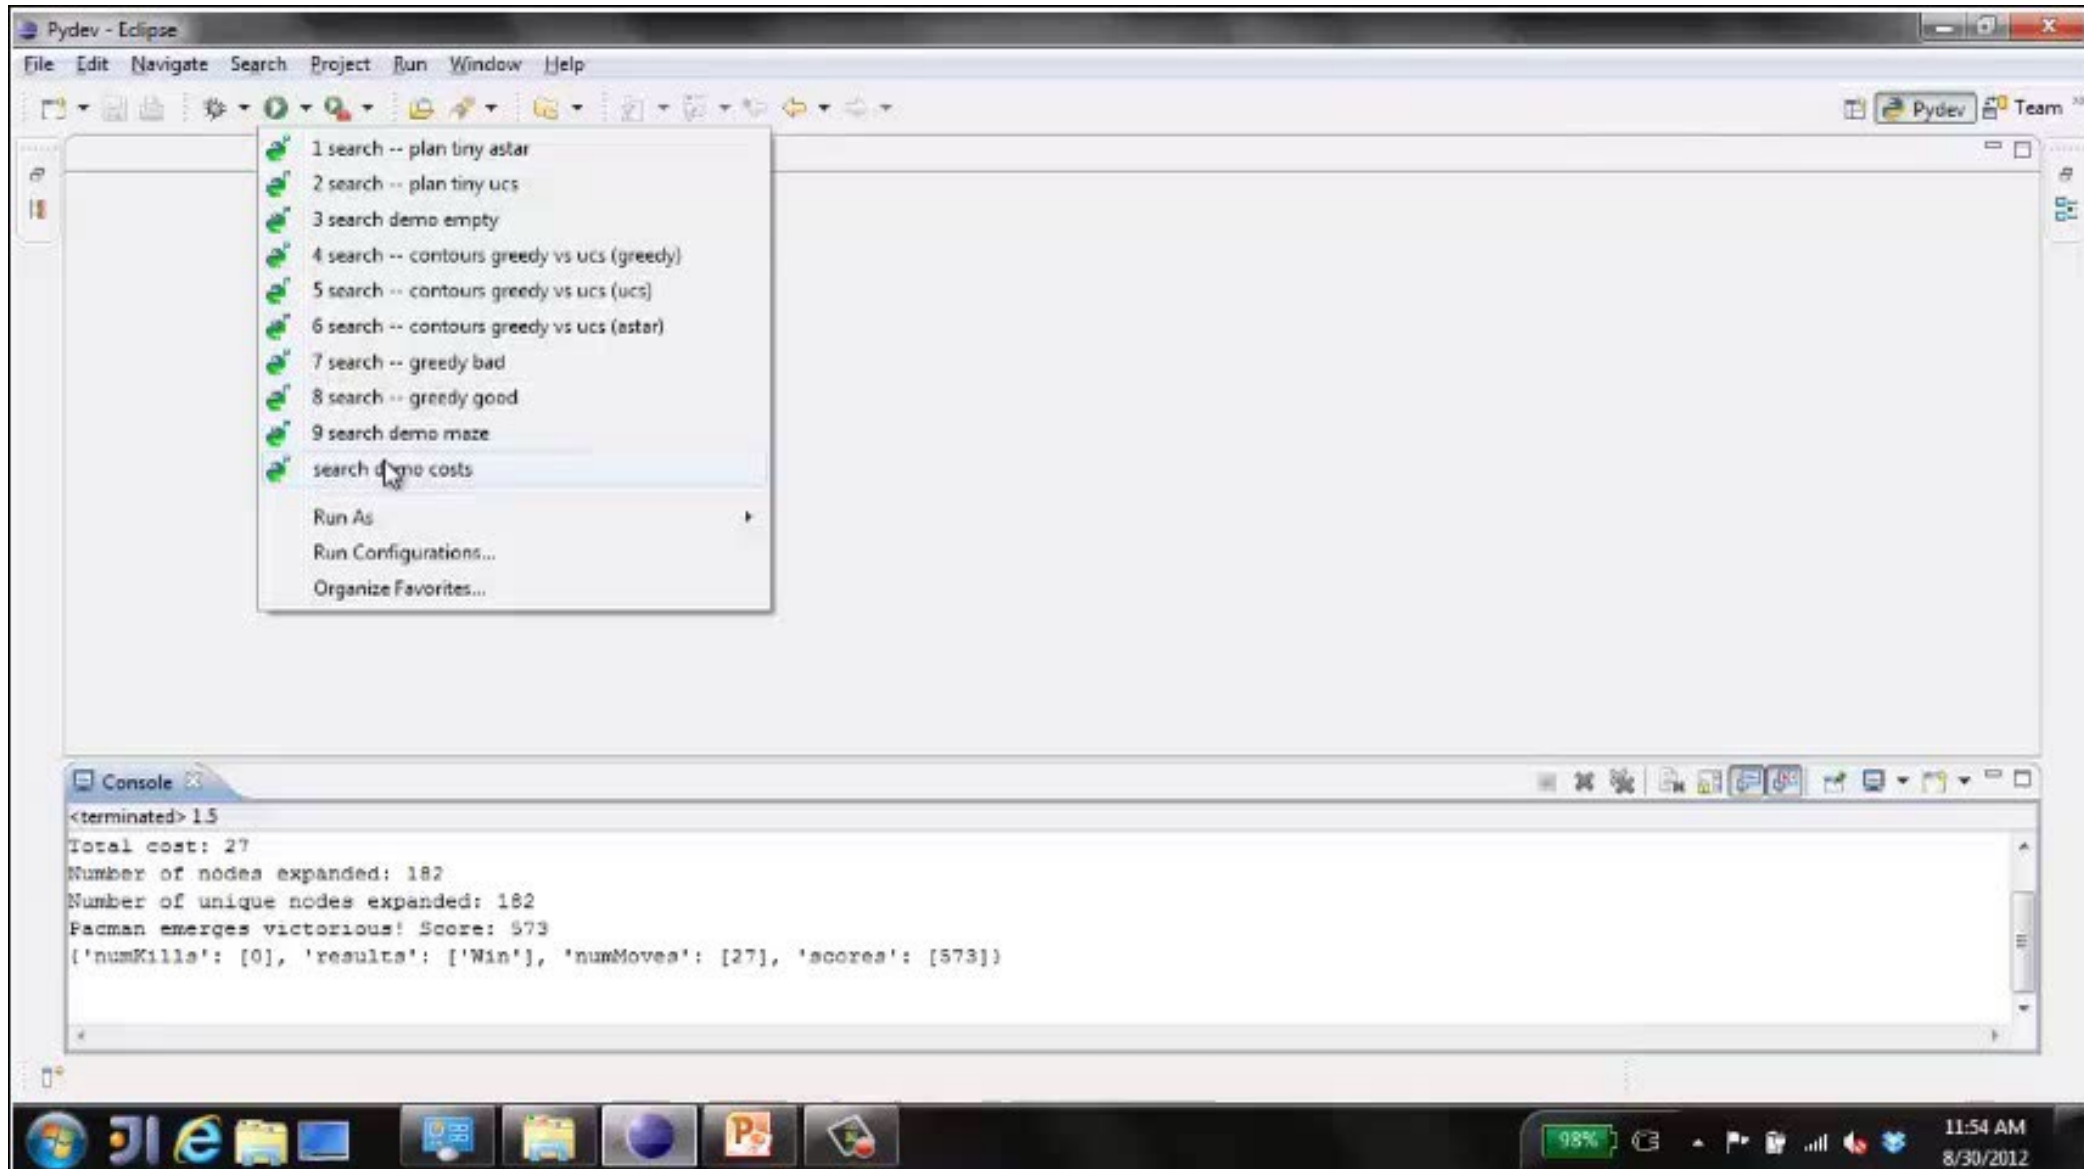


A* Applications

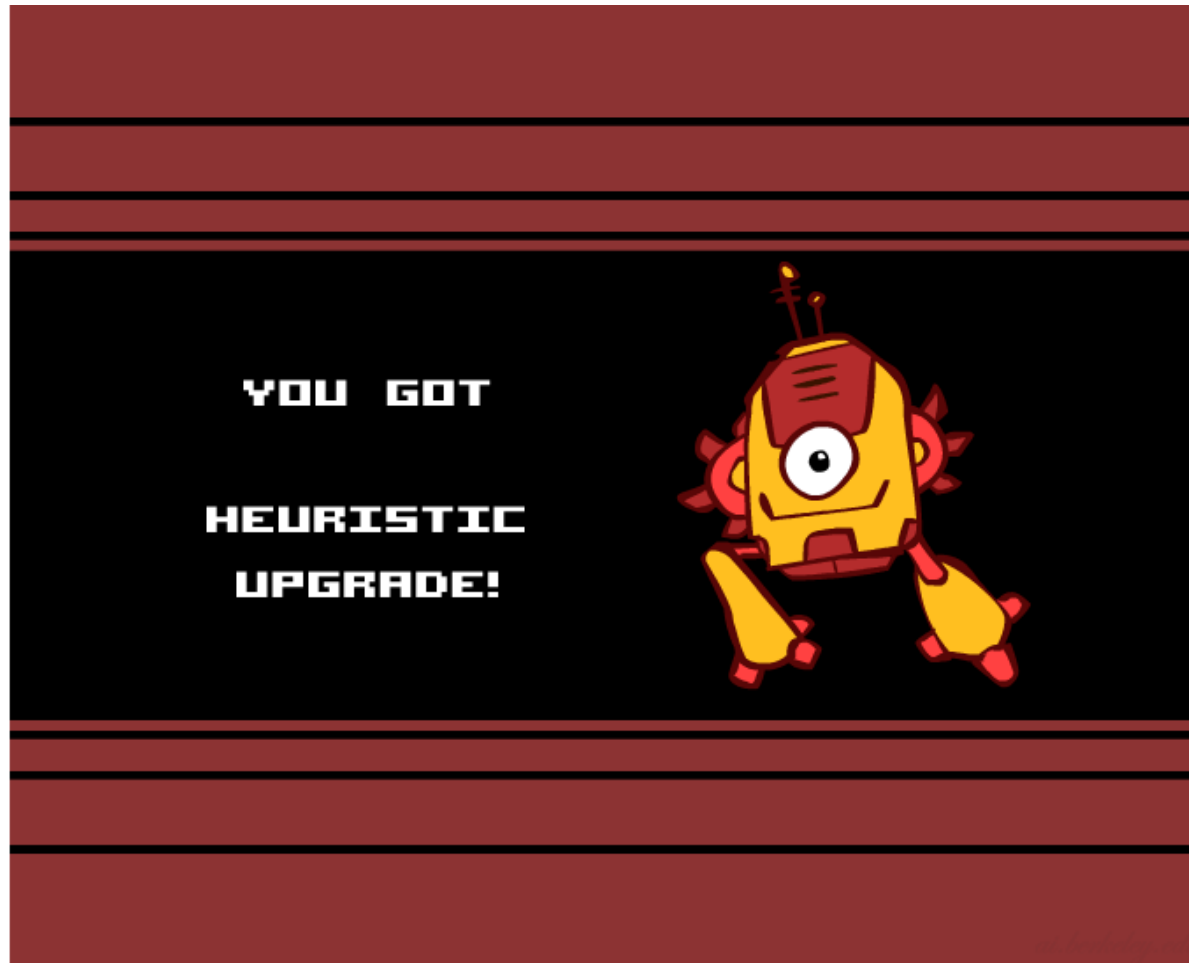
- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...



Video of Demo Empty Water Shallow/Deep – Guess Algorithm


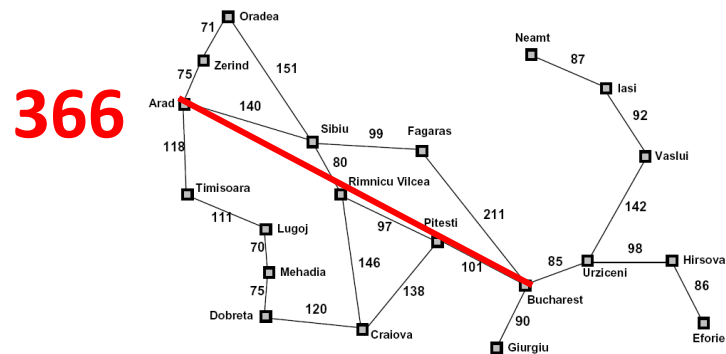


Creating Heuristics



Creating Admissible Heuristics


- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
 - Inadmissible heuristics are often useful too
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available



Example: 8 Puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State



Actions

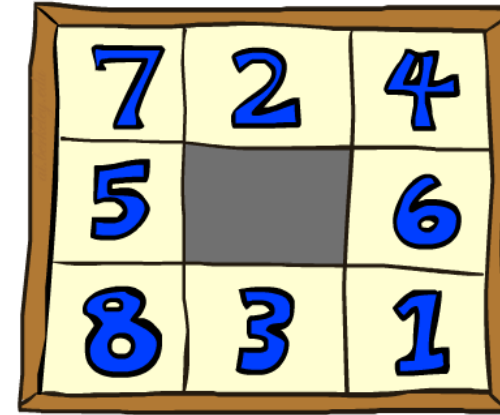
| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

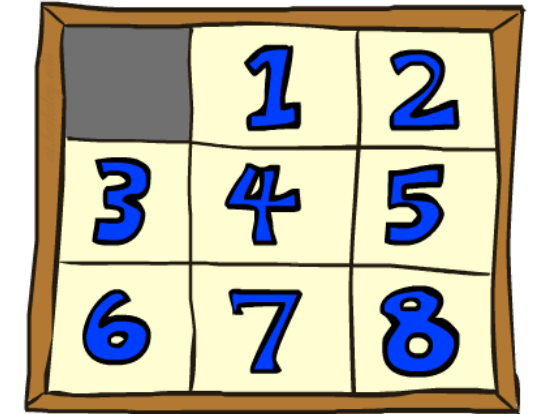
- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

8 Puzzle I


- Heuristic: Number of tiles misplaced
- $h(\text{start}) = 8$
- Is it admissible?
- This is a *relaxed-problem* heuristic



Start State



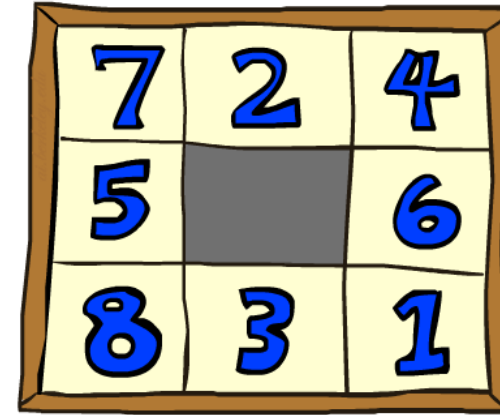
Goal State



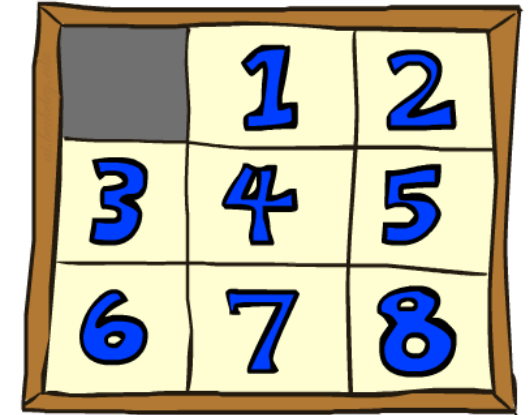
| Average nodes expanded when the optimal path has... | | | |
|---|------------|------------|-------------------|
| | ...4 steps | ...8 steps | ...12 steps |
| UCS | 112 | 6,300 | 3.6×10^6 |
| TILES | 13 | 39 | 227 |

8 Puzzle II

- Heuristic: total *Manhattan* distance
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$
- Is it admissible?
- *Relaxed-problem*: any tile could slide in any direction at any time, ignoring other tiles



Start State




Goal State

| Average nodes expanded when the optimal path has... | | | |
|---|------------|------------|-------------|
| | ...4 steps | ...8 steps | ...12 steps |
| TILES | 13 | 39 | 227 |
| MANHATTAN | 12 | 25 | 73 |

8 Puzzle III

- How about using the *actual cost* as a heuristic?


- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



- With A^* : a trade-off between quality of estimate and work per node



- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

Graph Search



Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.




Graph Search


- Idea: never **expand** a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set
- Can graph search wreck completeness? Why/why not?
- How about optimality?

A* Graph Search Gone Wrong?


State space graph



Search tree




Consistency of Heuristics



- Main idea: estimated heuristic costs \leq actual costs
 - Admissibility: heuristic cost \leq actual cost to goal
$$h(A) \leq \text{actual cost from A to G}$$
 - Consistency: heuristic “arc” cost \leq actual cost for each arc
$$h(A) - h(C) \leq \text{cost(A to C)}$$
 - Consistency implies admissibility
- Consequences of consistency:
 - The f value along a path never decreases, because:
$$h(A) \leq \text{cost(A to C)} + h(C)$$


Consistency of Heuristics

- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems
- Example: Manhattan distance or Euclidean distance for pathing




Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - See textbook for a proof
 - UCS optimal ($h = 0$ is consistent)



A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



Summary

- Why search
 - Agents that Plan Ahead
- Search Problems
 - state space, successor function
 - start state and goal test
- Uninformed Search Methods
 - DFS, BFS, UCS
- Informed Search
 - Greedy, A*
- Graph Search

