

Announcement


- Homework 1: search, CSP, adversarial search
 - Available in Blackboard -> Homework
 - Due: March 8, 11:59pm
- Programming Assignment 1A
 - Submission at AutoLab
 - Due: Mar 5, 11:59pm

Constraint Satisfaction Problems




AIMA Chapter 6

Structure




Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$, linear in n
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



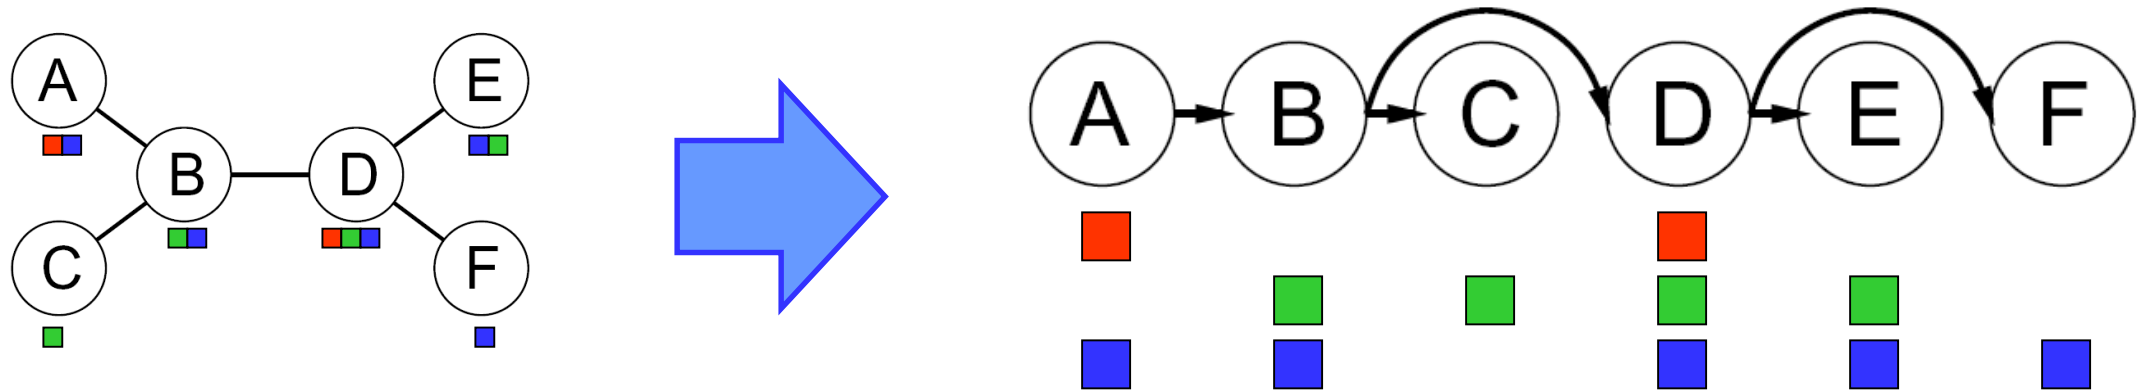
Tree-Structured CSPs



- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(n d^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to probabilistic reasoning (later)
- An example of the relation between syntactic restrictions and the complexity of reasoning

Tree-Structured CSPs


- Algorithm for tree-structured CSPs:
 - Order: Choose a root variable, order variables so that parents precede children



- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$
 - Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$
- Runtime: $O(n d^2)$


Tree-Structured CSPs

- Claim 1: After backward pass, all root-to-leaf arcs are consistent
- Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)

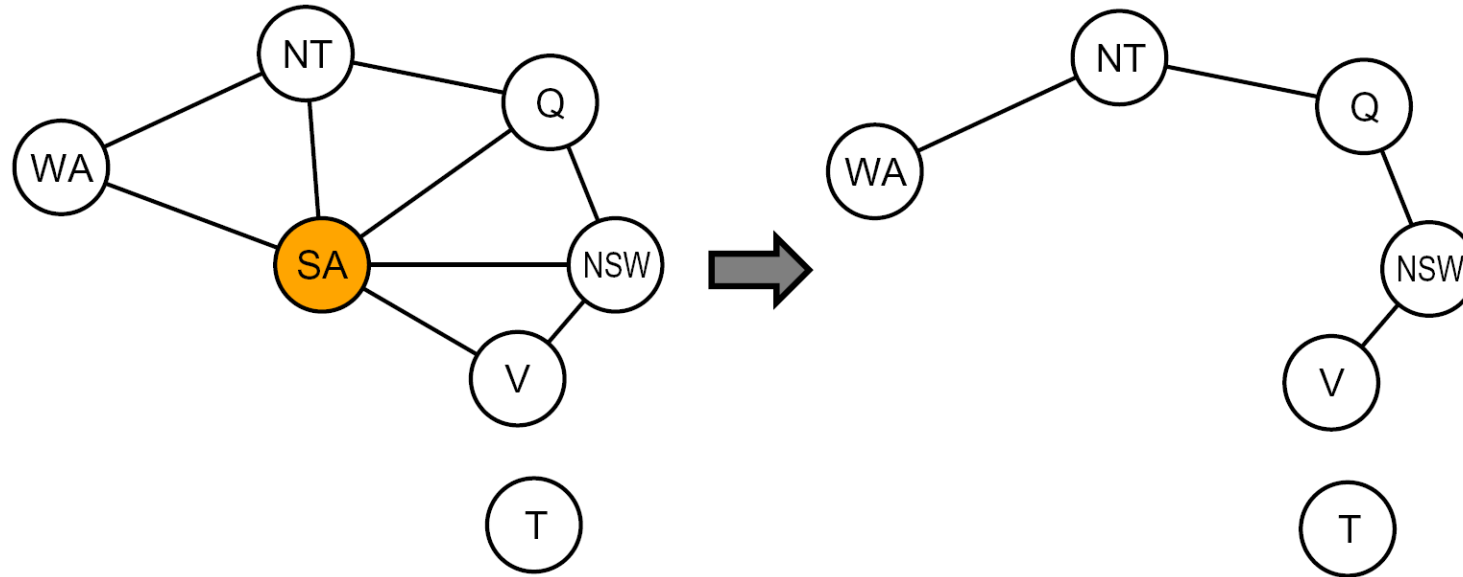


- Claim 2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
- Easy to prove
- Why doesn't this algorithm work with cycles in the constraint graph?
- Note: we'll see this basic idea again with Bayes' nets

Cutset Conditioning



Nearly Tree-Structured CSPs



- Cutset: a set of variables s.t. the remaining constraint graph is a tree
- Cutset conditioning: instantiate (in all ways) the cutset and solve the remaining tree-structured CSP
 - Cutset size c gives runtime $O(d^c (n-c) d^2)$, very fast for small c


Cutset Conditioning

Choose a cutset

Instantiate the cutset
(all possible ways)


Compute residual CSP
for each assignment

Solve the residual CSPs
(tree structured)




Finding Cutset

- Find the smallest cutset for the graph below.



- Finding the *smallest* cutset is NP-hard
- But there are efficient approximation algorithms


Iterative Improvement



Iterative Algorithms for CSPs

- Idea:

- Take a complete assignment with unsatisfied constraints
- *Reassign* variable values to minimize conflicts




- Algorithm: While not solved,

- Variable selection: randomly select any conflicted variable
- Value selection: min-conflicts heuristic:
 - Choose a value that violates the fewest constraints



Example: 4-Queens





- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n)$ = number of attacks

Demo – Iterative Improvement – Coloring

Performance

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$




Local Search



Local Search

- Goal: identification, optimization
- Local search: improve a single option until you can't make it better
- State: a complete assignment
- Successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)

Hill Climbing


- Simple, general idea:
 - Start wherever
 - Repeat: move to the best neighboring state
 - If no neighbors better than current, quit
- What's good about this approach?
 - Simple, fast
- What's bad about it?



Hill Climbing Diagram



Hill Climbing Quiz




Starting from X, where do you end up ?

Starting from Y, where do you end up ?


Starting from Z, where do you end up ?

Beam Search

- Like greedy hill climbing search, but keep K states at all times:



Greedy Search




Beam Search


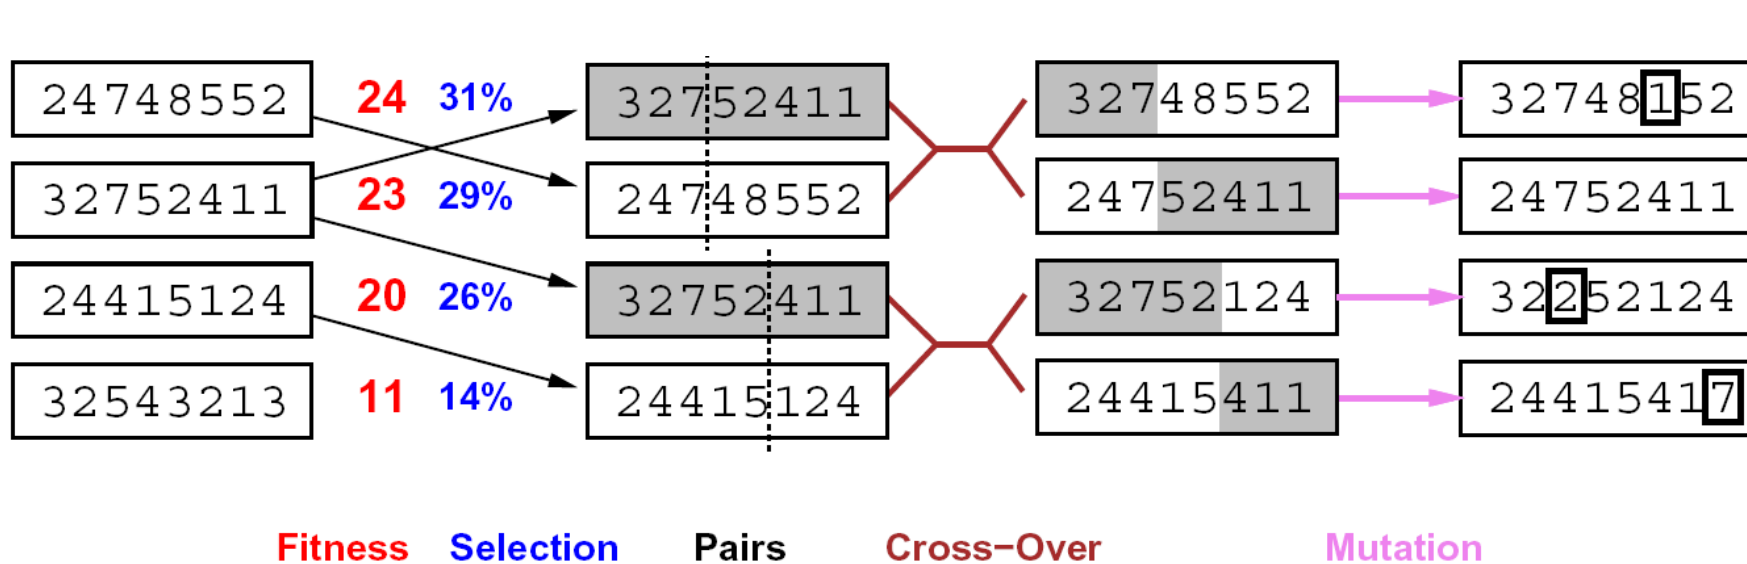
- The best choice in MANY practical settings
- Optimal?

Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves
 - Pick a random move
 - Always accept an uphill move
 - Accept a downhill move with probability $e^{-\Delta E / T}$
 - But make the probability smaller (by decreasing T) as time goes on
- Theoretical guarantee
 - If T decreased slowly enough, will converge to optimal state!
- Sounds like magic, but reality is reality:
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all



Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
 - Keep the best (or sample) N states at each step based on a fitness function
 - Pairwise crossover operators, with optional mutation to give variety

Summary: CSPs

- CSPs are a special kind of search problem:
 - States are partial assignments
 - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
 - Filtering
 - Forward Checking, Arc Consistency
 - Ordering
 - MRV, LCV
 - Structure
 - Tree structured, Cutset conditioning
- Iterative min-conflicts (local search) is often effective in practice

