

Announcement

- Programming Assignment 5
 - Due: May. 31, 11:59pm
- Homework 5
 - Due: May. 24, 11:59pm

Supervised Machine Learning




AIMA Chapter 18, 20

Machine Learning


- Up until now: how to use a model to make optimal decisions
 - Except reinforcement learning
- Machine learning: how to acquire a model from data / experience
- Related courses
 - SI151 Optimization and Machine Learning
 - CS282 Machine Learning
 - CS280 Deep Learning

Types of Learning


- Supervised learning 
 - Training data includes desired outputs
- Unsupervised learning
 - Training data does not include desired outputs
- Semi-supervised learning
 - Training data includes a few desired outputs
- Reinforcement learning
 - Rewards from sequence of actions

Supervised learning

- To learn an unknown *target function* f
- Input: a *training set* of *labeled examples* (x_j, y_j) where $y_j = f(x_j)$
- Output: *hypothesis* h that is “close” to f
- Types of supervised learning
 - Classification = learning f with discrete output value
 - Regression = learning f with real-valued output value
 - Structured prediction = learning f with structured output



Classification



Example: Spam Filter

- Input: an email
- Output: spam/ham
- Setup:
 - Get a large collection of example emails, each labeled “spam” or “ham” (by hand)
 - Want to learn to predict labels of new, future emails
- Features: The attributes used to make the ham / spam decision
 - Words: FREE!
 - Text Patterns: \$dd, CAPS
 - Non-text: SenderInContacts
 - ...



Dear Sir.

First, I must solicit your confidence in this transaction, this is by virtue of its nature as being utterly confidential and top secret. ...



TO BE REMOVED FROM FUTURE MAILINGS, SIMPLY REPLY TO THIS MESSAGE AND PUT "REMOVE" IN THE SUBJECT.

99 MILLION EMAIL ADDRESSES
FOR ONLY \$99



Ok, I know this is blatantly OT but I'm beginning to go insane. Had an old Dell Dimension XPS sitting in the corner and decided to put it to use, I know it was working pre being stuck in the corner, but when I plugged it in, hit the power nothing happened.

Example: Digit Recognition

- Input: images / pixel grids
- Output: a digit 0-9
- Setup:
 - Get a large collection of example images, each labeled with a digit
 - Want to learn to predict labels of new, future digit images
- Features: The attributes used to make the digit decision
 - Pixels: (6,8)=ON
 - Shape Patterns: NumComponents, AspectRatio, NumLoops
 - ...



0



1



2




1



??

Other Classification Tasks

- Medical diagnosis
 - input: symptoms
 - output: disease
- Automatic essay grading
 - input: document
 - output: grades
- Fraud detection
 - input: account activity
 - output: fraud / no fraud
- Email routing
 - input: customer complaint email
 - output: which department needs to ignore this email
- Fruit and vegetable inspection
 - input: image (or gas analysis)
 - output: moldy or OK
- ... many more



Naïve Bayes Classifier



Model-Based Classification

- Model-based approach

- Build a model (e.g. Bayes' net) where both the label and features are random variables
- Instantiate any observed features
- Query for the distribution of the label conditioned on the features

- Challenges


- What structure should the BN have?
- How should we learn its parameters?

Naïve Bayes for Digits

- Naïve Bayes: Assume all features are independent effects of the label

- Simple digit recognition version:


- One feature (variable) F_{ij} for each grid position $\langle i,j \rangle$
- Feature values are on / off, based on whether intensity is more or less than 0.5 in underlying image
- Each input maps to a feature vector, e.g.

 $\rightarrow \langle F_{0,0} = 0 \ F_{0,1} = 0 \ F_{0,2} = 1 \ F_{0,3} = 1 \ F_{0,4} = 0 \ \dots F_{15,15} = 0 \rangle$

- Here: lots of features, each is binary valued

- Naïve Bayes model: $P(Y|F_{0,0} \dots F_{15,15}) \propto P(Y) \prod_{i,j} P(F_{i,j}|Y)$

- What do we need to learn?



General Naïve Bayes


- A general Naive Bayes model:

$|Y|$ parameters

$$P(Y, F_1 \dots F_n) = P(Y) \prod_i P(F_i|Y)$$

$|Y| \times |F|^n$ values

$n \times |F| \times |Y|$
parameters



- We only have to specify how each feature depends on the class
- Total number of parameters is *linear* in n
- Model is very simplistic, but often works anyway

Inference for Naïve Bayes

- Goal: compute posterior distribution over label variable Y
 - Step 1: get joint probability of label and evidence for each label

$$P(Y, f_1 \dots f_n) = \begin{bmatrix} P(y_1, f_1 \dots f_n) \\ P(y_2, f_1 \dots f_n) \\ \vdots \\ P(y_k, f_1 \dots f_n) \end{bmatrix} \Rightarrow \begin{bmatrix} P(y_1) \prod_i P(f_i|y_1) \\ P(y_2) \prod_i P(f_i|y_2) \\ \vdots \\ P(y_k) \prod_i P(f_i|y_k) \end{bmatrix}$$

$$P(f_1 \dots f_n)$$

+ ↶


- Step 2: normalization

$$P(Y|f_1 \dots f_n)$$

General Naïve Bayes

- What do we need in order to use Naïve Bayes?
 - Inference method (we just saw this part)
 - Start with a bunch of probabilities: $P(Y)$ and the $P(F_i|Y)$ tables
 - Use standard inference to compute $P(Y|F_1...F_n)$
 - Nothing new here
 - Estimates of local conditional probability tables
 - $P(Y)$, the prior over labels
 - $P(F_i|Y)$ for each feature (evidence variable)
 - These probabilities are collectively called the *parameters* of the model and denoted by θ
 - Up until now, we assumed these appeared by magic, but...
 - ...they typically come from training data counts: we'll look at this soon

Example: Conditional Probabilities



Naïve Bayes for Text

- Bag-of-words Naïve Bayes:

- Features: W_i is the word at position i

$$P(Y, W_1 \dots W_n) = P(Y) \prod_i P(W_i | Y)$$

*Word at position
 i , not i^{th} word in
the dictionary!*

- Usually, each variable gets its own conditional probability distribution $P(W_i | Y)$
- Here
 - Each position is **identically distributed**
 - All positions share the same conditional probabilities $P(W | Y)$
- Called “**bag-of-words**” because model is insensitive to word order or reordering

Example: Spam Filtering

- **Model:** $P(Y, W_1 \dots W_n) = P(Y) \prod_i P(W_i|Y)$

$P(Y)$

ham : 0.66
spam: 0.33

$P(W|\text{spam})$

the : 0.0156
to : 0.0153
and : 0.0115
of : 0.0095
you : 0.0093
a : 0.0086
with: 0.0080
from: 0.0075
...

$P(W|\text{ham})$

the : 0.0210
to : 0.0133
of : 0.0119
2002: 0.0110
with: 0.0108
from: 0.0107
and : 0.0105
a : 0.0100
...

Spam Example


$$\begin{array}{c} P(Y) \\ P(W_1|Y) \\ P(W_2|Y) \end{array}$$

Word	P(w spam)	P(w ham)	Tot Spam	Tot Ham
(prior)	0.33333	0.66666	-1.1	-0.4
Gary	0.00002	0.00021	-11.8	-8.9
would	0.00069	0.00084	-19.1	-16.0
you	0.00881	0.00304	-23.8	-21.8
like	0.00086	0.00083	-30.9	-28.9
to	0.01517	0.01339	-35.1	-33.2
lose	0.00008	0.00002	-44.5	-44.0
weight	0.00016	0.00002	-53.3	-55.0
while	0.00027	0.00027	-61.5	-63.2
you	0.00881	0.00304	-66.2	-69.0
sleep	0.00006	0.00001	-76.0	-80.5

Log product of probabilities



$$P(\text{spam} \mid w) = 98.9$$

Training and Testing




Important Concepts

- Data: labeled instances, e.g. emails marked spam/ham
 - Training set
 - Held out set
 - Test set
- Experimentation cycle
 - Learn parameters (e.g. model probabilities) on training set
 - Tune hyperparameters on held-out set
 - Compute accuracy of test set (fraction of instances predicted correctly)
 - Very important: never “peek” at the test set!



Training




Parameter Estimation

- Estimating the distribution of a random variable
- *Elicitation*: ask a human (this is hard...)
- *Empirically*: use training data (learning!)
 - For each outcome x , look at the *empirical rate* of that value

$$P_{\text{ML}}(x) = \frac{\text{count}(x)}{\text{total samples}}$$

- Ex:
 - We've seen 1000 words from spam emails, among which we see "money" for 50 times
 - So we set $P(\text{money} \mid \text{spam}) = 0.05$
- This is the estimate that maximizes the *likelihood of the data*
 - Likelihood: conditional probability of the data given the parameters



Maximum Likelihood Estimation

- Coin flipping:
 - $P(\text{Heads}) = \theta$, $P(\text{Tails}) = 1 - \theta$
- Flips are *i.i.d.*
 - Independent events
 - Identically distributed according to unknown distribution
- Sequence \mathcal{D} of α_H Heads and α_T Tails

$$P(\mathcal{D} \mid \theta) = \theta^{\alpha_H} (1 - \theta)^{\alpha_T}$$

Maximum Likelihood Estimation

- **MLE:** Choose θ to maximize probability of \mathcal{D}

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} \ln P(\mathcal{D} \mid \theta) \\ &= \arg \max_{\theta} \ln \theta^{\alpha_H} (1 - \theta)^{\alpha_T}\end{aligned}$$

- Set derivative to zero, and solve!

$$\frac{d}{d\theta} \ln P(\mathcal{D} \mid \theta) = \frac{d}{d\theta} [\ln \theta^{\alpha_H} (1 - \theta)^{\alpha_T}]$$



$$= \frac{d}{d\theta} [\alpha_H \ln \theta + \alpha_T \ln(1 - \theta)]$$

$$= \alpha_H \frac{d}{d\theta} \ln \theta + \alpha_T \frac{d}{d\theta} \ln(1 - \theta)$$



$$= \frac{\alpha_H}{\theta} - \frac{\alpha_T}{1 - \theta} = 0$$

$$\boxed{\hat{\theta}_{MLE} = \frac{\alpha_H}{\alpha_H + \alpha_T}}$$

Generalization and Overfitting



Overfitting



Example: Overfitting

$P(\text{features}, C = 2)$


$P(C = 2) = 0.1$

$P(\text{on}|C = 2) = 0.8$

$P(\text{on}|C = 2) = 0.1$

$P(\text{off}|C = 2) = 0.1$

$P(\text{on}|C = 2) = 0.01$



$P(\text{features}, C = 3)$

$P(C = 3) = 0.1$


$P(\text{on}|C = 3) = 0.8$

$P(\text{on}|C = 3) = 0.9$

$P(\text{off}|C = 3) = 0.7$

$P(\text{on}|C = 3) = 0.0$

2 wins!!



Example: Overfitting

- Posterior determined by *relative* probabilities (odds ratios):


$$\frac{P(W|\text{ham})}{P(W|\text{spam})}$$

south-west	:	inf
nation	:	inf
morally	:	inf
nicely	:	inf
extent	:	inf
seriously	:	inf
...		

$$\frac{P(W|\text{spam})}{P(W|\text{ham})}$$

screens	:	inf
minute	:	inf
guaranteed	:	inf
\$205.00	:	inf
delivery	:	inf
signature	:	inf
...		

What went wrong here?



Generalization and Overfitting

- Using empirical rate will **overfit** the training data!
 - Just because we never saw a 3 with pixel (15,15) on during training doesn't mean we won't see it at test time
 - Just because we never saw a word in spam emails during training doesn't mean we won't see it at test time
 - Therefore, we can't give unseen events zero probability
 - More generally, rates in the training data may not exactly match rates at test time

Generalization and Overfitting

- Overfitting: learn to fit the training data very closely, but fit the test data poorly
 - Generalization: try to fit the test data as well
- Why does overfitting occur?
 - Training data is not representative of the true data distribution
 - Too few training samples
 - Training data is noisy
 - Too many attributes, some of them irrelevant to the classification task
 - The model is too expressive
 - Ex: the model is capable of memorizing all the spam emails in the training set

Generalization and Overfitting

- Avoid overfitting
 - Acquire more training data (not always possible)
 - Remove irrelevant attributes (not always possible)
 - Limit the model expressiveness by regularization, early stopping, pruning, etc.
- In our previous example, we may smooth the empirical rate to improve generalization


Laplace Smoothing

- Laplace's estimate:

- Pretend you saw every outcome once more than you actually did

$$\begin{aligned} P_{LAP}(x) &= \frac{c(x) + 1}{\sum_x [c(x) + 1]} \\ &= \frac{c(x) + 1}{N + |X|} \end{aligned}$$

- Can derive this estimate with *Dirichlet priors*



$$P_{ML}(X) =$$

$$P_{LAP}(X) =$$

Laplace Smoothing

- Laplace's estimate (extended):

- Pretend you saw every outcome k extra times


$$P_{LAP,k}(x) = \frac{c(x) + k}{N + k|X|}$$

- k is the **strength** of the prior
- What's Laplace with $k = 0$?

- Laplace for conditionals:

- Smooth each condition independently:

$$P_{LAP,k}(x|y) = \frac{c(x, y) + k}{c(y) + k|X|}$$



$$P_{LAP,0}(X) =$$

$$P_{LAP,1}(X) =$$


$$P_{LAP,100}(X) =$$

Linear Interpolation

- In practice, Laplace often performs poorly for $P(X|Y)$:
 - When $|X|$ is very large
 - When $|Y|$ is very large
- Another option: linear interpolation
 - Also get the empirical $P(X)$ from the data
 - Make sure the estimate of $P(X|Y)$ isn't too different from the empirical $P(X)$

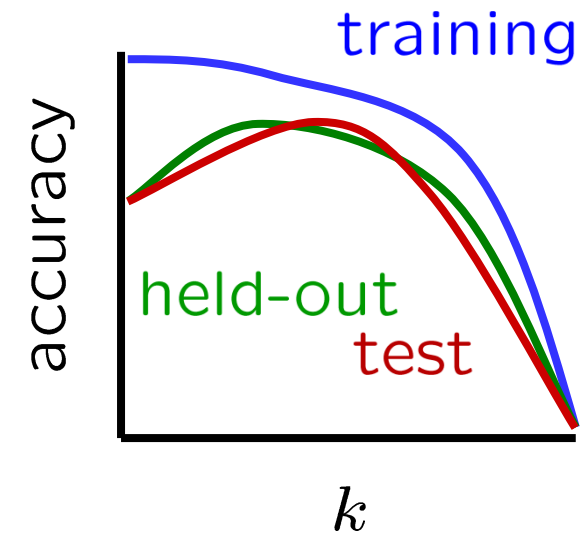
$$P_{LIN}(x|y) = \alpha \hat{P}(x|y) + (1.0 - \alpha) \hat{P}(x)$$

Tuning




Tuning on Held-Out Data

- Now we've got two kinds of unknowns
 - Parameters: the probabilities $P(X|Y)$, $P(Y)$
 - Hyperparameters: e.g. the amount / type of smoothing to do, k , α
- What should we learn where?
 - Learn parameters from training data
 - Tune hyperparameters on different data
 - Why?
 - For each value of the hyperparameter, train on the training data and test on the held-out data
 - Choose the best hyperparameter value and do a final test on the test data




Linear Classifiers




Feature Vectors

 x $f(x)$ y


```
Hello,  
  
Do you want free printr  
cartridges? Why pay more  
when you can get them  
ABSOLUTELY FREE! Just
```




```
# free      : 2  
YOUR_NAME   : 0  
MISPELLED   : 2  
FROM_FRIEND : 0  
...
```



SPAM
or
+





```
PIXEL-7,12 : 1  
PIXEL-7,13 : 0  
...  
NUM_LOOPS   : 1  
...
```



“2”

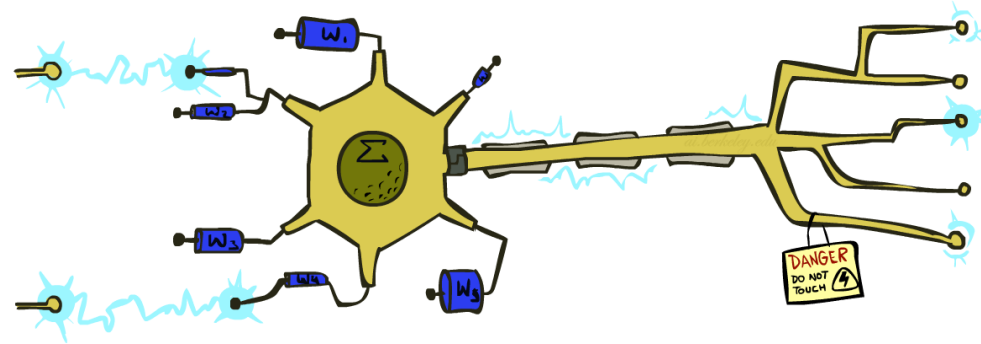
Some (Simplified) Biology

- Very loose inspiration: human neurons




Linear Classifiers

- Inputs are **feature values**
- Each feature has a **weight**
- Sum is the **activation**




$$\text{activation}_w(x) = \sum_i w_i \cdot f_i(x) = w \cdot f(x)$$

- Binary case: if the activation is:
 - Positive, output +1
 - Negative, output -1



Weights

- Binary case: compare features to a weight vector
- Learning: figure out the weight vector from examples





Binary Decision Rule

- In the space of feature vectors
 - Examples are points
 - Any weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$

w

BIAS	:	-3
free	:	4
money	:	2
...		


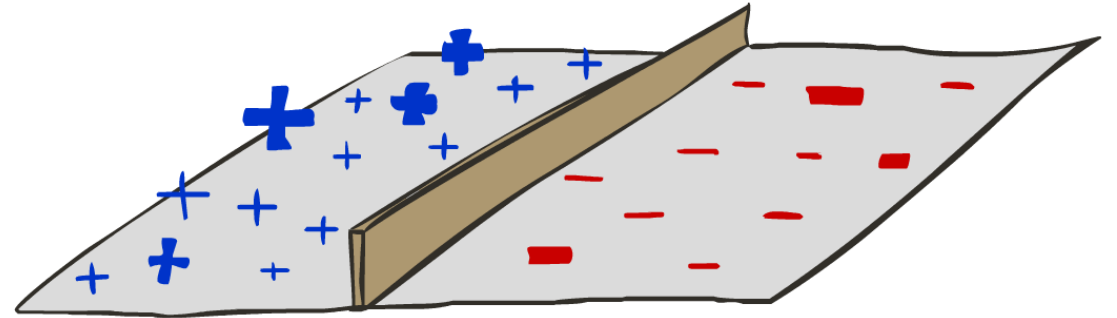


Binary Decision Rule

- In the space of feature vectors
 - Examples are points
 - Any weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$


w

BIAS	:	-3
free	:	4
money	:	2
...		




Binary Decision Rule

- In the space of feature vectors
 - Examples are points
 - Any weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$






w

BIAS	:	-3
free	:	4
money	:	2
...		



Learning: Binary Perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights
- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector




Learning: Binary Perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights

$$y = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector.

$$w = w + y^* \cdot f$$




Properties of Perceptrons


- Separability: true if some parameters get the training set perfectly correct
- Convergence: if the training is separable, perceptron will eventually converge (binary case)
- Mistake Bound: the maximum number of mistakes (binary case) related to the *margin* or degree of separability

$$\text{mistakes} < \frac{k}{\delta^2}$$

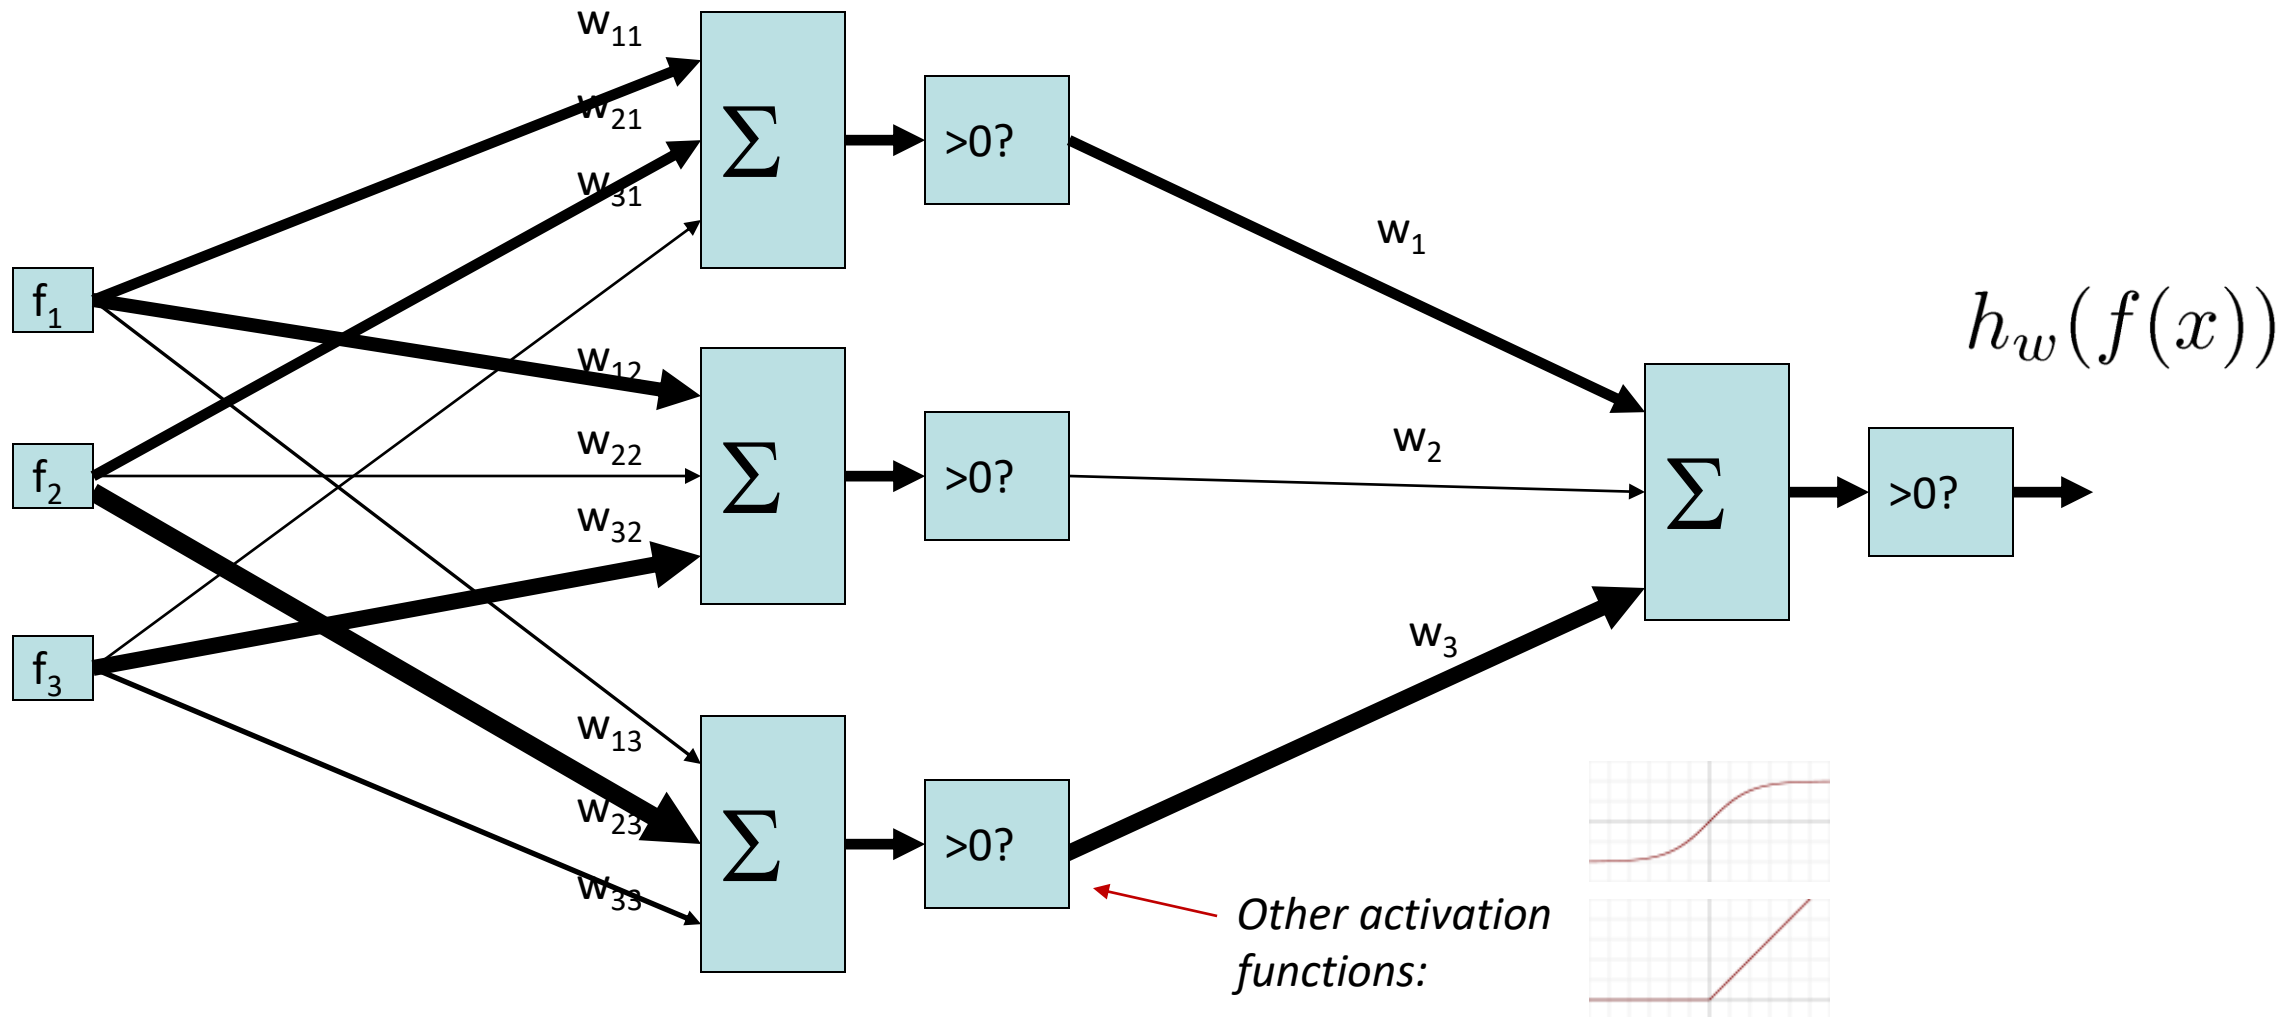
Separable



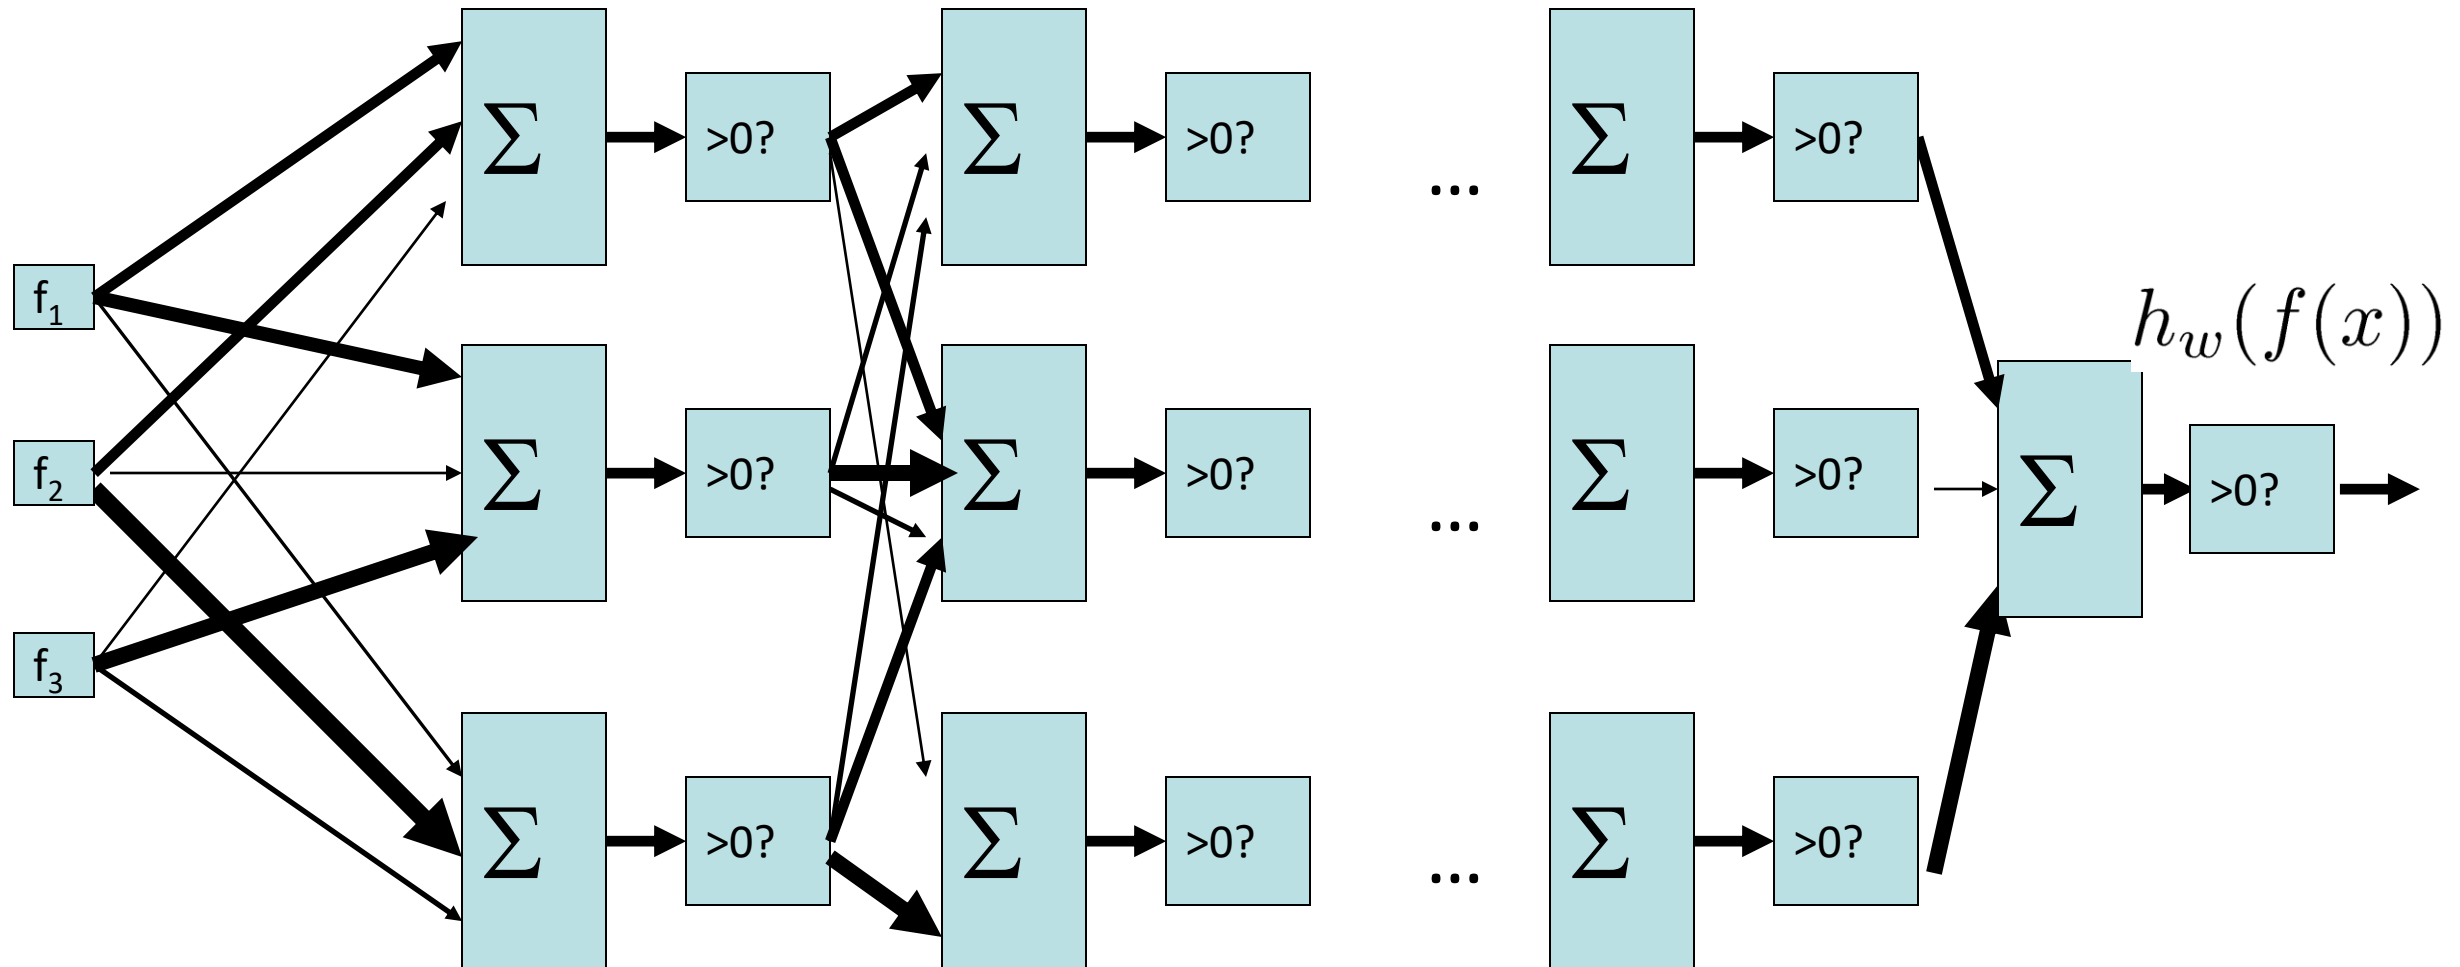
Non-Separable



Two-Layer Perceptron Network




Deep Neural Network



1-page Overview of Deep Learning

- Deep Learning
 - A large number of layers of neural networks
 - Ex. 1000 layers in ResNet
 - More complicated connections between layers
 - Ex. LSTM



1-page Overview of Deep Learning

- Deep Learning


Take CS280 Deep Learning!

- A large number of layers of neural networks
 - Ex. 1000 layers in ResNet
- More complicated connections between layers
 - Ex. LSTM
- Lots of new techniques and tricks
 - ReLU, Dropout, Batch Normalization, Adam, ...
- Big data
 - ImageNet (2009): 14 million images
 - NMT (a 2019 paper): 25 billion sentence pairs
- GPU parallelization
- Performance: superior to human experts in some tasks

More classification methods


- Naive Bayes
- Perceptron / Neural networks
- Decision trees / Random forest
- Support Vector Machines
- Nearest neighbors
- Model ensembles: bagging, boosting, etc.
-

Regression




Supervised learning

- To learn an unknown **target function** f
- Input: a **training set** of **labeled examples** (x_j, y_j) where $y_j = f(x_j)$
- Output: **hypothesis** h that is “close” to f
- Two types of supervised learning
 - Classification = learning f with discrete output value
 - Regression = learning f with real-valued output value



Linear Regression

Prediction: $h_w(x) = w_0 + w_1x$



Error on one instance: $|y - h_w(x)|$

Least squares: Minimizing squared error

- L2 loss function: sum of squared errors over all examples

$$L(\mathbf{w}) = \sum_i (y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2 = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- We want the weights \mathbf{w}^* that minimize loss
- Analytical solution: at \mathbf{w}^* the derivative of loss w.r.t. each weight is zero
 - \mathbf{X} is the data matrix (all the data, one example per row); \mathbf{y} is the vector of labels
 - $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

Least squares: Minimizing squared error

- $J(w) = \|Xw - y\|_2^2$ $\nabla J(w) = 2X^T(Xw - y)$

$$X^T X w - X^T y = 0$$

$$X^T X w = X^T y$$

$$w = (X^T X)^+ X^T y$$

Regularized Regression

- Overfitting is also possible in regression
 - Extreme case: n features, n training examples
- Regularization can be used to alleviate overfitting
- LASSO (Least Absolute Shrinkage and Selection Operator)

$$L(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_k |w_k|$$


- Ridge Regression

$$L(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_k w_k^2$$


Non-linear least squares

- No closed-form solution in general
- Numerical algorithms are typically used
 - Choose initial values for the parameters and then refine the parameters iteratively
 - Gradient descent
 - Gauss–Newton method
 - Limited-memory BFGS
 - Derivative-free methods
 - etc.


Overfitting in non-linear regression




Overfitting in non-linear regression



Overfitting in non-linear regression




Overfitting in non-linear regression



Overfitting in non-linear regression

Fit vs. complexity: a tradeoff

“*Ockham’s razor*”: prefer the *simplest* hypothesis consistent with the data



Summary

- Supervised learning:
 - Learning a function from labeled examples
- Classification: discrete-valued function
 - Naïve Bayes
 - Generalization and overfitting, smoothing
 - Perceptron
- Regression: real-valued function
 - Linear regression