



CS240 Algorithm Design and Analysis

Lecture 21

Randomized algorithms (Cont.)

Quan Li
Fall 2022
2022.11.29



Hash Tables





Hash Tables




- A hash table is a randomized data structure to efficiently implement a dictionary.
- Supports find, insert, and delete operations all in expected $O(1)$ time.
 - But in the worst case, all operations are $O(n)$.
 - The worst case is provably very unlikely to occur.
- A hash table does not support efficient min / max or predecessor / successor functions.
 - All these take $O(n)$ time on average.
- A practical, efficient alternative to binary search trees if only find, insert and delete needed.



Direct addressing



- Suppose we want to store (key, value) pairs, where keys come from a finite universe $U = \{0, 1, \dots, m-1\}$.
- Use an array of size m .
 - **insert**(k, v) Store v in array position k .
 - **find**(k) Return the value in array position k .
 - **delete**(k) Clear the value in array position k .
- All operations take $O(1)$ time.
- The problem is, if m is large, then we need to use a lot of memory.
 - Uses $O(|U|)$ space.
 - **Ex** For 32 bit keys, need 4 GB memory. For 64 bit keys, more memory than in world.
- If only need to store few values, lots of space wasted.



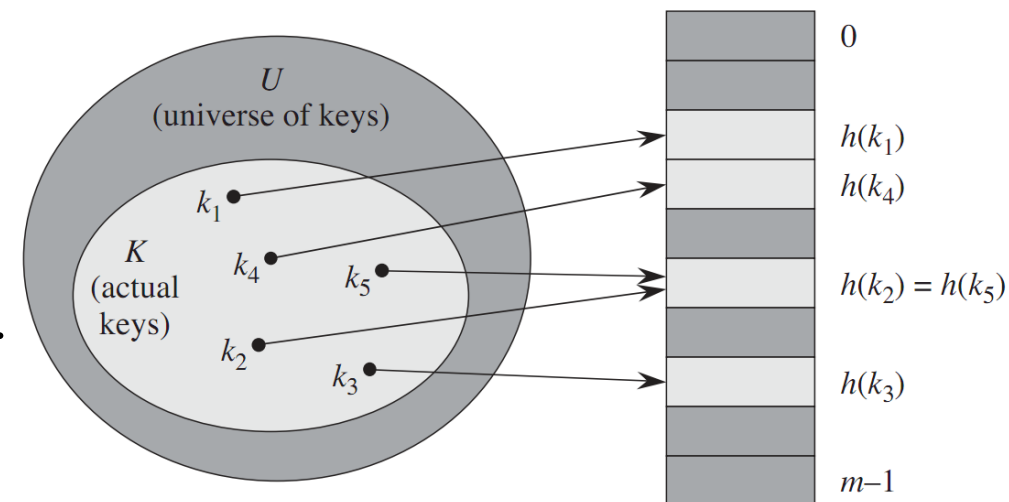
Source: Introduction to Algorithms, Cormen et al



Hash Table



- Similar to direct addressing but uses much less space.
- **Idea** Instead of storing directly at key's location, convert key to much smaller value, and store at this location.
- A hash table consists of the following.
 - A universe U of keys.
 - An array of T of size m .
 - A hashing function $h:U \rightarrow \{0,1,\dots,m-1\}$.
- We'll talk later about how to pick good hash functions.
- **insert(k, v)** Hash key to $h(k)$. Store v in $T[h(k)]$.
- **find(k)** Return the value in $T[h(k)]$
- **delete(k)** Delete the value in $T[h(k)]$
- Assuming $h(k)$ takes $O(1)$ time to compute, all ops still take $O(1)$ time. Uses $O(m)$ space.
- If $m \ll |U|$, then hashing uses much less space than direct addressing.
- However, our current scheme doesn't quite work, due to collisions.






Collisions



- We store a key at array position $h(k)$.
- But what if two keys hash to the same location, i.e., $k_1 \neq k_2$, but $h(k_1) = h(k_2)$?
 - This is called a collision.
- Collisions are unavoidable when $|U| > m$.
 - By Pigeonhole Principle, must exist at least two different keys in U that hash to same value.






Closed Addressing



- In closed addressing, every entry in hash table points to a linked list.
 - Keys that hash to the same location get added to the linked list.
 - For simplicity, we'll ignore values from now on and only focus on keys.
- **insert(k)** Add k to the linked list in $T[h(k)]$.
- **find(k)** Search the linked list in $T[h(k)]$ for k .
- **delete(k)** Delete k from the linked list in $T[h(k)]$.
- Suppose the longest list has length \hat{n} , and average length list is \bar{n} .
 - Each operation takes worst case $O(\hat{n})$ time.
 - An operation on a random key takes $O(\bar{n})$ time.






Load Factor



- The key to making closed addressing hashing fast is to make sure list lengths aren't too long.
- For this, we want the hash function to appear random.
 - Assume that any key is uniformly likely to be hashed to any table location.
- Suppose the hash table contains n items, and has size m .
- Then under the uniform hashing assumption, each table location has on average n/m keys.
 - Call $\alpha = n/m$ the load factor.
- So the average time for each operation is $O(\alpha)$.
- However, even with uniform hashing, in the worst case, all keys can hash to the same location. So, the worst-case performance is $O(n)$.






Picking a hash function



- We saw that we want hash functions to hash keys to “random” locations.
 - However, note that each hash function is itself a deterministic function, i.e. $h(k)$ always has the same value.
 - If $h(k)$ can produce different values, we can't find key k in the hash table anymore.
- It's hard to find such random hash functions, since we don't assume anything about the distribution of input keys.
- In practice, we use a number of heuristic functions.





Heuristic hash functions



- Assume the keys are natural numbers.
 - Convert other data types to numbers.
 - **Ex** To convert ASCII string to natural number, treat the string as a radix 128 number.
E.g. "pt" $\rightarrow (112 \times 128) + 116 = 14452$.
- **Division method** $h(k) = k \bmod m$
 - Often choose m a prime number not too close to a power of 2.
- **Multiplication method** $h(k) = \lfloor m (k A \bmod 1) \rfloor$, where A is some constant.
 - Knuth's suggestion is $A = \frac{\sqrt{5}-1}{2} \approx 0.618034 \dots$





Universal hashing



- As we said, regardless of the hash function, an adversary can choose a set of n inputs to make all operations $O(n)$ time.
- Universal hashing overcomes this using randomization.
 - No matter what the n input keys are, every operation takes $O(n/m)$ time in expectation, for a size m hash table.
 - Note $O(n/m)$ time is optimal.
- Instead of using a fixed hash function, universal hashing uses a random hash function, chosen from some set of functions H .
- Say H is a universal hash family if for any keys $x \neq y$

$$\Pr_{h \in H} [h(x) = h(y)] = 1/m$$

- So if we randomly choose a hash function from H and use it to hash any keys x, y , they have $1/m$ probability of colliding.
- Note the hash functions in H are not random. However, we choose which function to use from H randomly.






Universal hashing



- **Thm** Let H be a universal hash family. Let S be a set of n keys, and let $x \in S$. If $h \in H$ is chosen at random, then the expected number of $y \in S$ s.t. $h(x) = h(y)$ is n/m .
- **Proof** Say $S = \{x_1, \dots, x_n\}$.
 - Let X be a random variable equal to the number of $y \in S$ s.t. $h(x) = h(y)$.
 - Let $X_i = 1$ if $h(x_i) = h(x)$ and 0 otherwise.
 - $E[X_i] = \Pr_{h \in H}[h(x_i) = h(x)] \times 1 + \Pr_{h \in H}[h(x_i) \neq h(x)] \times 0 = 1/m$.
 - First equality follows by universal hashing property.
 - $E[X] = E[X_1] + \dots + E[X_n] = n/m$.





Constructing universal hash family 1

- Choose a prime number p such that $p > m$, and $p >$ all keys.
- Let $h_{ab}(k) = ((ak + b) \bmod p) \bmod m$.
- Let $H_{pm} = \{h_{ab} \mid a \in \{1, 2, \dots, p-1\}, b \in \{0, 1, \dots, p-1\}\}$.
- **Thm** H_{pm} is a universal hash family.
- **Proof** Let $x, y < p$ be two different keys. For a given h_{ab} let
$$r = (ax + b) \bmod p, \quad s = (ay + b) \bmod p$$
- We have $r \neq s$, because $r - s \equiv a(x - y) \bmod p \neq 0$, since neither a nor $x - y$ divide p .
- Also, each pair (a, b) leads to a different pair (r, s) , since
$$a = ((r - s)(x - y)^{-1} \bmod p), \quad b = (r - ax) \bmod p$$
 - Here, $(x - y)^{-1} \bmod p$ is the unique multiplicative inverse of $x - y$ in \mathbb{Z}_p^* .





Constructing universal hash family 2



- Since there are $p(p-1)$ pairs (a, b) and $p(p-1)$ pairs (r, s) with $r \neq s$, then a random (a, b) produces a random (r, s) .
- The probability x and y collide equals the probability $r \equiv s \pmod{m}$.
- For fixed r , number of $s \neq r$ s.t. $r \equiv s \pmod{m}$ is $(p-1)/m$.
- So for each r and random $s \neq r$, probability that $r \equiv s \pmod{m}$ is $((p-1)/m)/(p-1) = 1/m$.
- So $\Pr_{h_{ab} \in H_{pm}} [h_{ab}(x) = h_{ab}(y)] = 1/m$ and H_{pm} is universal.






Perfect hashing



- The hashing methods we've seen can ensure $O(1)$ expected performance but are $O(n)$ in the worst case due to collisions.
- However, if we have a fixed set of keys, perfect hashing can ensure no collisions at all.
 - Perfect hashing maintains a static set and allows $\text{find}(k)$ and $\text{delete}(k)$ in $O(1)$ time.
 - It doesn't support $\text{insert}(k)$.
- **Ex** The fixed set of keys may represent the file names on a non-writable DVD.






Perfect hashing



- Suppose we want to store n items with no collisions.
- Perfect hashing uses two levels of universal hashing.
 - The first layer hash table has size $m = n$.
 - Use first layer hash function h to hash key to a location in T .
 - Each location j in T points to a hash table S_j with hash function h_j .
 - If n_j keys hash to location j , the size of S_j is $m_j = n_j^2$.
- We'll ensure there are no collisions in the secondary hash tables S_1, \dots, S_m .
 - So all operations take worst case $O(1)$ time.
- Overall the space use is $O(m + \sum_{j=1}^m n_j^2)$.
 - We'll show this is $O(n) = O(m)$.
 - So perfect hashing uses same amount of space as normal hashing.

- $h(k) = ((3k + 42) \bmod 101) \bmod 9$
- $h_j(k) = ((a_j k + b_j) \bmod 101) \bmod m_j$






Avoiding collisions



- **Lemma** Suppose we store n keys in a hash table of size $m = n^2$ using universal hashing. Then with probability $\geq 1/2$ there are no collision.
- **Proof** There are $\binom{n}{2}$ pairs of keys that can collide.
 - Each collision occurs with probability $1/m = 1/n^2$, by universal hashing.
 - So the expected number of collisions is $\frac{\binom{n}{2}}{n^2} \leq \frac{1}{2}$.
 - By Markov's inequality the $\Pr[\# \text{ collisions} \geq 1] \leq E[\# \text{ collisions}] \leq 1/2$.
- When building each hash table S_j , there's $< 1/2$ probability of having any collisions.
 - If collisions occur, pick another random hash function from the universal family and try again.
 - In expectation, we try twice before finding a hash function causing no collisions.

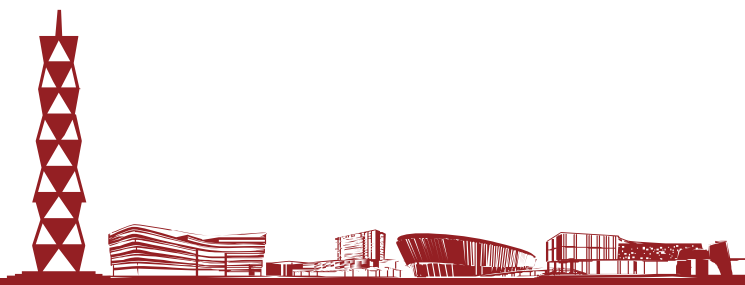


- **Lemma** Suppose we store n keys in a hash table of size $m=n$. Then the secondary hash tables use space $E[\sum_{j=0}^{m-1} n_j^2] < 2n$, where n_j is the number of keys hashing to location j .
- **Proof** $E[\sum_{j=0}^{m-1} n_j^2] = E[\sum_{j=0}^{m-1} (n_j + 2 \binom{n_j}{2})] = E[\sum_{j=0}^{m-1} n_j] + 2 E[\sum_{j=0}^{m-1} \binom{n_j}{2}]$
- $\sum_{j=0}^{m-1} \binom{n_j}{2}$ is the total number of pairs of hash keys which collide in the first level hash table.
 - By universal hashing, this equals $\binom{n}{2} \frac{1}{m} = \frac{n-1}{2}$.
- $E[\sum_{j=0}^{m-1} n_j] = n$.
- So $E[\sum_{j=0}^{m-1} n_j^2] = n + \frac{2(n-1)}{2} < 2n$.





Bloom Filters



Approximate Sets



- A Bloom filter is a data structure that can implement a set.
 - It only keeps track of which keys are present, not any values associated to keys.
 - It supports insert and find operations.
 - It doesn't support delete operations.
- Bloom filters use less memory than hash tables or other ways of implementing sets.
- However, Bloom filters are approximate.
 - It can produce false positives: it says an element is present even though it's not.
 - We can bound the probability of false positives.
 - But it doesn't produce false negatives: if it says an element isn't present, then it's not.





Bloom Filter Applications



- Suppose we have a big database and querying it to check if an item is present is expensive.
- We store the set of items in the database using a Bloom filter.
 - This tells us whether an item is in database or not.
- If filter says an item's not present, it's definitely not in the database.
 - So, no need to do an expensive query.
- If filter says an item is present, then either item is present, or there's false positive.
 - When we query the database, there's a small probability we waste time querying for a nonexistent item.
- Overall, we save time by checking Bloom filter first before querying database.







Bloom Filter Applications



- Consider a P2P network, where each node stores some files.
- If you want to get a file, you need to know which nodes have it.
- Keeping a list of all items stored at each node is too expensive.
- Instead, for every other node, keep a Bloom filter of its files.
- If filter says no for a node, it definitely doesn't have the file.
- If filter says yes, then either node has the file, or there's false positive and we make a useless request.
- Overall, we save space, and also won't waste much communication because we rarely make useless requests.






Bloom Filters



- A Bloom filter consists of
 - An array A of size m , initially all 0's.
 - k independent hash functions h_1, \dots, h_k , each mapping from keys to $\{1, \dots, m\}$.
- To store key x
 - Set $A[h_1(x)], A[h_2(x)], \dots, A[h_k(x)]$ all to 1.
 - Some locations can get set to 1 multiple times; that's fine.
- To check if key x is in the set
 - Read array locations $A[h_1(x)], A[h_2(x)], \dots, A[h_k(x)]$.
 - If all the values are 1, output "x is in set".
 - Otherwise, output "x is not in set".



A Bloom filter with $k=3$ hash functions storing 2 items.






Correctness



- Let's look at the correctness of the search function.
- If search for x returns no, then at least one of $A[h_1(x)], \dots, A[h_k(x)]$ equals 0.
 - So x cannot be in the set, because if x had been inserted into the set, then we would have $A[h_1(x)] = \dots = A[h_k(x)] = 1$.
 - So there are no false negatives.
- If search for x returns yes, then $A[h_1(x)] = \dots = A[h_k(x)] = 1$.
 - So either x was inserted into the set.
 - Or we inserted some keys that hashed to the same k locations as x .
 - So it looks as if x was inserted, even though it wasn't.
 - This is a false positive. We'll bound the probability this happens.

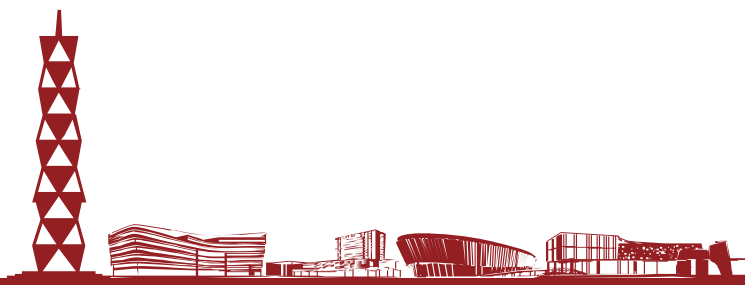




False Positive Probability 1



- False positive probability depends on k (number of hash functions), m (size of table) and n (number of keys inserted).
- Assume hash functions hash keys to random locations.
- When inserting one key, we set k random locations to 1.
- Fix any position i . Probability i is set to 1 by a hash function is $1/m$, so probability i stays 0 is $1 - 1/m$.
 - After k hashes, probability i still 0 is $(1 - 1/m)^k$.
 - To insert n items, we used nk hashes. So, probability i still 0 after all these is $p = (1 - 1/m)^{nk}$.
- We now use an approximation $\left(1 - \frac{1}{m}\right)^{nk} \approx e^{-\frac{nk}{m}}$.





False Positive Probability 2



- So, probability any position i is 1 after n keys inserted is $1 - p \approx 1 - e^{-\frac{nk}{m}}$.
- Since there are m positions in the array, assume there are $(1-p)m$ positions that are 1.
 - This isn't quite correct. The actual number of 1's in the array is a random variable, whose expectation is $(1-p)m$.
 - However, we can make the argument rigorous by showing that the actual number of 1's is $(1 - p)m \pm \sqrt{m \log m}$ with high probability.
- We only get a false positive if when we check k random locations, they're all 1.
 - Probability is $f = (1 - p)^k \approx \left(1 - e^{-\frac{nk}{m}}\right)^k$.

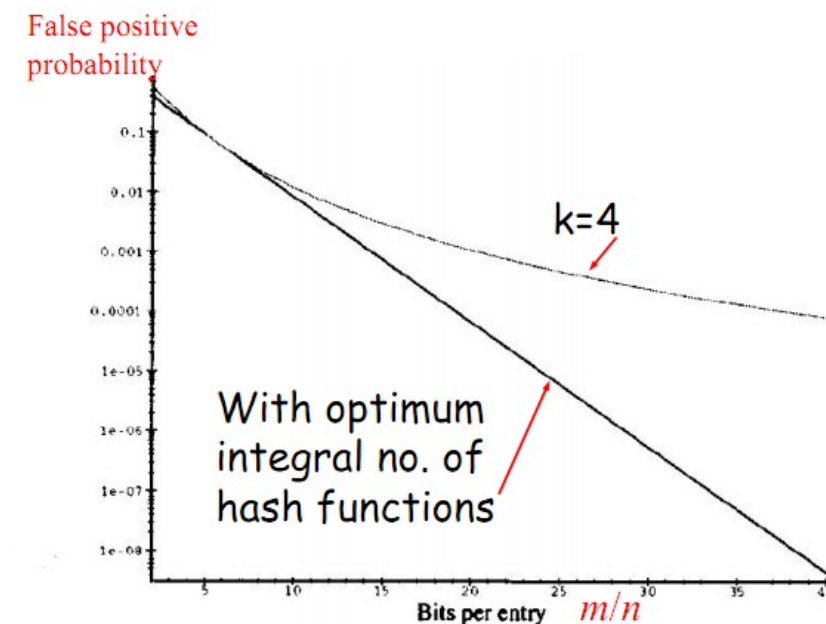




False Positive Probability 3



- Notice the false prob. $(1 - e^{-\frac{nk}{m}})^k$ is a function of k , the number of hash functions we use.
- We find k to minimize the false positive prob. by differentiating f wrt k and solving.
- The optimum k is $\frac{m \ln(2)}{n}$, which leads to $f = \left(\frac{1}{2}\right)^k \approx 0.6185^{\frac{m}{n}}$.
 - Notice that m/n is the average number of bits per item. So error rate decreases exponentially in space usage.






Improvements




- Right now, Bloom filters can't handle deletes.
 - Say keys k_1 , k_2 hash to two overlapping sets of locations. If you delete k_1 by setting some of its locations to 0, you could also delete k_2 .
- Deletes can be done by storing a count of how many keys hashed to that location, and inc / dec the counts when inserting or deleting.
 - But this uses more memory.
 - Also, what if the counts overflow?





String Equality and Fingerprinting






String Equality and Fingerprinting



- Alice and Bob both have copies of a database.
- They want to keep the database consistent, so they want to check if their copies are the same.
 - If you think of the databases as strings, they want to check if their strings are equal.
- But transferring the entire database is expensive.
- Instead, they calculate a small value called a fingerprint of their databases.
 - If the fingerprints are the different, then their databases are definitely different.
 - If the fingerprints are the same, then the databases are probably the same; but there's a small probability they're actually different.
- Transferring the fingerprint is much cheaper than the database.

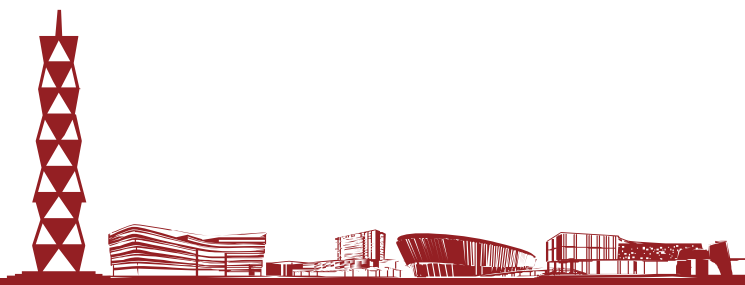




Fingerprinting



- Let Alice and Bob's databases be the bit sequences (a_1, \dots, a_n) and (b_1, \dots, b_n) .
- View these as n -bit integers $a = \sum_{i=1}^n a_i * 2^{i-1}$ and $b = \sum_{i=1}^n b_i * 2^{i-1}$.
- The fingerprint $F(a) = a \bmod p$, for a specially chosen prime number p .
 - Alice transfers $F(a)$ to Bob, and Bob compares it to his fingerprint $F(b) = b \bmod p$.
 - Since $F(a) < p$, transferring the fingerprint only takes $O(\log p)$ bits, instead of n .





Correctness



- No false positives (positive means " $a \neq b$ ").
 - If $F(a) \neq F(b)$, then $a \neq b$.
- False negatives are possible.
 - If $F(a)=F(b)$, then $a \bmod p = b \bmod p$.
 - So either $a=b$, or $a \neq b$ but p divides $(a-b)$.
- We can't avoid false negatives. But we can minimize the probability it occurs.
- Pick a random p .
 - If $a \neq b$, then probably p doesn't divide $(a-b)$, so probably $F(a) \neq F(b)$ and we'll detect a and b are different.
 - Bigger p decreases false negative probability.
 - But we don't want to make p too big, since we have to transfer $O(\log p)$ bits.

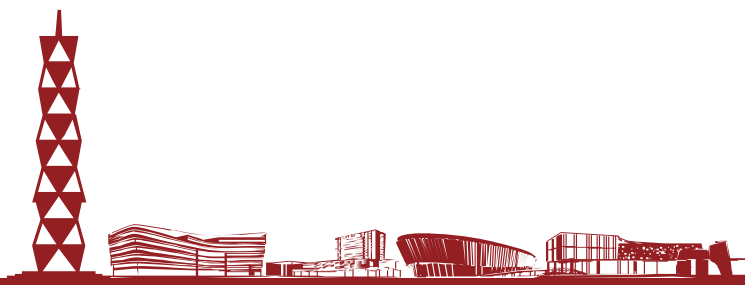




Correctness



- To analyze the false negative probability, we use two facts from number theory.
- **Lemma** Any number t has at most $\log_2(t)$ distinct prime divisors.
- **Proof** Each divisor is ≥ 2 , and their product is $\leq t$. If there were more than $\log_2(t)$ divisors, their product would be $> 2^{\log_2(t)} = t$, contradiction.
- Recall $a = \sum_{i=1}^n a_i * 2^{i-1}$ and $b = \sum_{i=1}^n b_i * 2^{i-1}$
- So $a-b < 2^n$, and so $a-b$ has at most n distinct prime divisors.





Correctness



- **Prime Number Theorem** Given any number t , the number of primes smaller than t is $\sim t / \ln(t)$.
- The PNT allows us to efficiently generate a random prime.
 - Picking a number less than t at random, it has a $1/\ln(t)$ probability of being prime.
 - We can check if a number is prime using the Rabin-Miller primality test.
 - If number is prime, it always passes the test.
 - If number is composite, there's small probability it's declared a prime.
 - Run the test few more times to exponentially decrease false positive probability.
 - So with high probability, we can tell if a number is prime.





Correctness



- Let $t = n^2 \ln(n)$. The number of primes less than t is $\approx \frac{t}{\ln(t)} = \frac{n^2 \ln(n)}{2 \ln(n) + \ln \ln(n)} = O(n^2)$.
- Pick a random prime p less than t .
- We get a false negative if $a \neq b$ but p divides $(a-b)$.
 - We saw earlier that $a-b$ has $< n$ prime divisors, and p must be one of these.
 - But p is randomly chosen from $O(n^2)$ primes less than t .
 - So false negative probability $\leq n/O(n^2) = O(1/n)$.
- We transfer $\log(p) \leq \log(t) = O(\log n)$ bits.
- Transferring $O(\log n)$ bits gets $O(1/n)$ probability of error. If we want perfect accuracy, we need to transfer the entire database, $O(n)$ bits.





Next Time: Randomized algorithms (Cont.)

