

Online Lecture Notes

Prof. Boris Houska

April 26, 2022

1 Numerical Methods for Linear Equation Systems

In the last lecture we have learned about Gauss elimination, which proceeds by computing upper- and lower triangular matrices R and L associated to a given invertible square matrix A such that

$$A = LR$$

This method can be refined for sparse matrices A , too, but exploiting the zeros in the coefficient scheme explicitly (also we may permute rows and columns of A appropriately). A prototype example for a sparse LR-decomposition algorithm is obtained by having a closer look at tri-diagonal matrices of the form

$$A = \begin{pmatrix} a_1 & b_1 & & \\ c_1 & a_2 & \ddots & \\ & \ddots & \ddots & b_{n-1} \\ & & c_{n-1} & a_n \end{pmatrix} \in \mathbb{R}^{n \times n}.$$

All the “empty spaces” are filled with zeros that do not have to be stored explicitly. The main observation now is that L and R inherit the sparsity of A (in this case directly). In detail, this means that L and R should be matrices of the form

$$L = \begin{pmatrix} 1 & & & \\ \gamma_1 & 1 & & \\ & \ddots & \ddots & \\ & & \gamma_{n-1} & 1 \end{pmatrix} \quad \text{and} \quad R = \begin{pmatrix} \alpha_1 & b_1 & & \\ & \alpha_2 & \ddots & \\ & & \ddots & b_{n-1} \\ & & & \alpha_n \end{pmatrix}$$

Notice that we can work out a recursion for computing the coefficients α_i and γ_i . Essentially, this is in complete analogy to the dense Gauss elimination with the only difference being that we exploit the particular sparsity pattern of A .

For instance, in order to see we can write

$$\begin{aligned}
R = L^{-1}A &= \begin{pmatrix} 1 & & & \\ \gamma_1 & 1 & & \\ & \ddots & \ddots & \\ & & \gamma_{n-1} & 1 \end{pmatrix}^{-1} \begin{pmatrix} a_1 & b_1 & & \\ c_1 & a_2 & \ddots & \\ & \ddots & \ddots & b_{n-1} \\ & & c_{n-1} & a_n \end{pmatrix} \\
&= \begin{pmatrix} 1 & & & \\ -\gamma_1 & 1 & & \\ & \ddots & \ddots & \\ & & -\gamma_{n-1} & 1 \end{pmatrix} \begin{pmatrix} a_1 & b_1 & & \\ c_1 & a_2 & \ddots & \\ & \ddots & \ddots & b_{n-1} \\ & & c_{n-1} & a_n \end{pmatrix} \quad (1) \\
&= \begin{pmatrix} a_1 & b_1 & & \\ c_1 - \gamma_1 a_1 & a_2 - \gamma_1 b_1 & \ddots & \\ & \ddots & \ddots & b_{n-1} \\ & & c_{n-1} - \gamma_{n-1} a_{n-1} & a_n - \gamma_{n-1} b_{n-1} \end{pmatrix}
\end{aligned}$$

Since we want to ensure that this matrix is upper triangular, we need to choose the coefficients γ_i such that

$$\forall i \in \{1, \dots, n-1\}, \quad c_i - \gamma_i a_i = 0 \quad \implies \quad \gamma_i = \frac{c_i}{a_i}$$

Additionally, we find that $\alpha_1 = a_1$ and $\alpha_i = a_i - \gamma_{i-1} b_{i-1}$ for $i \in \{2, \dots, n\}$. With this we have found a way to compute a sparse LR-decomposition of the tri-diagonal matrix A with computational complexity

$$O(n) .$$

Notice that this is much faster than using a dense LR-decomposition solver, which would have run-time $O(n^3)$.

Summary: in practice it is very important to exploit the sparsity pattern of the matrix A when implementing an LR-decomposition, since exploiting the pattern of A can improve the run-time by orders of magnitude! This is especially important when n is large!

1.1 QR-Decomposition

Recall that if the matrix is given columnwise, $A = (a_1, a_2, a_3, \dots, a_n) \in \mathbb{R}^{n \times n}$, we can run a Gram-Schmidt algorithm to find an orthonormal basis if the span of the vectors $a_1, a_2, a_3, \dots, a_n$. The corresponding orthonormal vectors q_1, q_2, \dots, q_n are constructed recursively by the Gram-Schmidt algorithm such that

$$a_1 = (a_1^\top q_1) q_1 \quad (2)$$

$$a_2 = (a_2^\top q_1) q_1 + (a_2^\top q_2) q_2 \quad (3)$$

$$\vdots \quad (4)$$

$$a_n = (a_n^\top q_1) q_1 + (a_n^\top q_2) q_2 + \dots + (a_n^\top q_n) q_n \quad (5)$$

Notice that the coefficients $r_{i,j} = a_i^\top q_j$ are computed by the Gram-Schmidt recursion, too. This means that we can write the above equations in the form

$$a_1 = r_{1,1}q_1 \quad (6)$$

$$a_2 = r_{1,2}q_1 + r_{2,2}q_2 \quad (7)$$

$$\vdots \quad (8)$$

$$a_n = r_{1,n}q_1 + r_{2,n}q_2 + \dots + r_{n,n}q_n \quad (9)$$

Notice that we can write these equation more elegantly in matrix form

$$\underbrace{(a_1, a_2, \dots, a_n)}_{=A} = \underbrace{(q_1, q_2, \dots, q_n)}_{=Q} \underbrace{\begin{pmatrix} r_{1,1} & r_{1,2} & \dots & r_{1,n} \\ & r_{2,2} & \dots & r_{2,n} \\ & & \ddots & \vdots \\ & & & r_{n,n} \end{pmatrix}}_{=R}$$

Thus, in summary, the Gram-Schmidt algorithm automatically generates a QR-decomposition of the form

$$A = QR,$$

where R is an upper triangular matrix and Q is an ortogonal matrix, since

$$Q^\top Q = (q_1, q_2, \dots, q_n)^\top (q_1, q_2, \dots, q_n) = \begin{pmatrix} q_1^\top q_1 & q_1^\top q_2 & \dots & q_1^\top q_n \\ q_2^\top q_1 & q_2^\top q_2 & \dots & q_2^\top q_n \\ \vdots & \vdots & \ddots & \vdots \\ q_n^\top q_1 & q_n^\top q_2 & \dots & q_n^\top q_n \end{pmatrix} = I.$$

The corresponding linear equation system of the form

$$Ax = b$$

can then be solved by first storing the QR -decomposition of A and then using that

$$Ax = b \implies QRx = b \implies Q^\top QRx = Q^\top b \implies Rx = Q^\top b$$

where the latter equation can be solved by backward substitution in $O(n^2)$ operations, since R is upper triangular matrix.

Summary: the QR decomposition is in many ways very similar to the LR decomposition, but one difference is that QR decomposition can also be computed for degenerate (as well as non-square) matrices A , since the Gram-Schmidt algorithm is completely generic—we can always detect an orthonogal basis of the span of the columns of A no matter what the dimensional of this span is! This means that in practice QR decompositions are eventually a bit more well-conditioned than LR decompositions, since they can be used on degerate matrices, too.

1.2 Cholesky Decomposition

Recall first that a matrix A is called symmetric positive definite if

$$A = A^\top \quad \text{and} \quad \forall v \in \mathbb{R}^n, \quad v^\top A v > 0.$$

Also recall that the eigenvalues of a symmetric matrix are real-valued and that A is positive definite if all eigenvalues of A are strictly positive. The goal of the following considerations is to develop efficient algorithms for finding a Cholesky decomposition of the form

$$A = LDL^\top$$

with L being a lower triangular matrix with ones on the diagonal and D a diagonal matrix. Notice that such a decomposition exists, since we could also start with a LR-decomposition of A ,

$$A = LR$$

and then denote by D the part of R . Then we must have $R = DL^\top$ due to symmetry. But alternatively, we can find L and D by a direct comparison of coefficients, which leads to Cholesky's algorithm. Let us work this out by setting

$$L = \begin{pmatrix} 1 & & & \\ L_{2,1} & 1 & & \\ \vdots & & \ddots & \\ L_{n,1} & \dots & L_{n,n-1} & 1 \end{pmatrix} \quad \text{and} \quad D = \begin{pmatrix} D_1 & & & \\ & D_2 & & \\ & & \ddots & \\ & & & D_n \end{pmatrix}.$$

This yields

$$\begin{aligned} LDL^\top &= \begin{pmatrix} 1 & & & \\ L_{2,1} & 1 & & \\ \vdots & & \ddots & \\ L_{n,1} & \dots & L_{n,n-1} & 1 \end{pmatrix} \begin{pmatrix} D_1 & & & \\ & D_2 & & \\ & & \ddots & \\ & & & D_n \end{pmatrix} \begin{pmatrix} 1 & & & \\ L_{2,1} & 1 & & \\ \vdots & & \ddots & \\ L_{n,1} & \dots & L_{n,n-1} & 1 \end{pmatrix}^\top \\ &= \begin{pmatrix} D_1 & & & \\ D_1 L_{2,1} & D_2 & & \\ \vdots & & \ddots & \\ D_1 L_{n,1} & \dots & D_{n-1} L_{n,n-1} & D_n \end{pmatrix} \begin{pmatrix} 1 & L_{2,1} & \dots & L_{n,1} \\ & 1 & & \\ & & \ddots & L_{n,n-1} \\ & & & 1 \end{pmatrix} \\ &= \begin{pmatrix} D_1 & DL_{2,1} & & \\ D_1 L_{2,1} & D_1 L_{2,1}^2 + D_2 & & \text{sym} \\ \vdots & & \ddots & \\ D_1 L_{n,1} & \dots & \dots & D_1 L_{n,1}^2 + D_2 L_{n,2}^2 + \dots + D_{n-1} L_{n,n-1}^2 + D_n \end{pmatrix} \end{aligned}$$

We can compare the coefficients of this matrix with the coefficients of A . This leads to the equations

$$A_{1,1} = D_1 \quad (10)$$

$$A_{2,1} = D_1 L_{1,2} \quad (11)$$

$$A_{2,2} = D_1 L_{2,1}^2 + D_2 \quad (12)$$

$$\vdots \quad (13)$$

Notice that this equation system can be solved recursively with respect to the coefficients of D and L . This yields

$$D_1 = A_{1,1} \quad (14)$$

$$L_{1,2} = \frac{A_{2,1}}{D_1} \quad (15)$$

$$D_2 = A_{2,2} - D_1 L_{2,1}^2 \quad (16)$$

$$\vdots \quad (17)$$

This leads to the general recursion

$$D_j = A_{j,j} - \sum_{k=1}^{j-1} D_k L_{j,k}^2 \quad \text{and} \quad L_{i,j} = \frac{1}{D_j} \left(A_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k} D_k \right)$$

which can be evaluated in $O(n^3)$ operations, but the interesting aspect is that the symmetry of the matrix A is exploited explicitly by the recursion.

1.3 Summary

In this lecture we have discussed three types of numerical algorithm for solving linear equation systems:

1. LR -decomposition (Gauss elimination): applicable to invertible square matrices
2. QR -decomposition (Gram-Schmidt algorithm): applicable to any matrix
3. LDL^\top -decomposition (Cholesky algorithm): applicable to symmetric positive definite matrices.