

Part II: Least Squares

LS Solution

Theorem 1. A vector \mathbf{x}_{LS} is an optimal solution to the LS problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2$$

if and only if it satisfies

$$\mathbf{A}^T \mathbf{A} \mathbf{x}_{\text{LS}} = \mathbf{A}^T \mathbf{y}. \quad (*)$$

- the optimality condition in $(*)$ is true for any \mathbf{A} , not just full-column rank \mathbf{A}
- suppose that \mathbf{A} has full-column rank
 - $(*)$ is a symmetric PD linear system
 - the Gram matrix $\mathbf{A}^T \mathbf{A}$ is nonsingular (easy to verify)
 - the solution to $(*)$ is uniquely given by $\mathbf{x}_{\text{LS}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$
- $(*)$ is called the **normal equations**
- the same result holds for the complex case, viz., $\mathbf{A}^H \mathbf{A} \mathbf{x}_{\text{LS}} = \mathbf{A}^H \mathbf{y}$
- LS is a unconstrained optimization problem with a quadratic objective
- **Linear regression** is a linear approach to modelling the relationship between a dependent variable and one or more independent variables (i.e., data fitting) which can be estimated using LS method (or based LS loss function).

LS Solution

- there are many ways to prove Theorem 1, such as by the projection theorem, by optimization, or by singular value decomposition
 - projection theorem
 - optimization
 - singular value decomposition (cf. [Singular Value Decomposition Topic](#))
 - more...

LS and Projection Theorem

- Theorem 1 can be shown using the projection theorem
- let \mathbf{x}_{LS} be an LS solution, and observe that

$$\Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \arg \min_{\mathbf{z} \in \mathcal{R}(\mathbf{A})} \|\mathbf{z} - \mathbf{y}\|_2^2 = \mathbf{A}\mathbf{x}_{\text{LS}}$$

- by the projection theorem (Theorem 2 in [Basic Concepts Topic](#)), we have

$$\begin{aligned} \Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \mathbf{A}\mathbf{x}_{\text{LS}} &\iff \mathbf{z}^T(\mathbf{A}\mathbf{x}_{\text{LS}} - \mathbf{y}) = 0 \text{ for all } \mathbf{z} \in \mathcal{R}(\mathbf{A}) \\ &\iff \mathbf{x}^T \mathbf{A}^T(\mathbf{A}\mathbf{x}_{\text{LS}} - \mathbf{y}) = 0 \text{ for all } \mathbf{x} \in \mathbb{R}^n \\ &\iff \mathbf{A}^T(\mathbf{A}\mathbf{x}_{\text{LS}} - \mathbf{y}) = \mathbf{0} \\ &\iff \mathbf{A}^T \mathbf{A}\mathbf{x}_{\text{LS}} = \mathbf{A}^T \mathbf{y} \end{aligned}$$

Orthogonal Projections

- the projections of \mathbf{y} onto $\mathcal{R}(\mathbf{A})$ and $\mathcal{R}(\mathbf{A})^\perp$ (for full column-rank \mathbf{A}) are, resp.,

$$\Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \mathbf{A}\mathbf{x}_{\text{LS}} = \mathbf{A}(\mathbf{A}^T\mathbf{A})^{-1}\mathbf{A}^T\mathbf{y}$$

$$\Pi_{\mathcal{R}(\mathbf{A})^\perp}(\mathbf{y}) = \mathbf{y} - \Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = (\mathbf{I} - \mathbf{A}(\mathbf{A}^T\mathbf{A})^{-1}\mathbf{A}^T)\mathbf{y}$$

- the **orthogonal projector** of \mathbf{A} is defined as

$$\mathbf{P}_\mathbf{A} = \mathbf{A}(\mathbf{A}^T\mathbf{A})^{-1}\mathbf{A}^T$$

the **orthogonal complement projector** of \mathbf{A} is defined as

$$\mathbf{P}_\mathbf{A}^\perp = \mathbf{I} - \mathbf{A}(\mathbf{A}^T\mathbf{A})^{-1}\mathbf{A}^T = \mathbf{I} - \mathbf{P}_\mathbf{A}.$$


- obviously, we want to write $\Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \mathbf{P}_\mathbf{A}\mathbf{y}$, $\Pi_{\mathcal{R}(\mathbf{A})^\perp}(\mathbf{y}) = \mathbf{P}_\mathbf{A}^\perp\mathbf{y}$
- note: a more general definition for orthogonal projectors for general \mathbf{A} will be studied in **Singular Value Decomposition Topic**

Orthogonal Projections

- properties of \mathbf{P}_A (same properties apply to \mathbf{P}_A^\perp):
 - \mathbf{P}_A is idempotent; i.e., $\mathbf{P}_A^2 = \mathbf{P}_A \mathbf{P}_A = \mathbf{P}_A$
 - $\mathbf{P}_A = \mathbf{P}_A^T$ for real A ($\mathbf{P}_A = \mathbf{P}_A^H$ for complex A)
- additional properties that will be revealed in later lectures:
 - the eigenvalues of \mathbf{P}_A are either zero or one (cf. [Eigendecomposition Topic](#))
 - \mathbf{P}_A can be written as $\mathbf{P}_A = \mathbf{U}_1 \mathbf{U}_1^T = \mathbf{P}_{\mathbf{U}_1}$ for some semi-orthogonal \mathbf{U}_1 (cf. [Singular Value Decomposition Topic](#))
 - * we can also prove it here:
 - there always exists a semi-orthogonal \mathbf{U}_1 such that $\mathcal{R}(\mathbf{A}) = \mathcal{R}(\mathbf{U}_1)$
 - $\Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \Pi_{\mathcal{R}(\mathbf{U}_1)}(\mathbf{y}) = \mathbf{U}_1 \mathbf{U}_1^T \mathbf{y}$
 - as $\Pi_{\mathcal{R}(\mathbf{A})}(\mathbf{y}) = \Pi_{\mathcal{R}(\mathbf{U}_1)}(\mathbf{y})$ holds for any \mathbf{y} , or $(\mathbf{P}_A - \mathbf{U}_1 \mathbf{U}_1^T) \mathbf{y} = \mathbf{0}$ for any \mathbf{y} , we must have $\mathbf{P}_A = \mathbf{U}_1 \mathbf{U}_1^T$
 - * suppose $\mathbf{U}_1 \in \mathbb{R}^{m \times n}$, $\Pi_{\mathcal{R}(\mathbf{U}_1)}(\mathbf{y}) = \mathbf{U}_1 \mathbf{U}_1^T \mathbf{y} = \sum_{i=1}^n (\mathbf{u}_{1i}^T \mathbf{y}) \mathbf{u}_{1i}$

More on Projections

- **Definition:** a square matrix \mathbf{P} is called a **projection matrix (projector)** if it is idempotent, i.e, $\mathbf{P}^2 = \mathbf{P}$.
 - easy to understand from a geometric view
 - projection onto $\mathcal{R}(\mathbf{P})$
 - **complement projector:** $\mathbf{I} - \mathbf{P}$; projection onto $\mathcal{R}(\mathbf{I} - \mathbf{P}) = \mathcal{N}(\mathbf{P})$
 - if $\mathbf{P} \in \mathbb{R}^{m \times m}$, $\mathbb{R}^m = \mathcal{R}(\mathbf{P}) \oplus \mathcal{N}(\mathbf{P})$
- in practice, when we say projection, it mostly refers to **orthogonal projection**
 - $\mathcal{R}(\mathbf{P})^\perp = \mathcal{N}(\mathbf{P}^T) = \mathcal{N}(\mathbf{P})$ (the complement is the orthogonal complement)
- a projection matrix that is not an orthogonal projection matrix is called an **oblique projection matrix**



Pseudo-Inverse

- the [pseudo-inverse](#) of a full-column-rank \mathbf{A} is defined as

$$\mathbf{A}^\dagger = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T.$$

- \mathbf{A}^\dagger satisfies $\mathbf{A}^\dagger \mathbf{A} = \mathbf{I}$, but not necessarily $\mathbf{A} \mathbf{A}^\dagger = \mathbf{I}$
- $\mathbf{A}^\dagger \mathbf{y}$ is the LS solution
- the orthogonal projector of \mathbf{A} becomes $\mathbf{P}_\mathbf{A} = \mathbf{A} \mathbf{A}^\dagger$
- note: a more general definition of the pseudo-inverse for general \mathbf{A} will be studied later (cf. [Singular Value Decomposition Topic](#))

LS and Convex Optimization

- we can also prove the LS optimality condition (Theorem 1) by optimization
- the **gradient** of a continuously differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

where $\frac{\partial f}{\partial x_i}$ is the partial derivative

- **Fact:** consider an unconstrained optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is continuously differentiable

- suppose f is **convex** (we skip the def. here). A point \mathbf{x}^* is an optimal solution if and only if $\nabla f(\mathbf{x}^*) = \mathbf{0}$
- for non-convex f , any point $\hat{\mathbf{x}}$ satisfying $\nabla f(\hat{\mathbf{x}}) = \mathbf{0}$ is a stationary point

LS and Convex Optimization

- **Fact:** consider a quadratic function

$$f(\mathbf{x}) = \frac{1}{2}\mathbf{x}^T \mathbf{R} \mathbf{x} + \mathbf{q}^T \mathbf{x} + c,$$

where $\mathbf{R} \in \mathbb{S}^{n \times n}$.

- $\nabla f(\mathbf{x}) = \mathbf{R} \mathbf{x} + \mathbf{q}$
- f is convex if \mathbf{R} is positive semidefinite (PSD); for now it suffices to know that if \mathbf{R} takes the form $\mathbf{R} = \mathbf{A}^T \mathbf{A}$ for some \mathbf{A} , it is PSD (easy to verify)
- the LS objective function is

$$f(\mathbf{x}) = \|\mathbf{y} - \mathbf{A} \mathbf{x}\|_2^2 = \mathbf{x}^T \mathbf{A}^T \mathbf{A} \mathbf{x} - 2(\mathbf{A}^T \mathbf{y})^T \mathbf{x} + \|\mathbf{y}\|_2^2.$$

Using the above optimization facts, \mathbf{x}_{LS} is an LS optimal solution if and only if $\mathbf{A}^T \mathbf{A} \mathbf{x}_{\text{LS}} - \mathbf{A}^T \mathbf{y} = \mathbf{0}$.

- LS problem is one **quadratic programming (QP)** problem
- the normal equation is equivalent to the first-order optimality (or KKT) condition

LS and Convex Optimization

- using optimization results is handy in some (actually, many) cases
- example ([Tikhonov regularization](#)): consider a regularized LS problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2 + \lambda \|\mathbf{x}\|_2^2, \quad \text{for some constant } \lambda > 0.$$

- ℓ_2 -norm enforces total smoothness
- solution by optimization: $\nabla f(\mathbf{x}) = 2\mathbf{A}^T \mathbf{A}\mathbf{x} - 2\mathbf{A}^T \mathbf{y} + 2\lambda \mathbf{x}$. Thus the optimal solution is

$$\mathbf{x}_{\text{RLS}} = (\mathbf{A}^T \mathbf{A} + \lambda \mathbf{I})^{-1} \mathbf{A}^T \mathbf{y}$$

- solution by the projection thm., in contrast: have to rewrite the problem as

$$\min_{\mathbf{x} \in \mathbb{R}^n} \left\| \begin{bmatrix} \mathbf{y} \\ \mathbf{0} \end{bmatrix} - \begin{bmatrix} \mathbf{A} \\ \sqrt{\lambda} \mathbf{I} \end{bmatrix} \mathbf{x} \right\|_2^2,$$

and use the projection theorem to get the same result.

- LS with Tikhonov regularization is commonly used for solving underdeter. linear systems; it can make ill-conditioned (i.e., rank-deficient or close-to-rank-deficient) LS problem to be well-conditioned; LS + Tikhonov reg. = ridge regression model
- if there are \mathbf{x} that satisfy $\mathbf{A}\mathbf{x} = \mathbf{b}$, this will chose the solution with least norm

How to obtain the Solution to a LS?

- direct methods for solving LS
 - method of normal equations
 - QR decomposition
- iterative methods for solving LS
 - gradient descent
 - coordinate descent
 - more...

Direct Methods via Method of Normal Equations

- for LS problem with $\mathbf{A} \in \mathbb{R}^{m \times n}$ of full column rank, the solution is

$$\mathbf{x}_{\text{LS}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

- naïve computation, complexity: $\mathcal{O}(mn^2 + n^3)$

- For example: solving LS via the normal equations

$$\mathbf{A}^T \mathbf{A} \mathbf{x}_{\text{LS}} = \mathbf{A}^T \mathbf{y}$$

- solving the above symmetric PD linear system (cf. [Linear Systems Topic](#))
 - * compute the lower triangular portion of $\mathbf{C} = \mathbf{A}^T \mathbf{A}$, $\mathcal{O}(mn^2)$
 - * form the matrix-vector product $\mathbf{d} = \mathbf{A}^T \mathbf{y}$, $\mathcal{O}(mn)$
 - * compute the Cholesky factorization $\mathbf{C} = \mathbf{G}\mathbf{G}^T$, $\mathcal{O}(n^3/3)$
 - * solve $\mathbf{G}\mathbf{z} = \mathbf{d}$ and $\mathbf{G}^T \mathbf{x}_{\text{LS}} = \mathbf{z}$, $\mathcal{O}(n^2)$

Direct Methods via QR Decomposition

- LS can be solved by the QR decompositions
- will discuss this in detail in [Topic: Orthogonalization and QR Decomposition](#)

Iterative Methods for Solving LS

- in the direct methods for solving LS, we need to solve

$$(\mathbf{A}^T \mathbf{A}) \mathbf{x}_{\text{LS}} = \mathbf{A}^T \mathbf{y},$$

and that requires $\mathcal{O}(n^3)$

– we also need to compute $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{y}$; their complexities are $\mathcal{O}(mn^2)$ and $\mathcal{O}(mn)$, resp.

- $\mathcal{O}(n^3)$ is expensive for very large n
- **Question:** can we have cheaper LS solutions, perhaps with some compromise of the solution accuracies?

Gradient Descent

- consider a general unconstrained optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

where f is continuously differentiable

- **Gradient Descent (GD)**: given a starting point $\mathbf{x}^{(0)}$, do

$$\mathbf{x}^{(k)} = \mathbf{x}^{(k-1)} - \mu \nabla f(\mathbf{x}^{(k-1)}), \quad k = 1, 2, \dots$$

where $\mu > 0$ is a step size

- take an optimization course to get more details! It is known that
 - for convex f and under some appropriate choice of μ , GD converges to an optimal solution
 - for non-convex f and under some appropriate choice of μ , GD converges to a stationary point

Gradient Descent

- GD for LS:

$$\mathbf{x}^{(k)} = \mathbf{x}^{(k-1)} - 2\mu(\mathbf{A}^T \mathbf{A} \mathbf{x}^{(k-1)} - \mathbf{A}^T \mathbf{y}), \quad k = 1, 2, \dots$$

- complexity for dense \mathbf{A}
 - computing $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{y}$: $\mathcal{O}(mn^2)$ and $\mathcal{O}(mn)$, resp. (same as before)
 - * $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{y}$ are cached for subsequent use in gradient descent
 - complexity of each iteration: $\mathcal{O}(n^2)$
- complexity for sparse \mathbf{A} (solving (large) sparse LS is an important topic)
 - computing $\mathbf{A}^T \mathbf{y}$: $\mathcal{O}(\text{nnz}(\mathbf{A}))$
 - complexity of each iteration: $\mathcal{O}(n + \text{nnz}(\mathbf{A}))$
 - * $\mathbf{A}^T \mathbf{A}$ is not necessarily sparse, so we do $\mathbf{A} \mathbf{x}^{(k-1)}$ and then $\mathbf{A}^T(\mathbf{A} \mathbf{x}^{(k-1)})$

Gradient Descent

- gradient descent is easy to understand, but there are better algorithms...
- further reading: the conjugate gradient method; see, e.g.,
https://stanford.edu/class/ee364b/lectures/conj_grad_slides.pdf

Online LS

- let $\bar{\mathbf{a}}_i \in \mathbb{R}^n$ denote the i th row of \mathbf{A} , then

$$\|\mathbf{A}\mathbf{x} - \mathbf{y}\|_2^2 = \sum_{t=1}^m |\bar{\mathbf{a}}_t^T \mathbf{x} - y_t|^2$$

- the LS formulation can be written as

$$\min_{\mathbf{x} \in \mathbb{R}^n} \sum_{t=1}^m |\bar{\mathbf{a}}_t^T \mathbf{x} - y_t|^2$$

- the LS we learnt is a batch process; i.e., solve one \mathbf{x} given the whole (\mathbf{A}, \mathbf{y}) ; the afore-mentioned GD method is also hence referred to as batch GD
- there are many applications where new $(\bar{\mathbf{a}}_t, y_t)$ appears as time goes, and we want the process to be adaptive or in real time; i.e., \mathbf{x} is updated with t
- alternatively, we want something cheaper than gradient descent

Incremental Gradient Descent

- consider an optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} \sum_{t=1}^m f_t(\mathbf{x})$$

where every f_t is continuously differentiable

- Incremental Gradient Descent:

$$\mathbf{x}_t = \mathbf{x}_{t-1} - \mu \nabla f_t(\mathbf{x}_{t-1}), \quad t = 1, 2, \dots$$

– also called online gradient descent, stochastic gradient descent (SGD), least mean squares (LMS) (in 70's), ...

- incremental gradient descent for LS:

$$\mathbf{x}_t = \mathbf{x}_{t-1} + 2\mu(y_t - \bar{\mathbf{a}}_t^T \mathbf{x}_{t-1})\bar{\mathbf{a}}_t$$

- complexity: $\mathcal{O}(n)$
 - commonly used in large-scale optimization like learning neural networks

Recursive LS

- Recursive LS (RLS) formulation:

$$\mathbf{x}_t = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \sum_{i=1}^t \lambda^{t-i} |\bar{\mathbf{a}}_i^T \mathbf{x} - y_i|^2$$

where $0 < \lambda \leq 1$ is a prescribed constant and is called the forgetting or discounting factor

- weigh the importance of $|\bar{\mathbf{a}}_i^T \mathbf{x} - y_i|^2$ w.r.t. time t ; the present is most important; distant pasts are insignificant; how much we remember the pasts depends on λ

- at first look, the RLS solution is $\mathbf{x}_t = \mathbf{R}_t^{-1} \mathbf{q}_t$, where

$$\mathbf{R}_t = \sum_{i=1}^t \lambda^{t-i} \bar{\mathbf{a}}_i \bar{\mathbf{a}}_i^T, \quad \mathbf{q}_t = \sum_{i=1}^t \lambda^{t-i} y_i \bar{\mathbf{a}}_i$$

- a recursive formula for \mathbf{x}_t can be derived by using the Woodbury matrix identity and by using the problem structures carefully

Woodbury Matrix Identity

For $\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}$ of appropriate dimensions, we have

$$(\mathbf{A} - \mathbf{B}\mathbf{D}^{-1}\mathbf{C})^{-1} = \mathbf{A}^{-1} + \mathbf{A}^{-1}\mathbf{B}(\mathbf{D} - \mathbf{C}\mathbf{A}^{-1}\mathbf{B})^{-1}\mathbf{C}\mathbf{A}^{-1},$$

assuming that the inverses above exist.

- for the RLS problem, it is sufficient to know the special case

$$(\mathbf{A} + \mathbf{b}\mathbf{b}^T)^{-1} = \mathbf{A}^{-1} - \frac{1}{1 + \mathbf{b}^T\mathbf{A}^{-1}\mathbf{b}}\mathbf{A}^{-1}\mathbf{b}\mathbf{b}^T\mathbf{A}^{-1}$$

Recursive LS

- it can be verified that $\mathbf{R}_t = \lambda \mathbf{R}_{t-1} + \bar{\mathbf{a}}_t \bar{\mathbf{a}}_t^T$, $\mathbf{q}_t = \lambda \mathbf{q}_{t-1} + y_t \bar{\mathbf{a}}_t$
- by the Woodbury matrix identity,

$$\mathbf{R}_t^{-1} = (\lambda \mathbf{R}_{t-1} + \bar{\mathbf{a}}_t \bar{\mathbf{a}}_t^T)^{-1} = \frac{1}{\lambda} \mathbf{R}_{t-1}^{-1} - \frac{1}{1 + \frac{1}{\lambda} \bar{\mathbf{a}}_t^T \mathbf{R}_{t-1}^{-1} \bar{\mathbf{a}}_t} \left(\frac{1}{\lambda} \mathbf{R}_{t-1}^{-1} \bar{\mathbf{a}}_t \right) \left(\frac{1}{\lambda} \mathbf{R}_{t-1}^{-1} \bar{\mathbf{a}}_t \right)^T$$

- let $\mathbf{P}_t = \mathbf{R}_t^{-1}$ and $\mathbf{g}_t = \frac{1}{1 + \frac{1}{\lambda} \bar{\mathbf{a}}_t^T \mathbf{R}_{t-1}^{-1} \bar{\mathbf{a}}_t} \left(\frac{1}{\lambda} \mathbf{R}_{t-1}^{-1} \bar{\mathbf{a}}_t \right)$. We have

$$\mathbf{g}_t = \frac{1}{1 + \frac{1}{\lambda} \bar{\mathbf{a}}_t^T \mathbf{P}_{t-1} \bar{\mathbf{a}}_t} \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)$$

$$\mathbf{P}_t = \frac{1}{\lambda} \mathbf{P}_{t-1} - \mathbf{g}_t \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)^T$$

$$\begin{aligned} \mathbf{x}_t = \mathbf{P}_t \mathbf{q}_t &= \mathbf{P}_{t-1} \mathbf{q}_{t-1} - \lambda \mathbf{g}_t \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)^T \mathbf{q}_{t-1} + \frac{1}{\lambda} y_t \mathbf{P}_{t-1} \bar{\mathbf{a}}_t - y_t \mathbf{g}_t \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)^T \bar{\mathbf{a}}_t \\ &= \mathbf{x}_{t-1} - (\bar{\mathbf{a}}_t^T \mathbf{x}_{t-1}) \mathbf{g}_t + y_t \mathbf{g}_t \\ &= \mathbf{x}_{t-1} + (y_t - \bar{\mathbf{a}}_t^T \mathbf{x}_{t-1}) \mathbf{g}_t \end{aligned}$$

Recursive LS

- summary of the RLS recursion:

$$\mathbf{g}_t = \frac{1}{1 + \frac{1}{\lambda} \bar{\mathbf{a}}_t^T \mathbf{P}_{t-1} \bar{\mathbf{a}}_t} \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)$$

$$\mathbf{P}_t = \frac{1}{\lambda} \mathbf{P}_{t-1} - \mathbf{g}_t \left(\frac{1}{\lambda} \mathbf{P}_{t-1} \bar{\mathbf{a}}_t \right)^T$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + (y_t - \bar{\mathbf{a}}_t^T \mathbf{x}_{t-1}) \mathbf{g}_t$$

- complexity: $\mathcal{O}(n)$
- remarks:
 - comparison with incremental gradient descent: it replaces \mathbf{g}_t with $2\mu \bar{\mathbf{a}}_t$
 - the above RLS recursion may be numerically unstable as empirical results suggested (further reading: [\[Liavas-Regalia'99\]](#)); modified RLS schemes were developed to mend this issue

Coordinate Descent

- the problem is to solve

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

```
input: a starting point  $\mathbf{x}^{(0)}$ 
for  $k = 0, 1, 2, \dots$ 
     $x_1^{(k+1)} = \arg \min_{x_1 \in \mathbb{R}} f(x_1, x_2^{(k)}, \dots, x_n^{(k)})$ 
     $x_2^{(k+1)} = \arg \min_{x_2 \in \mathbb{R}} f(x_1^{(k+1)}, x_2, \dots, x_n^{(k)})$ 
     $\vdots$ 
     $x_n^{(k+1)} = \arg \min_{x_n \in \mathbb{R}} f(x_1^{(k+1)}, x_2^{(k+1)}, \dots, x_n)$ 
end
```

- a.k.a. nonlinear Gauss-Seidel
- It is known that **[Tseng'01]**
 - Convergence guarantees toward a local optimal (minimal) point for smooth functions or separable functions
 - No convergence toward a minimum for non-separable and non-smooth functions: some points (coordinatewise minimum) get stuck

Coordinate Descent

- CD for LS. notice

$$f(\mathbf{x}) = \|\mathbf{Ax} - \mathbf{y}\|_2^2 = \sum_{i=1}^m \|\mathbf{a}_i x_i - \mathbf{y}\|_2^2 = f(x_1, \dots, x_n)$$

and

$$\nabla f(\mathbf{x}) = 2\mathbf{A}^T(\mathbf{Ax} - \mathbf{y}) = 2 \begin{bmatrix} \mathbf{a}_1^T(\mathbf{Ax} - \mathbf{y}) \\ \vdots \\ \mathbf{a}_n^T(\mathbf{Ax} - \mathbf{y}) \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

- minimize w.r.t x_i for $i = 1, \dots, n$, with fixed x_j ($j \neq i$)

$$0 = \frac{\partial f}{\partial x_i} = \mathbf{a}_i^T(\mathbf{Ax} - \mathbf{y}) = \mathbf{a}_i^T(\mathbf{a}_i x_i + \sum_{j \neq i} \mathbf{a}_j x_j - \mathbf{y})$$

we have

$$x_i = \frac{\mathbf{a}_i^T(\mathbf{y} - \sum_{j \neq i} \mathbf{a}_j x_j)}{\|\mathbf{a}_i\|_2^2}$$

Coordinate Descent

```
input: a starting point  $\mathbf{x}^{(0)}$ 
for  $k = 0, 1, 2, \dots$ 
  for  $i = 1, 2, \dots, n$ 
    
$$x_i^{(k+1)} = \frac{\mathbf{a}_i^T (\mathbf{y} - \sum_{j=1}^{i-1} \mathbf{a}_j x_j^{(k+1)} - \sum_{j=i+1}^n \mathbf{a}_j x_j^{(k)})}{\|\mathbf{a}_i\|_2^2}$$

  end
end
```

- equivalent to solving linear system $\mathbf{A}^T \mathbf{A} \mathbf{x} = \mathbf{A}^T \mathbf{y}$ via CD (cf. [Linear Sys. Topic](#))
- clever update scheme can be developed with low memory impact
- CD is a [non-gradient optimization](#) (or a [derivative-free optimization](#)) method; it can be extremely fast
- possibly visit the coordinate in arbitrary order (cycle, random, more refined methods, etc.)
- [Block coordinate descent \(BCD\)](#): update not only one coordinate, but a bunch of them, i.e., optimizing according to problem architecture; a.k.a. block nonlinear Gauss-Seidel

More Iterative Methods for LS

- accelerated GD
- conjugate GD
- Newton's method
- Gauss-Newton method (line search)
- the Levenberg-Marquardt method (trust region)
- ...

Beyond LS


- The LS problem can be represented as

$$\min_{\mathbf{x} \in \mathbb{R}^n} \sum_{i=1}^m \ell(\bar{\mathbf{a}}_i^T \mathbf{x} - y_i)$$


where $\ell(z) = |z|^2$ denotes the **loss function** for measuring the badness of fit

- **Question:** why don't we use other loss functions?
 - we can indeed use other loss functions, such as
 - * 1-norm loss: $\ell(z) = |z|$ (least absolute deviations (LAD))
 - * Huber loss: $\ell(z) = \begin{cases} \frac{1}{2}|z|^2, & |z| \leq 1 \\ |z| - \frac{1}{2}, & |z| > 1 \end{cases}$
 - * power- p loss: $\ell(z) = |z|^p$, with $p < 1$
 - corresponding to different prior distributions for noise \mathbf{v} in $\mathbf{y} = \mathbf{A}\mathbf{x} + \mathbf{v}$
 - the above loss functions are more robust against outliers, but
 - they require optimization and don't result in a clean closed-form solution as LS (a method to solve them is **iteratively reweighted least squares (IRLS)**; in each iteration, a (weighted) LS is solved via successive linear approximation (SLA))

Illustration of Loss Functions



Curve Fitting Example



“True” curve: the true $f(x)$, $p = 5$. The points at $x = -0.3$ and $x = 0.4$ are outliers, and they do not follow the true curve. The 1-norm loss problem is solved by a convex optimization tool.

More on LS

more topics related to LS in future lectures (cf. [LS Revisited and Sparse Opt. Topic](#))

- linear LS (ordinary, weighted, generalized...) vs. nonlinear LS (neural networks...)
- regularized LS
 - penalized LS (e.g., ℓ_0 /best subset, ℓ_1 /lasso, ℓ_2 /ridge, $\ell_1+\ell_2$ /elastic net, ...)
 - constrained LS (e.g., non-negative LS, bounded-variable LS, linearly constrained LS, LS with simplex constraints, LS with norm ball constraint, ...)
- underdetermined linear system of equations
 - find the minimum ℓ_2 solution of an underdetermined linear system
 - find the minimum ℓ_0 solution of an underdetermined linear system
 - find the minimum ℓ_1 solution of an underdetermined linear system
 - majorization-minimization for ℓ_2 - ℓ_1 minimization
 - dictionary learning and frame learning
- LS with errors in \mathbf{A}
 - total LS
 - robust LS, and its equivalence to regularized LS

Part III: Matrix Factorization

Matrix Factorization

There are also many applications in which we deal with a representation of multiple given \mathbf{y}_i 's via

$$\mathbf{y}_i = \mathbf{A}\mathbf{b}_i + \mathbf{v}_i, \quad i = 1, \dots, n,$$

where $\mathbf{A} \in \mathbb{R}^{n \times k}$, $\mathbf{b}_i \in \mathbb{R}^k$, $i = 1, \dots, n$; \mathbf{v}_i 's are noise. In particular, both \mathbf{b}_i 's and \mathbf{A} are to be determined.

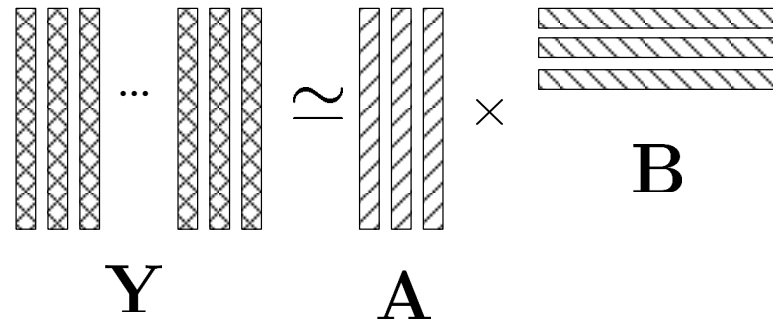
- for example, in basis representation, we want to jointly learn the dictionary from data

Matrix Factorization

Problem: given $\mathbf{Y} \in \mathbb{R}^{m \times n}$ and a positive integer $k < \min\{m, n\}$, solve

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times k}, \mathbf{B} \in \mathbb{R}^{k \times n}} \|\mathbf{Y} - \mathbf{AB}\|_F^2$$

where $\|\mathbf{Y} - \mathbf{AB}\|_F^2 = \sum_{i=1}^n \|\mathbf{y}_i - \mathbf{A}\mathbf{b}_i\|_2^2 = \sum_{i=1}^m \|\bar{\mathbf{y}}_i - \mathbf{B}^T \bar{\mathbf{a}}_i\|_2^2 = \sum_{i,j} |y_{ij} - [\mathbf{AB}]_{ij}|^2$ with $\bar{\mathbf{y}}_i \in \mathbb{R}^n, \bar{\mathbf{a}}_i \in \mathbb{R}^k$ denoting the i th row of \mathbf{Y}, \mathbf{A} , respectively.



$$\mathbf{Y} \approx \mathbf{A} \mathbf{B}$$

- matrix factorization (MF) is also called **low-rank matrix factorization** or **low-rank matrix approximation**: let $\mathbf{Z} = \mathbf{AB}$. It has $\text{rank}(\mathbf{Z}) \leq k$.
- like in LS, we may often want to add constraints and/or penalties in MF problems, like orthogonality constraint, non-negative constraint, linear constraint, sparsity constraint, etc.

Principal Component Analysis

Aim: given a collection of data points $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^m$, perform a low-dimensional representation

$$\mathbf{y}_i = \mathbf{A}\mathbf{b}_i + \mathbf{c} + \mathbf{v}_i, \quad i = 1, \dots, n,$$


where $\mathbf{A} \in \mathbb{R}^{m \times k}$ is a basis matrix; $\mathbf{b}_i \in \mathbb{R}^k$ is the coefficient for \mathbf{y}_i ; $\mathbf{c} \in \mathbb{R}^m$ is the base or mean in statistics terms; \mathbf{v}_i is noise or modeling error.

- Principal component analysis (PCA):

- choose $\mathbf{c} = \frac{1}{n} \sum_{i=1}^n \mathbf{y}_i$
- let $\bar{\mathbf{y}}_i = \mathbf{y}_i - \mathbf{c}$, and solve

$$\min_{\mathbf{A}, \mathbf{B}} \|\bar{\mathbf{Y}} - \mathbf{A}\mathbf{B}\|_F^2$$

- we may also want a semi-orthogonal \mathbf{A}
- in PCA problem, minimizing squared distances equals maximizing variance



Principal Component Analysis

- applications: dimensionality reduction, visualization of high-dimensional data, compression, extraction of meaningful features from data,...
- an example:
 - senate voting: <http://livebooklabs.com/keepies/c5a5868ce26b8125>


Topic Modeling

Aim: discover thematic information, or topics, from a (often large) collection of documents, such as books, articles, news, blogs,...

- **bag-of-words representation:** represent each document as a vector of word counts

... In fact, we will soon see that the **implementation** of **SDR** can be very easy, which allows **signal processing** practitioners to quickly test the viability of **SDR** in their applications. Several highly successful **applications** will be showcased as **examples**

a document



bag of words



$y =$

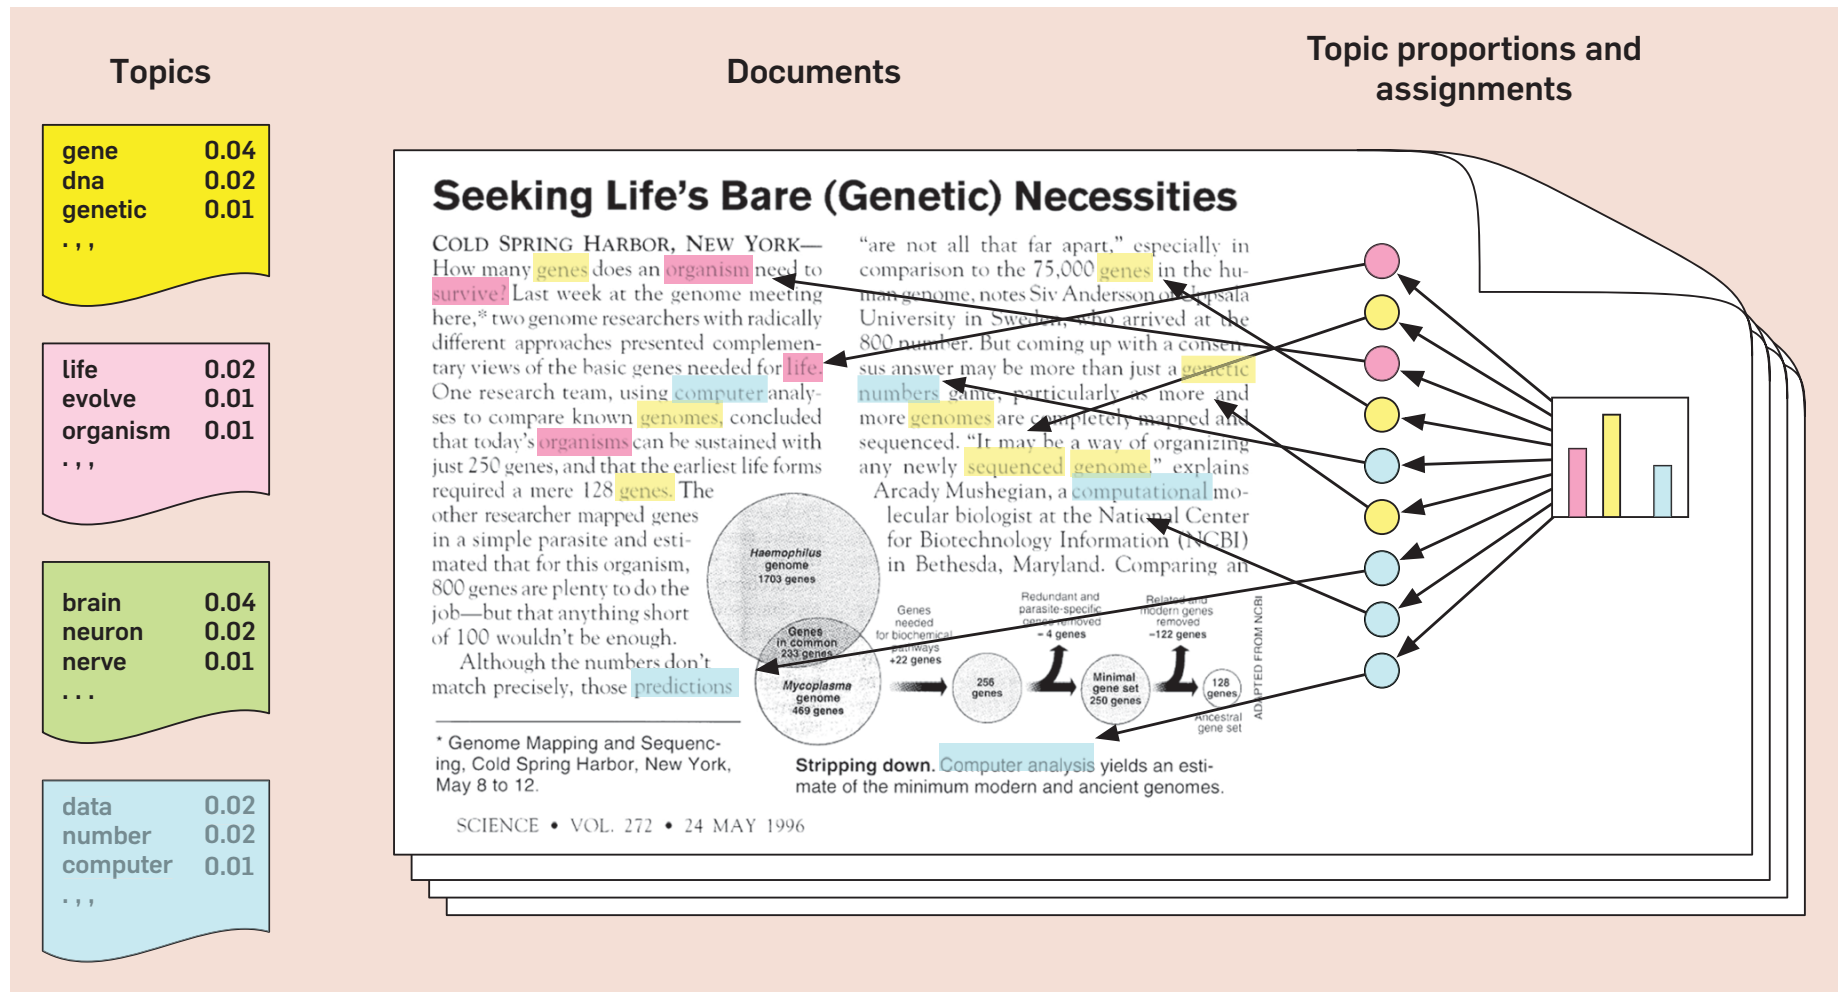
count	term
0	efficiency
2	applications
2	SDR
0	communications
1	example
1	signal processing
⋮	⋮
1	implementation

bag-of-words representation

Topic Modeling

- let n be the number of documents
- let $\mathbf{y}_i \in \mathbb{R}^m$ be the bag-of-words representation of the i th document, $i = 1, \dots, n$
- let $\mathbf{Y} = [\mathbf{y}_1, \dots, \mathbf{y}_n] \in \mathbb{R}^{m \times n}$, called the term-document matrix
- hypotheses: **[Turney-Pantel'10]**
 - if documents have similar columns vectors in \mathbf{Y} , or similar usage of words, they tend to have similar meanings
 - the topic of a document will probabilistically influence the author's choice of words when writing the document

Topic Modeling



Source: [Blei'12].

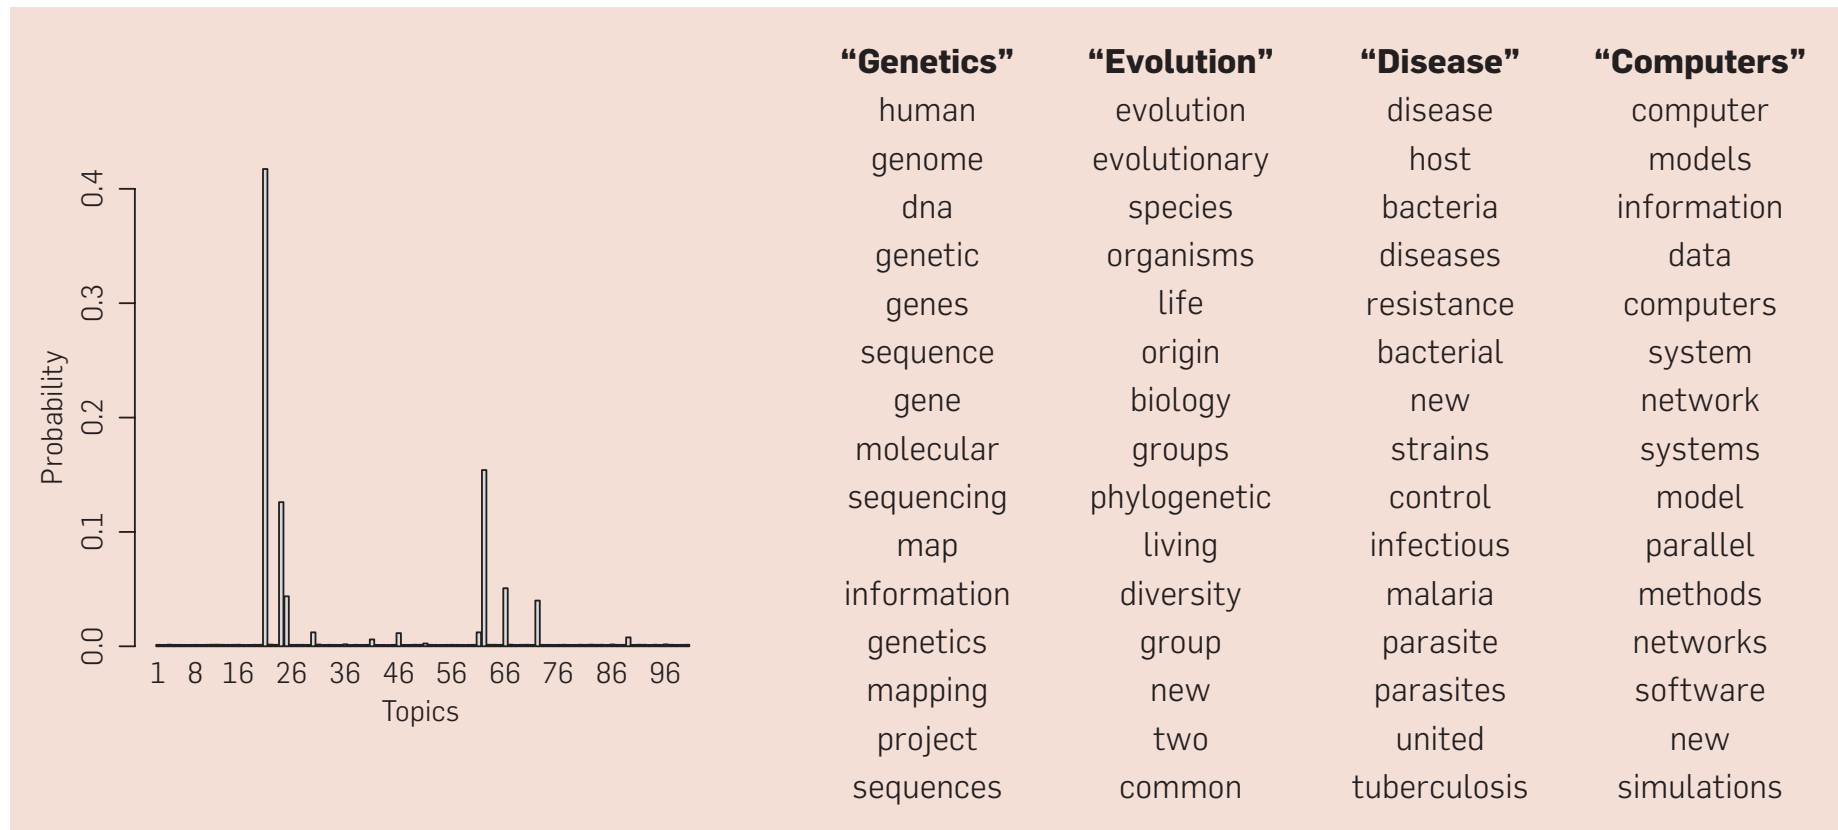
Topic Modeling

- **Problem:** apply matrix factorization to a term-document matrix \mathbf{Y}

$$\mathbf{Y} \approx \mathbf{A} \times \mathbf{B}$$

- \mathbf{A} is called a term-topic matrix, \mathbf{B} is called a topic-document matrix
- **Interpretation:**
 - each column \mathbf{a}_i of \mathbf{A} should represent a theme topic, e.g., local affairs, foreign affairs, politics, sports... in a collection of newspapers
 - as $\mathbf{y}_i \approx \mathbf{A}\mathbf{b}_i$, each document is postulated as a linear combination of topics
 - matrix factorization aims at discovering topics from the documents

Topic Modeling



Topics found in a real set of documents. Source: [\[Blei'12\]](#). The document set consists of 17,000 articles from the journal *Science*. The topics are discovered using a technique called *latent Dirichlet allocation*, which is not the same as, but has strong connections to, matrix factorization.

Topic Modeling

- topic modeling via matrix factorization has been used in, or is tightly connected to
 - information retrieval, natural language processing, machine learning
 - document clustering, classification and retrieval
 - latent semantic analysis, latent semantic indexing: finding similarities of documents, finding similarities of terms (are “cars,” “Lamborghini,” and “Ferrari” related?)
- though not considered in this course, it seems better to also model \mathbf{A} , \mathbf{B} as element-wise non-negative—this will lead to *non-negative matrix factorization*
- further reading: [\[Turney-Pantel’10\]](#)
 - as an aside, it mentions a related application where computers can achieve a score of 92.5% on multiple-choice synonym questions from TOEFL, whereas the average human score is 64.5%

Matrix Factorization

The matrix factorization problem

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times k}, \mathbf{B} \in \mathbb{R}^{k \times n}} \|\mathbf{Y} - \mathbf{AB}\|_F^2$$

- has non-unique factors
 - suppose $(\mathbf{A}^*, \mathbf{B}^*)$ is an optimal solution to the problem, and let $\mathbf{Q} \in \mathbb{R}^{k \times k}$ be any nonsingular matrix. Then $(\mathbf{A}^* \mathbf{Q}^{-1}, \mathbf{Q} \mathbf{B}^*)$ is also an optimal solution.
 - the non-uniqueness of (\mathbf{A}, \mathbf{B}) makes the above matrix factorization formulation a bad formulation for problems such as topic modeling
- is non-convex, but can be solved by singular value decomposition (beautifully) (cf. [Singular Value Decomposition Topic](#))
- can also be handled by LS

Matrix Factorization

- **Alternating LS (ALS):** given a starting point $(\mathbf{A}^{(0)}, \mathbf{B}^{(0)})$, do

$$\mathbf{A}^{(i+1)} = \arg \min_{\mathbf{A} \in \mathbb{R}^{m \times k}} \|\mathbf{Y} - \mathbf{A}\mathbf{B}^{(i)}\|_F^2$$

$$\mathbf{B}^{(i+1)} = \arg \min_{\mathbf{B} \in \mathbb{R}^{k \times n}} \|\mathbf{Y} - \mathbf{A}^{(i+1)}\mathbf{B}\|_F^2$$

for $i = 0, 1, 2, \dots$, and stop when a stopping rule is satisfied.

- let's make a mild assumption that $\mathbf{A}^{(i)}, \mathbf{B}^{(i)}$ have full rank at every i . Then,
$$\mathbf{A}^{(i+1)} = \mathbf{Y}(\mathbf{B}^{(i)})^T(\mathbf{B}^{(i)}(\mathbf{B}^{(i)})^T)^{-1}, \quad \mathbf{B}^{(i+1)} = ((\mathbf{A}^{(i+1)})^T\mathbf{A}^{(i+1)})^{-1}(\mathbf{A}^{(i+1)})^T\mathbf{Y}$$
- a special case of **alternating minimization (AM, AltMin)** or **BCD**
- ALS is guaranteed to converge an optimal solution to $\min_{\mathbf{A}, \mathbf{B}} \|\mathbf{Y} - \mathbf{A}\mathbf{B}\|_F^2$ under some mild assumptions **[Udell-Horn-Zadeh-Boyd'16]**
 - note: this result is specific and does not directly carry forward to other related problems such as low-rank matrix completion
- you can also apply GD, SGD, alternating GD, etc.

Low-Rank Matrix Completion

- let $\mathbf{Y} \in \mathbb{R}^{m \times n}$ be a matrix with missing entries, i.e., the values y_{ij} 's are known only for $(i, j) \in \Omega$ where Ω is an index set that indicates the available entries
- **Aim:** recover the missing entries of \mathbf{Y}
- application: recommender system, data science
- example: movie recommendation (further reading: [\[Koren-Bell-Volinsky'09\]](#))
 - \mathbf{Y} records how user i likes movie j
 - \mathbf{Y} has lots of missing entries; a user doesn't watch all movies
 - \mathbf{Y} may be assumed to have low rank; research shows that only a few factors affect users' preferences.

$$\mathbf{Y} = \begin{matrix} & \text{movies} \\ \begin{bmatrix} 2 & 3 & 1 & ? & ? & 5 & 5 \\ 1 & ? & 4 & 2 & ? & ? & ? \\ ? & 3 & 1 & ? & 2 & 2 & 2 \\ ? & ? & ? & 3 & ? & 1 & 5 \end{bmatrix} & \text{users} \end{matrix}$$

Low-Rank Matrix Completion

- **Problem:** given $\{y_{ij}\}_{(i,j) \in \Omega}$, Ω and a positive integer k , solve

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times k}, \mathbf{B} \in \mathbb{R}^{k \times n}} \sum_{(i,j) \in \Omega} |y_{ij} - [\mathbf{AB}]_{ij}|^2$$

- ALS can be applied; more tedious to write out the LS solutions than the previous matrix factorization problem but not any harder in principle
- supposingly a very difficult problem, but
- methods like ALS were found to work by means of empirical studies
- recent theoretical research suggests that matrix completion may not be that hard under some assumptions, e.g., ALS can give good results [\[Sun-Luo'16\]](#)

Low-Rank Matrix Completion

- an ALS alternative to matrix completion (easier to program):
 - consider an equivalent reformulation of the matrix completion problem

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times k}, \mathbf{B} \in \mathbb{R}^{k \times n}, \mathbf{R} \in \mathbb{R}^{m \times n}} \|\mathbf{Y} + \mathbf{R} - \mathbf{AB}\|_F^2 \quad \text{s.t. } r_{ij} = 0, (i, j) \in \Omega$$

- do alternating optimization

$$\mathbf{A}^{(i+1)} = \arg \min_{\mathbf{A} \in \mathbb{R}^{m \times k}} \|\mathbf{Y} - \mathbf{AB}^{(i)} + \mathbf{R}^{(i)}\|_F^2$$


$$\mathbf{B}^{(i+1)} = \arg \min_{\mathbf{B} \in \mathbb{R}^{k \times n}} \|\mathbf{Y} - \mathbf{A}^{(i+1)}\mathbf{B} + \mathbf{R}^{(i)}\|_F^2$$

$$\mathbf{R}^{(i+1)} = \arg \min_{\mathbf{R} \in \mathbb{R}^{m \times n}} \|\mathbf{Y} - \mathbf{A}^{(i+1)}\mathbf{B}^{(i+1)} + \mathbf{R}\|_F^2$$

the first two are LS as before; the third has a closed form

$$r_{ij}^{(i+1)} = \begin{cases} 0, & (i, j) \in \Omega \\ -[\mathbf{Y} - \mathbf{A}^{(i+1)}\mathbf{B}^{(i+1)}]_{i,j}, & (i, j) \notin \Omega \end{cases}$$

Toy Demonstration of Low-Rank Matrix Completion



Left: An incomplete image with 40% missing pixels. Right: the matrix completion result of the algorithm shown on last page. $k = 120$.

References

- [Stoica-Moses'97]** P. Stoica and R. L. Moses, *Introduction to Spectral Analysis*, Prentice Hall, 1997.
- [Aharon-Elad-Bruckstein'06]** M. Aharon, M. Elad, and A. Bruckstein, " K -SVD: An algorithm for designing overcomplete dictionaries for sparse representation," *IEEE Trans. Image Process.*, vol. 54, no. 11, pp. 4311–4322, 2006.
- [Golub-Van Loan'13]** G. H. Golub and C. F. Van Loan, *Matrix Computations*, 4th edition, JHU Press, 2013.
- [Sayed-Tarighat-Khajehnouri'05]** A. H. Sayed, A. Tarighat, and N. Khajehnouri. "Network-based wireless location," *IEEE Signal Process. Mag.*, vol. 22, no. 4, pp. 24–40, 2005.
- [Tseng'01]** P. Tseng. "Convergence of a block coordinate descent method for nondifferentiable minimization," *J. Optim. Theory Appl.*, vol. 109, no. 3, pp. 475–494, 2001.
- [Turney-Pantel'10]** P. D. Turney and P. Pantel, "From frequency to meaning: Vector space models of semantics," *Journal of Artificial Intelligence Research*, vol. 37, pp. 141–188, 2010.
- [Blei'12]** D. Blei, "Probabilistic topic models," *Communications of the ACM*, vol. 55, no. 4, pp. 77–84, 2012.

[Udell-Horn-Zadeh-Boyd'16] M. Udell, C. Horn, R. Zadeh, and S. Boyd, “Generalized low rank models”, *Foundations and Trends in Machine Learning*, 2016.

[Koren-Bell-Volinsky'09] B. Koren, R. Bell, and C. Volinsky, “Matrix factorization techniques for recommender systems,” *IEEE Computer*, vol. 42 no. 8, pp. 30–37, 2009.

[Sun-Luo'16] R. Sun and Z.-Q. Luo, “Guaranteed matrix completion via non-convex factorization.” *IEEE Trans. Inform. Theory*, vol. 62, no. 11, pp. 6535–6579, 2016.

[Liavas-Regalia'99] A. P. Liavas and P. A. Regalia, “On the numerical stability and accuracy of the conventional recursive least squares algorithm.” *IEEE Trans. Signal Process.*, vol. 47, no. 1, pp. 88-96, 1999.